

BIOS 7747: Machine Learning for Biomedical Applications

Feature exploration, pre-processing and normalization

Antonio R. Porras (antonio.porras@cuanschutz.edu)

Department of Biostatistics and Informatics

Colorado School of Public Health

University of Colorado Anschutz Medical Campus

Outline

- ❑ Feature exploration and data cleaning
- ❑ Feature distributions
- ❑ Feature interactions
- ❑ Mitigating outlier effects and multicollinearity
- ❑ Feature scaling

Feature exploration and data cleaning

□ Types of features:

Numerical

- Continuous
- Discrete

Categorical

- Ordinal
- Nominal



Feature exploration and data cleaning

❑ Encoding categorical variables

- Label encoding (`sklearn.preprocessing.LabelEncoder`)

Cancer stage	Encoded value
Stage I	0
Stage II	1
Stage III	2

Often not recommended


- Distances cannot be assumed in most ordinal variables
- Order cannot be assumed in nominal variables

Feature exploration and data cleaning

□ Encoding categorical variables

- One hot encoding (`sklearn.preprocessing.OneHotEncoder`)

Sample	Cancer stage
0	Stage III
1	Stage I
2	Stage II



Sample	Stage I	Stage II	Stage III
0	0	0	1
1	1	0	0
2	0	1	0

It can increase substantially the number of features:

$$M = M_{numerical} + \sum_{\forall i \in categorical} N_{classes}^i$$

Feature exploration and data cleaning

□ Encoding categorical variables

- Binary encoding: every feature is coded as a binary number with a fixed number of digits. Each digit is a feature to consider in the model.

Sample	Cancer stage
0	Stage III
1	Stage I
2	Stage II



Sample	Feature 1	Feature 2
0	1	0
1	0	0
2	0	1

It requires less features than one-hot-encoding

$$M = M_{numerical} + \sum_{\forall i \in \text{categorical}} \log_2(N_{classes}^i)$$

Feature exploration and data cleaning

□ Encoding categorical variables

- Feature hashing: Convert labels to “words” with predefined fixed size

1	A,B,C,D,E,F,G		1	2	3	4
2	H,I,J,K,L,M,N	Mary →	1	1	1	1
3	O,P,Q,R,S,T	John →	0	3	1	0
4	U,V,W,X,Y,Z	Jennifer →	3	4	1	0

Great to standardize representations using low number of variables

Feature collision may happen

Feature exploration and data cleaning

❑ Missing values

- Common reasons:
 - Incomplete data (e.g., prefer not to answer, data transfer errors)
 - Human error (e.g., forgot to annotate, incorrect annotations)
 - Study design (e.g., data does not apply)
- Classification [D.B. Rubin, 1976]
 - Missing completely at random (MCAR): missing values are not related to the observations (the probability of having missing value is equation for all samples). Unusual in biomedicine.
 - Missing at random (MAR): the probability of missing value is a function of another variable (e.g., male are less likely to answer about mental health questions in a survey).
 - Missing not at random (MNAR): there is no insight about the probability of missing data.

Feature exploration and data cleaning

❑ Missing values

- **Sample dropping** (Pandas' *dropna()* function):
 - MCAR: it may not affect predictions when “sufficient” data are available
 - But it may result in insufficient data
 - MAR: it can introduce biases that could potentially be identified
 - Example: male do not have mental illness problems
 - MNAR: it can introduce biases that are hard to identify

Feature exploration and data cleaning

❑ Missing values

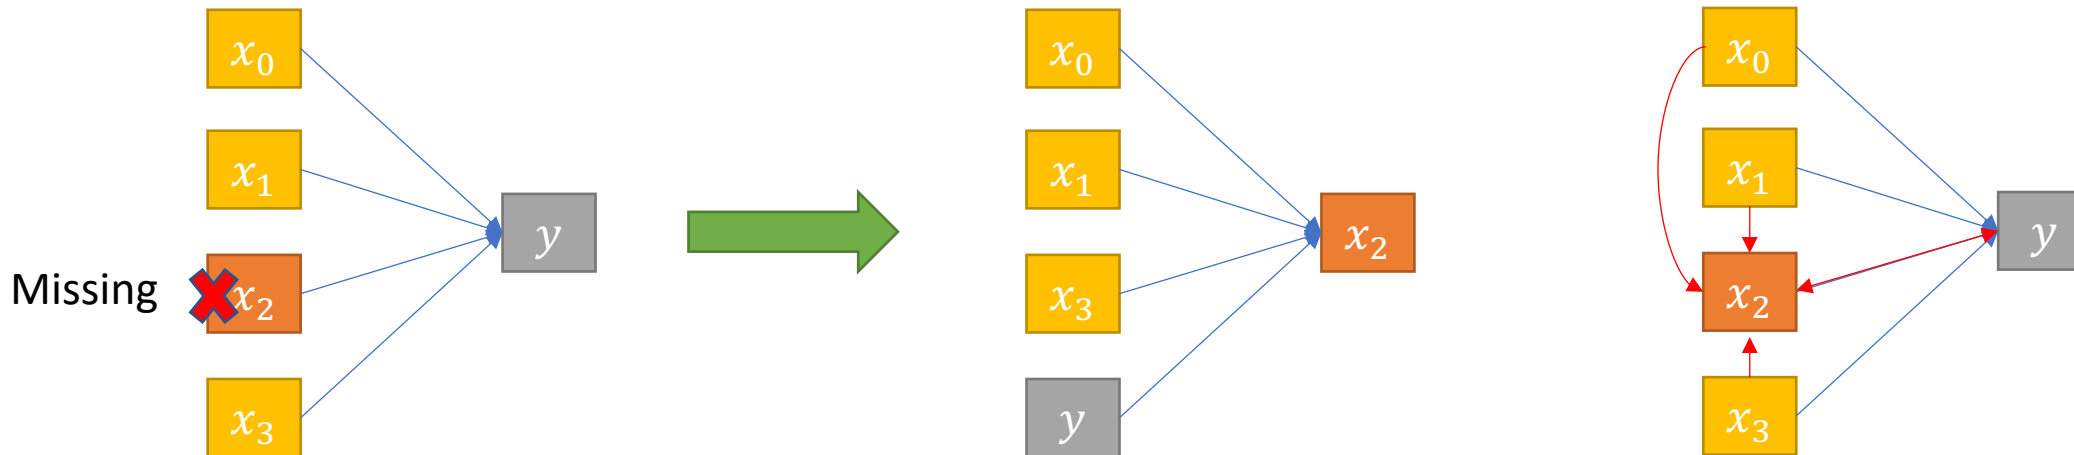
- **Data imputation** (`sklearn.impute`)
 - Normal imputation: most likely value assuming a normal distribution
 - Numerical data: mean / median value
 - `Categorical data: mode` missing feature values
 - Class label imputation: normal imputation using same-class samples
 - Model-based imputation: two-step approach
 1. Train a model to predict the missing value using samples with non-missing values
 2. Predict the missing values

Feature exploration and data cleaning

❑ Missing values

- **Data imputation**

- Model-based imputation (ext.): regression



Feature exploration and data cleaning

❑ Missing values

- **Data imputation**

- Model-based imputation (ext.):

- Multiple imputation by chained equations (MICE)

1. Perform a simple imputation (e.g., mean, median, mode)

2. For each cycle c :

- For each variable v with missing values (ascending order based on number of missing values):

- 2.a. Use all variables except for v to predict the dependent missing values y

- 2.b. Update previously missing values for v using a model trained with all other variables

3. Evaluate convergence criterion (e.g., number of iterations, convergence of variable distributions...)

- K-Nearest Neighbors and other clustering techniques

Feature exploration and data cleaning

❑ Missing values

- **Incorporating missing values in model**

- Adding a present/missing value feature k

$$y = \theta_0 + \theta_1 x_0 + \theta_2 x_1 + k\theta_2 x_2 + \theta_3 x_3$$

- The zero contribution to the gradient **from missing parameters** in a group of samples may bias the optimization algorithm

Note: missing values may not be missing

- Datasets often present values of 0 or default value as n/a or empty.
- Comprehensive visual data exploration is important before **automated analysis**.

Feature exploration and data cleaning

□ Data exploration (with Pandas): data summary

```
print(dataFrame)
```

	Age (years)	Sex	Racial group	Cell profile 0	Cell profile 1	Cell profile 2	...
0	51.060702	Female	Black or African American	-41.684111	-119.691309	-160.581327	...
1	46.496877	Female	Hispanic or Latino	-57.294304	166.842440	-144.137162	...
2	83.191342	Female	Asian	177.426379	114.040486	47.796610	...
3	42.471701	Female	White	NaN	39.509062	-268.640377	...
4	37.555152	Female	Asian	-2.271115	68.110832	-68.743788	...
...
2495	39.671317	Female	American Indian or Alaska Native	-103.507276	-72.794797	90.986559	...
2496	55.985015	Female	White	NaN	231.690782	78.141183	...
2497	39.839834	Male	Hispanic or Latino	98.922336	-50.937941	-90.234747	...
2498	59.924294	Female	Hispanic or Latino	125.041301	14.066663	-247.036367	...
2499	61.561259	Male	NaN	76.367017	-215.978229	-291.415638	...

```
dataFrame.describe()
```

	Age (years)	Cell profile 0	Cell profile 1	Cell profile 2	Cell profile 3	Cell profile 4
count	2500.000000	2475.000000	2500.000000	2500.000000	2500.000000	2500.000000
mean	59.364493	11.592020	12.645618	6.084614	22.849973	-6.115149
std	14.332899	125.409011	202.773280	188.869888	121.565582	145.307406
min	35.005334	-699.398400	-1121.705481	-1036.079878	-655.762851	-818.681653
25%	46.669997	-59.710493	-100.819220	-95.765523	-46.907599	-90.511291
50%	59.104247	11.611980	13.869728	6.137660	19.613346	-5.406596
75%	71.965853	83.803713	126.177779	109.555465	90.510083	73.471722
max	84.956439	685.537250	1157.416566	1049.429225	707.637951	793.509556

Feature exploration and data cleaning

□ Data exploration (with Pandas): missing values

```
→ print(dataFrame.isna())
```

	Age (years)	Sex	Racial group	Cell profile 0	Cell profile 1	Cell profile 2	...
0	False	False	False	False	False	False	...
1	False	False	False	False	False	False	...
2	False	False	False	False	False	False	...
3	False	False	False	True	False	False	...
4	False	False	False	False	False	False	...
...
2495	False	False	False	False	False	False	...
2496	False	False	False	True	False	False	...
2497	False	False	False	False	False	False	...
2498	False	False	False	False	False	False	...
2499	False	False	True	False	False	False	...

```
→ [2500 rows x 24 columns]
print(dataFrame.isna().any())
```

Age (years)	False
Sex	False
Racial group	True
Cell profile 0	True
Cell profile 1	False
Cell profile 2	False
Cell profile 3	False
Cell profile 4	False
Cell profile 5	True
Cell profile 6	True
Cell profile 7	True
Cell profile 8	False
Cell profile 9	False
Cell profile 10	False
Cell profile 11	False
Cell profile 12	False
Cell profile 13	False
Cell profile 14	False
Cell profile 15	False
Cell profile 16	False
Cell profile 17	False
Cell profile 18	False
Cell profile 19	False
Survival time (years)	False

`.isna().any()`

Feature exploration and data cleaning

❑ Data exploration (with Pandas): missing values

```
dataFrame.dropna(inplace=True)
print(dataFrame.describe())
```

	Age (years)	Cell profile 0	Cell profile 1	Cell profile 2	Cell profile 3	Cell profile 4	...
count	2179.000000	2179.000000	2179.000000	2179.000000	2179.000000	2179.000000	...
mean	59.608130	11.848821	11.494482	5.294306	22.612309	-6.371349	...
std	14.337864	123.608102	198.813711	188.152667	120.201394	147.168326	...
min	35.005334	-682.812718	-1121.705481	-1036.079878	-655.762851	-818.681653	...
25%	46.862880	-59.852472	-100.427905	-98.030660	-46.908169	-90.370241	...
50%	59.434286	12.025879	13.868784	8.107831	19.205531	-5.133992	...
75%	72.120810	82.371832	125.082277	106.775255	90.573344	73.937939	...
max	84.956439	657.486647	1157.416566	1049.429225	707.637951	793.509556	...

```
dataFrame['Cell profile 0'].fillna(np.mean(dataFrame['Cell profile 0']), inplace=True)
print(dataFrame.isna().any())
```


Age (years)	False
Sex	False
Racial group	True

Cell profile 0	False
Cell profile 1	False
Cell profile 2	False
Cell profile 3	False
Cell profile 4	False
Cell profile 5	True
Cell profile 6	True
Cell profile 7	True
Cell profile 8	False
Cell profile 9	False
Cell profile 10	False
Cell profile 11	False
Cell profile 12	False
Cell profile 13	False
Cell profile 14	False
Cell profile 15	False
Cell profile 16	False
Cell profile 17	False
Cell profile 18	False
Cell profile 19	False

Feature exploration and data cleaning

□ Data exploration (with Pandas): categorical variables

```
print(dataFrame['Sex'].unique())  
['Female' 'Male']  
print(dataFrame['Racial group'].unique())  
['Black or African American' 'Hispanic or Latino' 'Asian' 'White'  
 'American Indian or Alaska Native'  
 'Native Hawaiian or other Pacific Islander' nan]
```



Feature distributions

Information Visualization (BIOS 7719)

□ Numerical features

- Quantitative values: mean, median, IQR, standard deviation, range... (numpy)

- Visualizations (matplotlib.pyplot as plt):


```
plt.hist(dataFrame['Cell profile 0'], bins=50)  
plt.show()
```

- Testing normality:

- Shapiro-Wilk test (scipy.stats) or other available tests

```
print(scipy.stats.shapiro(dataFrame['Cell profile 0']))  
ShapiroResult(statistic=0.9458088278770447, pvalue=1.3563603337555129e-27)
```


- Often data are only normally distributed for specific subgroups of samples



Feature distributions

□ Categorical features

```
raceLabels = list(set(dataFrame['Racial group']))  
raceCounts = [list(dataFrame['Racial group']).count(c) for c in raceLabels]  
raceLabels[raceLabels==np.nan] = 'nan' # Converting to string  
plt.barh(raceLabels, raceCounts)  
plt.show()
```



Feature distributions


❑ Outliers

- Observations that do not follow the overall patterns in the population
- They can bias the training process and lead towards suboptimal models
- Types:
 - Natural: realistic/plausible observations that are uncommon
 - Error:
 - Data entry (human)
 - Measurement (instrument)
 - Experimental (extraction or execution)
 - Sampling (source of information)


Feature distributions

❑ Outliers

- First step for outlier identification is visualization



```
plt.hist(dataFrame['Cell profile 0'], bins=50)
plt.show()
```



```
plt.boxplot(dataFrame['Cell profile 0'])
plt.show()
```

Feature distributions

□ Identifying outliers

- **z-score**: assumes normal data distributions. A threshold must be established, usually ≥ 3 .

$$z = \frac{x - \mu}{\sigma}$$

- **Thomson's Tau test**:

Rejection interval:

$$\frac{t_{\alpha/2}(n-1)}{\sqrt{n}\sqrt{n-2+t_{\alpha/2}^2}}$$

- **Tukey's "fences"**: $[Q_1 - k * (Q_3 - Q_1), Q_3 + k * (Q_3 - Q_1)]$
 - For $k = 1.5$: outliers. For 3: extreme values.

- Clustering methods: e.g., **KNN, DBScan**

[Martin et al, A density-based algorithm for discovering clusters in large spatial databases with noise. AAAI Press. pp. 226–231, 1996]

Feature distributions


❑ Outliers

- Distinction between natural and error outliers in biomedicine is very important
- Underrepresented cases often appear as natural outliers of the normative population
 - There is a high risk of creating biased methods that discriminate specific populations
- Outlier treatment:
 - Natural outliers: highly encouraged to consider in model training. Consider creating separate models.
 - Error outliers:
 - Remove samples
 - Use data imputation techniques

Feature distributions

❑ Outliers

- There are no robust methods to identify the outlier type and ensure a model free of biases
 - Understanding the dataset is essential



College students?

Basketball team?

Feature interactions

□ Feature interactions

- Multicollinearity: when an “independent” feature can be predicted to a degree from other independent features.
- Why is it a problem?

- Linear regression:
$$\mathbf{X} = \begin{pmatrix} 1 & x_0^{(0)} & \dots & x_{M-1}^{(0)} \\ \vdots & & & \vdots \\ 1 & x_0^{(N-1)} & \dots & x_{M-1}^{(N-1)} \end{pmatrix}$$
 If columns are correlated, the rank of \mathbf{X} is lower than $M + 1$. Hence, $\mathbf{X}\mathbf{X}^T$ does not have an inverse.

- Individual relationships between the dependent and independent variables cannot be recovered.

$$f(x_0, x_1) = \alpha x_0 + \beta x_1 \text{ and } x_0 = \gamma x_1, \text{ then: } f(x_0, x_1) = \alpha x_0 + \beta x_1 = \frac{1}{N} \alpha x_0 + N \beta x_1$$

Feature interactions

□ Detection:

- Coefficients associated with correlated variables usually have high standard errors
- No significant contribution (or extreme contribution) of one variable to the regression model may indicate a collinearity (note that it may also mean lack of correlation or extreme correlation with predicted variable)
- Variance inflation factor for predictor j : $VIF_j = \frac{1}{1-R_j^2}$
 - A value over 5 may indicate collinearity

[O'Brien, R. M. (2007). "A Caution Regarding Rules of Thumb for Variance Inflation Factors". *Quality & Quantity*. **41** (5): 673–690]
- High condition number
 - Two or more features with similarly high variance may suggest which are the correlated predictors.
- Correlation matrix:
 - It can only evaluate pair-wise relationships and multicollinearity often involves several features.

Feature interactions

□ What to do?

- Multicollinearity does not necessarily bias the predictions but their explanation (the contribution from the colinear variables).
- Feature dropping: may cause loss of information (lower predictive accuracy) in exchange for more significant coefficient values.
- Data transformations:
 - Transform the data to a new space where features are uncorrelated.
 - Example: principal component regression:
 - We will see PCA later in the course


Mitigating multicollinearity and outlier values

□ Regularization:


- Can reduce the effect of both outliers and multicollinearity
- L2-regularization (aka Ridge)

$$\frac{\partial}{\partial \theta} J(\theta) = \frac{\partial}{\partial \theta} \frac{1}{2} ((X\theta - y)^T (X\theta - y) + \alpha \theta^T \theta)$$
$$\theta = (X^T X + \alpha I_M)^{-1} X^T y$$

Promotes low parameter values, which tends to prevent extreme effects from independent variables



Without regularization



With L2-regularization

Mitigating multicollinearity and outlier values

□ Regularization:


- L1-regularization (aka Lasso)
 - Its origin comes from the introduction of a **soft** threshold to parameters estimates in linear regression

$$\theta = \arg \min \left\{ \frac{1}{2} (\mathbf{X}\theta - \mathbf{y})^T (\mathbf{X}\theta - \mathbf{y}) \right\} \text{ s.t. } \|\theta\|_1 \leq t$$



$$\theta = \arg \min \left\{ \frac{1}{2} (\mathbf{X}\theta - \mathbf{y})^T (\mathbf{X}\theta - \mathbf{y}) + \alpha \|\theta\|_1 \right\}$$

Lagrangian form (we will see Lagrange multipliers later in the course)



Circles:

- Purple: training
- Green: test

Lines:

- Purple: MSE
- Orange: L1

Mitigating multicollinearity and outlier values

Regularization

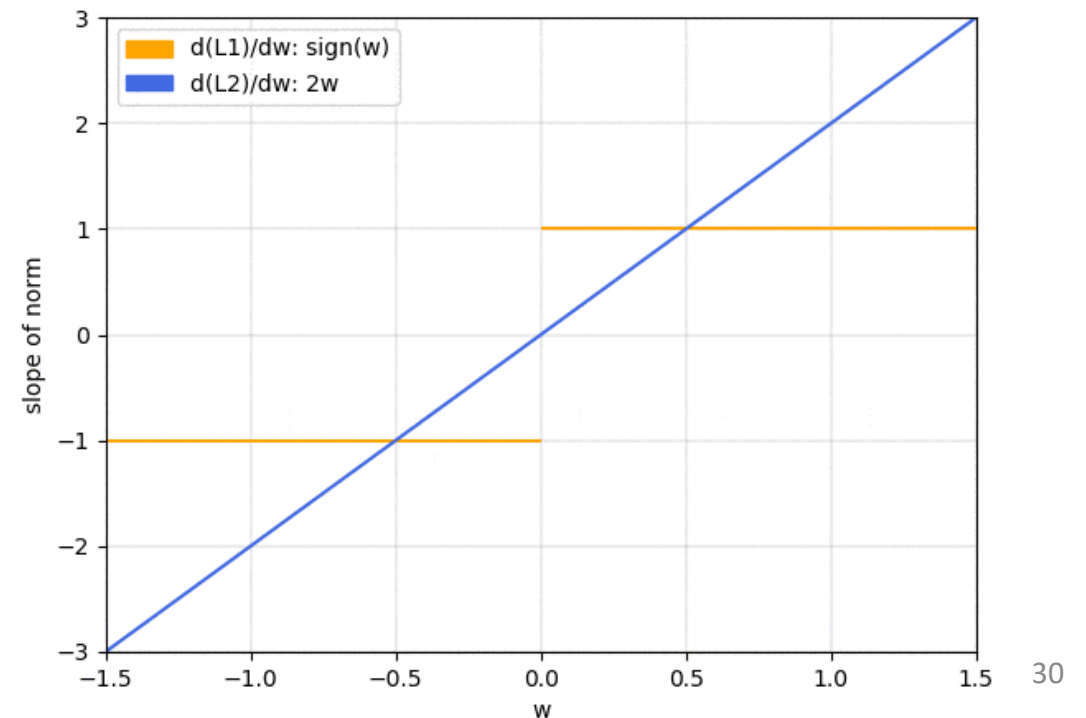
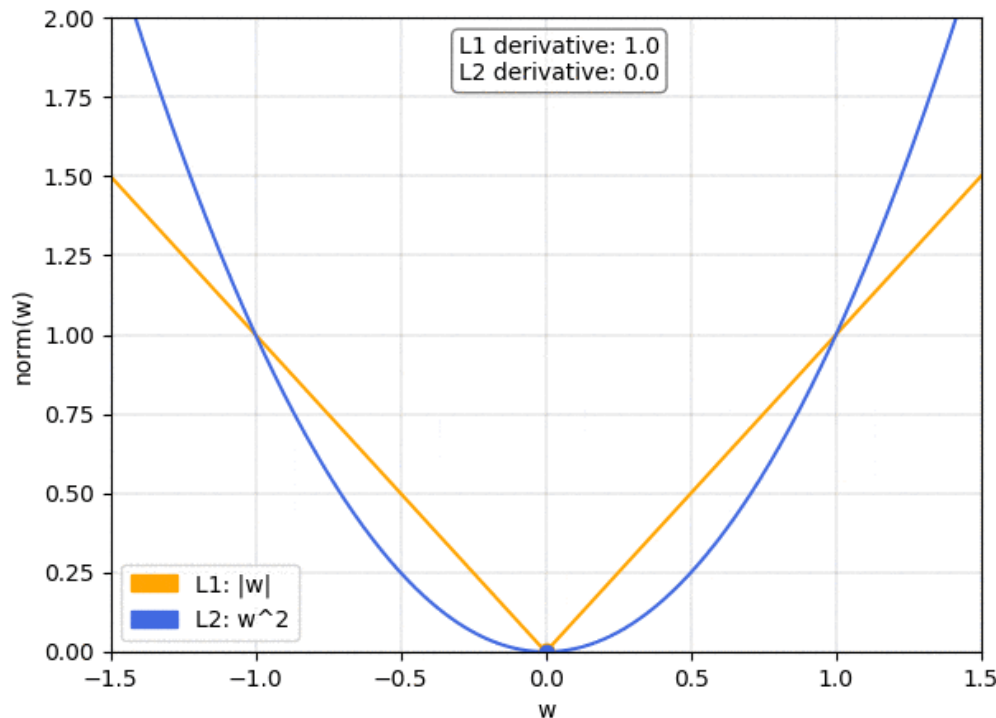
- Gradients: L2 vs. L1 regularization term

L2

$$\frac{1}{2} \alpha \frac{\partial \sum \theta_i^2}{\partial \theta_j} = \alpha \theta_j$$

L1

$$\alpha \frac{\partial \sum |\theta_i|}{\partial \theta_j} = \alpha \text{sign}(\theta_j)$$



Mitigating multicollinearity and outlier values



□ Regularization

- L2: the gradient depends on the value of each coefficient
 - Coefficients with higher values will provide gradients with higher magnitude.
 - Gradient descent optimization will prioritize decreasing the value of coefficients with higher values
- L1: the gradient is constant and the same for all coefficients
 - Coefficients with lower values tend to be zeroed first
 - Zeroing coefficients is equivalent to feature selection... but coefficients with lower magnitude will tend to be eliminated first.

Mitigating multicollinearity and outlier values

Regularization

- Gradients: L2 vs. L1 regularization term



Feature scaling

□ Feature scaling: converting features values to present specific ranges or distributions

□ Why?

- Parameter updates: $\theta_{t+1} = \theta_t + \delta\theta_t = \theta_t - \alpha \frac{\partial}{\partial \theta_t} J(\theta_t)$

- MSE: $\theta_0 \leftarrow \theta_0 - \alpha \sum_{i=0}^{N-1} (f(\mathbf{x}^{(i)}; \theta) - y^{(i)})$

$$\theta_1 \leftarrow \theta_1 - \alpha \sum_{i=0}^{N-1} (f(\mathbf{x}^{(i)}; \theta) - y^{(i)}) \mathbf{x}_0^{(i)}$$

$$\theta_2 \leftarrow \theta_2 - \alpha \sum_{i=0}^{N-1} (f(\mathbf{x}^{(i)}; \theta) - y^{(i)}) \mathbf{x}_1^{(i)}$$

→ Different magnitudes mean different rates of change

- Features with higher magnitude tend to bias optimization algorithms (not only regression models) towards solutions that prioritize the adjustment of model parameters associated with them

Feature scaling

- Feature normalization (or scaling normalization): converts features to specific ranges

$$x_{scaled} = \frac{x - x_{min}}{x_{max} - x_{min}}$$

`sklearn.preprocessing.MinMaxScaler`

- Feature standardization (or z-score normalization): converts features to have zero-mean and unit standard deviation

$$x_{scaled} = \frac{x - \bar{x}}{\sigma_x}$$

`sklearn.preprocessing.StandardScaler`

Feature scaling

❑ Normalization vs. standardization:

- Homogeneous ranges vs. homogeneous variance across features
- Normalized scales vs. normalized distributions
- Unknown distributions vs. assumed Gaussian distributions
- Bounded vs. unbounded
- Highly sensitive vs. robust to outliers

❑ In general, standardization is preferred on methods that assume normal distributions (e.g., Lasso regression), while normalization is preferred when features present uniform distributions and datasets don't present extreme outliers (e.g., KNN)

Takeaway points

- ❑ First, explore your dataset visually and try to understand it as much as possible
- ❑ Proper data encoding and management of missing values and outliers has a significant impact in any machine learning models
- ❑ Exploring automated data summaries (e.g., using Pandas) and plotting feature distributions is the first step to identify if certain assumptions can be made
- ❑ Try to remove feature dependencies when possible. Slightly lower accuracy may be preferred to improve confidence and reproducibility
- ❑ If the dataset may be underpowered or there is evidence of partial multicollinearity and/or possible outliers, consider model regularization
- ❑ Normalize or standardize your data