

# BIOS 7747: Machine Learning for Biomedical Applications

## Dimensionality reduction

Antonio R. Porras ([antonio.porras@cuanschutz.edu](mailto:antonio.porras@cuanschutz.edu))

Department of Biostatistics and Informatics

Colorado School of Public Health

University of Colorado Anschutz Medical Campus

# Outline

- ❑ The curse of dimensionality
- ❑ Dimensionality reduction: principal component analysis
  - Linear geometric transformations
  - Eigenvectors and eigenvalues
  - Multivariate Gaussian transformations
  - Eigen-decomposition of the covariance matrix
- ❑ Supervised dimensionality reduction: linear discriminant analysis

# The curse of dimensionality

## □ Hughes phenomenon or peaking paradox

- Model performance (on test) increases with number of features up to an optimal performance, and decreases after that

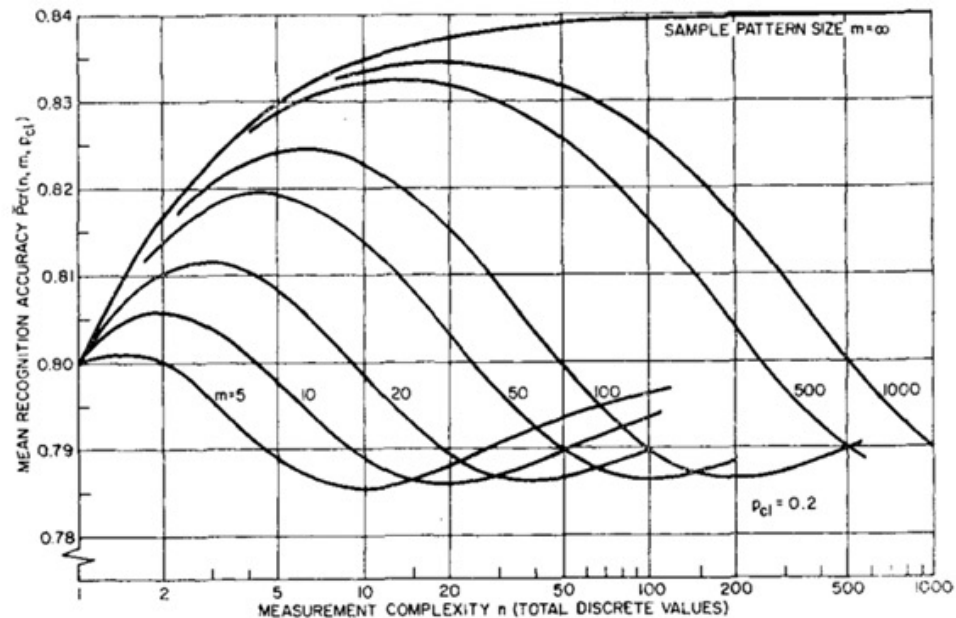
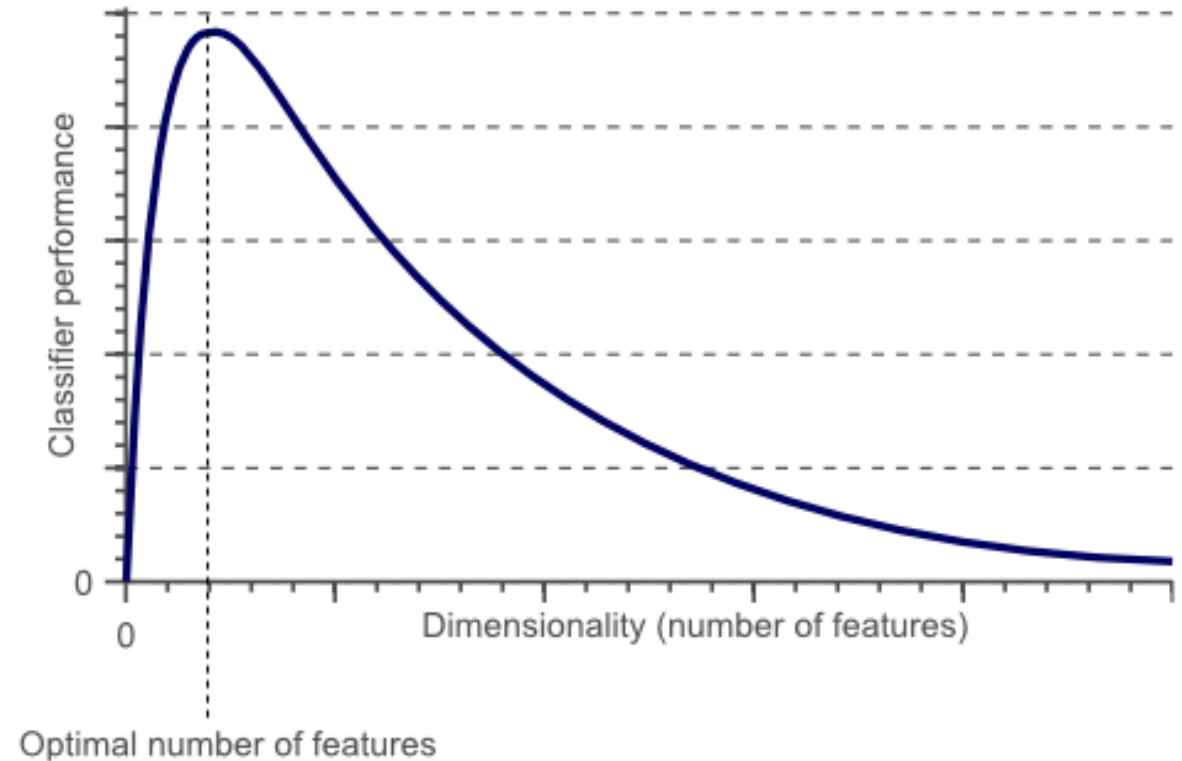


Fig. 4. Finite data set accuracy ( $p_{c1} = \frac{1}{5}$ ).



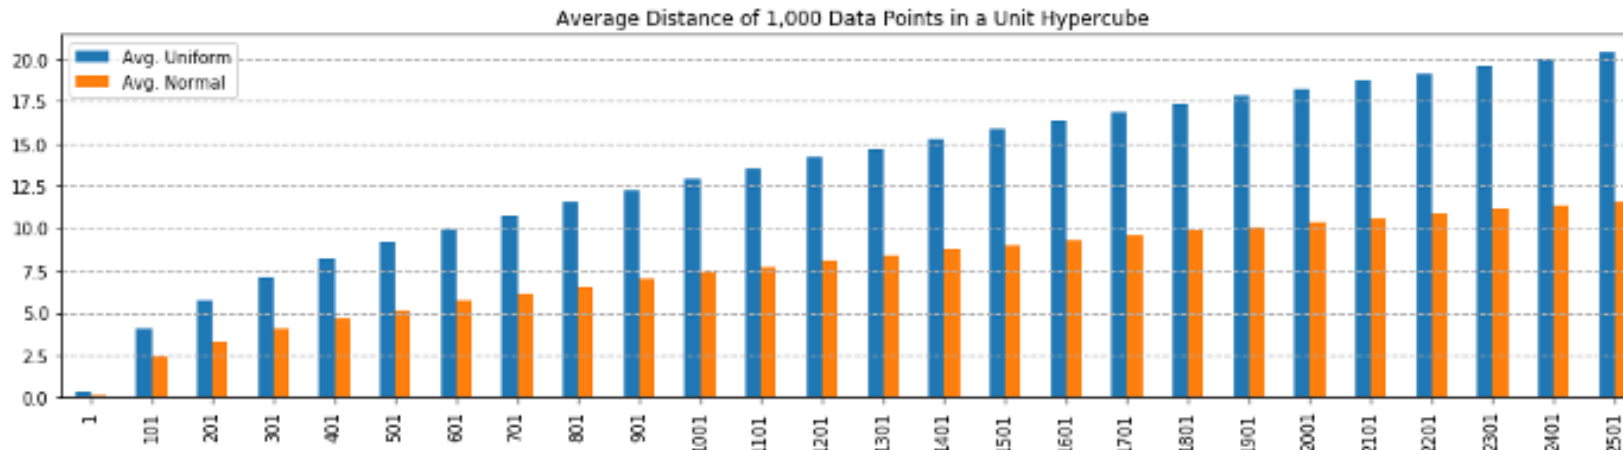
G. Hughes, "On the mean accuracy of statistical pattern recognizers," in *IEEE Transactions on Information Theory*, vol. 14, no. 1, pp. 55-63, January 1968, doi: 10.1109/TIT.1968.1054102.

# The curse of dimensionality

- Increasing the number of features increases data sparsity

Average Euclidean distance between two samples for M features

$$d_{p,q} = \sqrt{\sum_{i=1}^M (p_i - q_i)^2}$$



- As number of features increase, larger sample sizes are needed to preserve average distances (to fill in the empty space)

# The curse of dimensionality

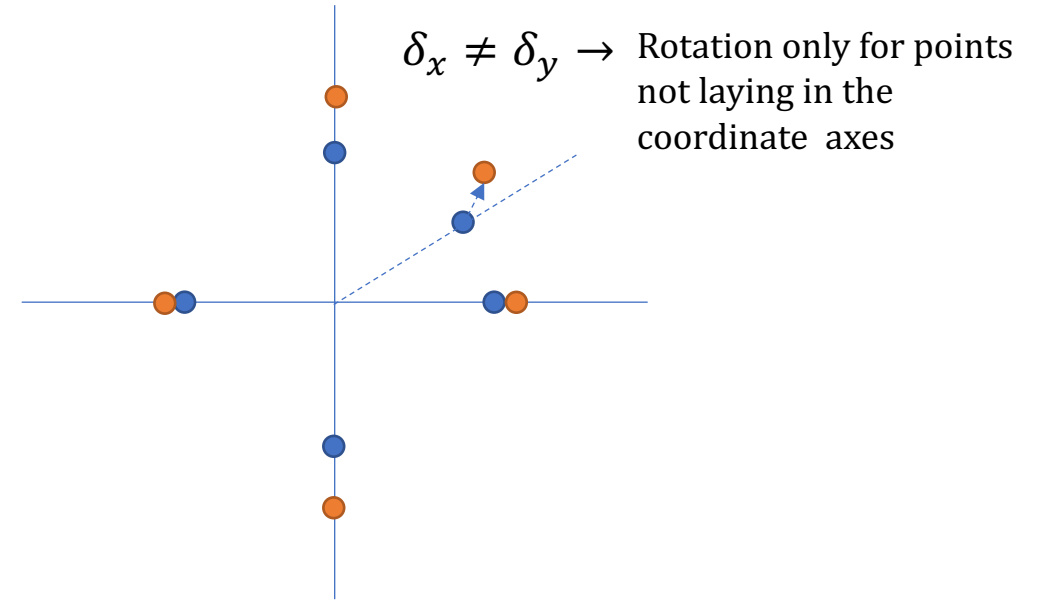
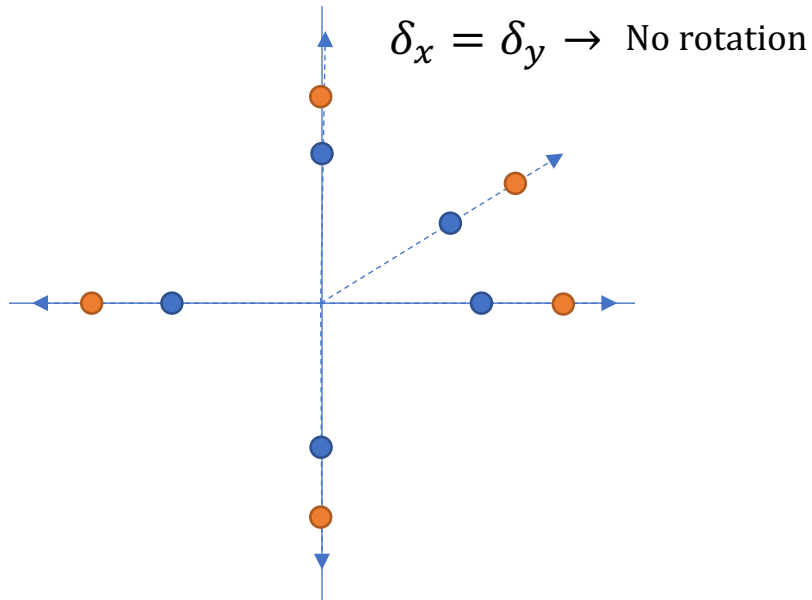
- ❑ High number of features usually mean higher likelihood of creating an overfit model
- ❑ Dimensionality reduction:
  - Feature selection: identify most relevant features to a machine learning task
  - Feature extraction: create a lower-dimensional representations of the available features that is meaningful to a machine learning task

# Linear geometric transformations

## □ Linear geometric transformations

$$T(x) = Ax$$

$$A = \begin{pmatrix} 1 + \delta_x & 0 \\ 0 & 1 + \delta_y \end{pmatrix}$$



For any scaling-only matrix  $\Lambda = \begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix}$ , points in the axes are scaled but not rotated

- Scale on axis 1:  $\begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix} \begin{pmatrix} 1 \\ 0 \end{pmatrix} = \lambda_1 \begin{pmatrix} 1 \\ 0 \end{pmatrix} = \lambda_1 e_1 = \Lambda e_1$
- Scale on axis 2:  $\begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix} \begin{pmatrix} 0 \\ 1 \end{pmatrix} = \lambda_2 \begin{pmatrix} 0 \\ 1 \end{pmatrix} = \lambda_2 e_2 = \Lambda e_2$

Any diagonal matrix maps any vector parallel to a basis vector into another vector that is also parallel to that basis vector

# Eigenvectors and eigenvalues

□ For any diagonal matrix  $\mathbf{\Lambda} = \begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix}$ , points in the axes are scaled but not rotated

- Scale on axis 1:  $\begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix} \begin{pmatrix} 1 \\ 0 \end{pmatrix} = \lambda_1 \begin{pmatrix} 1 \\ 0 \end{pmatrix} = \lambda_1 \mathbf{e}_1 = \mathbf{\Lambda} \mathbf{e}_1$

- Scale on axis 2:  $\begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix} \begin{pmatrix} 0 \\ 1 \end{pmatrix} = \lambda_2 \begin{pmatrix} 0 \\ 1 \end{pmatrix} = \lambda_2 \mathbf{e}_2 = \mathbf{\Lambda} \mathbf{e}_2$



Eigenvalue equation

$$\mathbf{\Lambda} \mathbf{e}_i = \lambda_i \mathbf{e}_i$$

eigenvector

eigenvalue

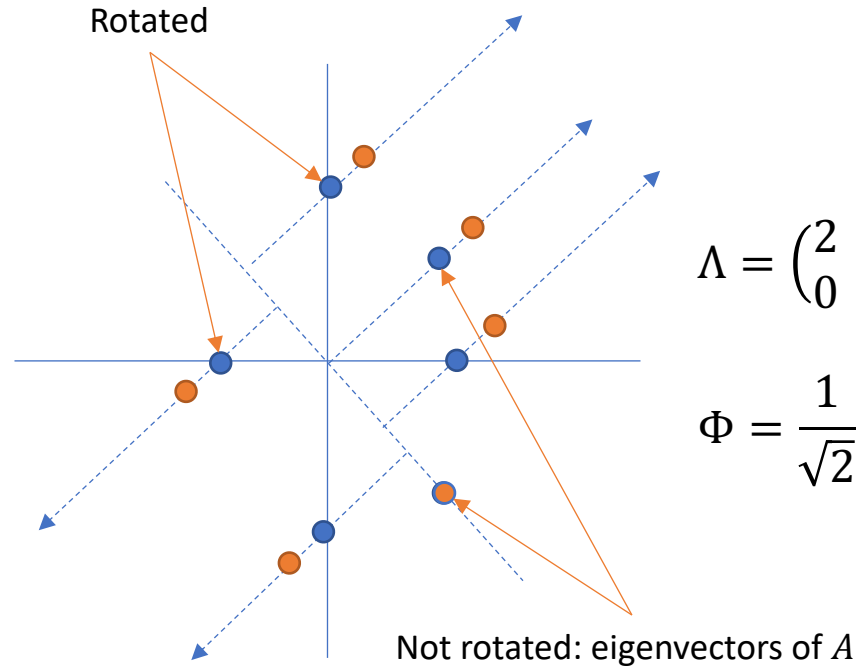
□ Eigenvector of a matrix: any vector  $\mathbf{v}$  that is mapped by the matrix into a parallel vector  $\lambda \mathbf{v}$

□ Eigenvalue of a matrix: the scale factor of an eigenvector

□ A matrix in  $\mathbb{R}^n$  has  $n$  pairs of eigenvectors and eigenvalues

# Eigenvectors and eigenvalues

Rotated scale transformations:  $T(x) = Ax$        $A = \Phi\Lambda\Phi^T$



$$\Lambda = \begin{pmatrix} 2 & 0 \\ 0 & 1 \end{pmatrix}$$

$$\Phi = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & -1 \\ 1 & 1 \end{pmatrix} \text{ [45 degrees]}$$

$$Ae_i = \lambda_i e_i \quad \leftarrow \text{eigenvector}$$

$$\Phi\Lambda\Phi^T e_i = \lambda_i e_i$$

$$\Phi^T \Phi\Lambda\Phi^T e_i = \Phi^T \lambda_i e_i$$

$$\Lambda\Phi^T e_i = \Phi^T (\lambda_i e_i)$$

$$\Lambda\Phi^T e_i = \lambda_i \Phi^T e_i$$

$$\text{Basis vector: } b_i = \Phi^T e_i \\ e_i = \Phi b_i$$

Eigenvectors are rotated standard basis vectors

$$\begin{aligned} \bullet \quad \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & -1 \\ 1 & 1 \end{pmatrix} \begin{pmatrix} 1 \\ 0 \end{pmatrix} &= \frac{1}{\sqrt{2}} \begin{pmatrix} -1 \\ 1 \end{pmatrix} = e_1 \\ \bullet \quad \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & -1 \\ 1 & 1 \end{pmatrix} \begin{pmatrix} 0 \\ 1 \end{pmatrix} &= \frac{1}{\sqrt{2}} \begin{pmatrix} 1 \\ 1 \end{pmatrix} = e_2 \end{aligned}$$

Each column is  
an eigenvector!

Any symmetric transformation matrix A can be written as:

$$A = \Phi\Lambda\Phi^T = (e_1, e_2 \dots) \begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix} (e_1, e_2 \dots)^T$$



# Eigenvectors and eigenvalues

## □ Calculating eigenvalues and eigenvectors: 2D example

$$\mathbf{A}\mathbf{e}_i = \lambda_i \mathbf{e}_i$$

Solutions for which  $\mathbf{e}_i \neq \mathbf{0}$ :  $\det(\mathbf{A}) = \begin{vmatrix} a - \lambda_i & b \\ c & d - \lambda_i \end{vmatrix} = (a - \lambda_i)(d - \lambda_i) - cd = 0$

$$\mathbf{A}\mathbf{e}_i = \lambda_i \mathbf{I}\mathbf{e}_i$$

$$\det(\mathbf{A} - \lambda_i \mathbf{I}) = 0$$

$$(\mathbf{A} - \lambda_i \mathbf{I})\mathbf{e}_i = 0$$

$$\lambda_i^2 - \lambda_i(a + d) + (ad - bc) = 0$$

Characteristic equation of matrix A

$$\lambda_i = \frac{1}{2}(a + d) \pm \frac{1}{2}\sqrt{(a - d)^2 + 4bc}$$

→ Eigenvalues can be real or complex!

In a symmetric matrix, since  $b = c$ , then:  $\lambda_i = \frac{1}{2}(a + d) \pm \frac{1}{2}\sqrt{(a - d)^2 + 4b^2}$ . All eigenvalues are real.

## □ Eigen-decomposition: process of decomposing $\mathbf{A}$ in $\mathbf{A} = \mathbf{\Phi}\mathbf{\Lambda}\mathbf{\Phi}^T$

$$\mathbf{A}\mathbf{\Phi} = \mathbf{\Phi}\mathbf{\Lambda}\mathbf{\Phi}^T\mathbf{\Phi}$$

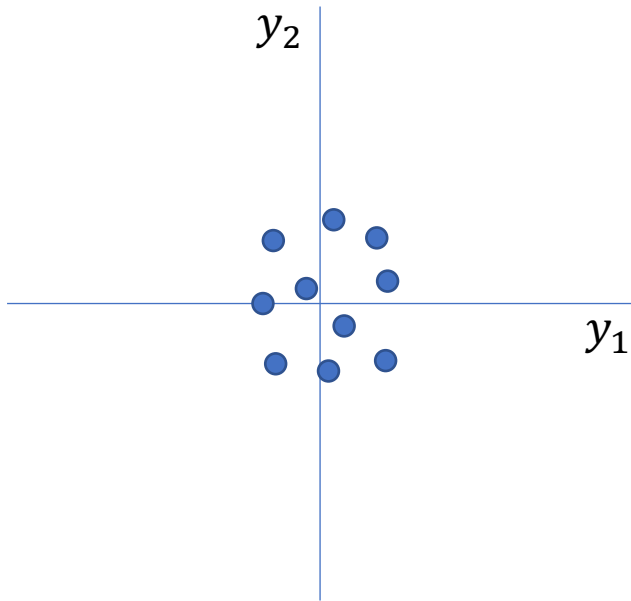
$$\mathbf{A}\mathbf{\Phi} = \mathbf{\Phi}\mathbf{\Lambda}$$

↕ Equivalent

$$\mathbf{A}\mathbf{e}_i = \lambda_i \mathbf{e}_i$$

# Multivariate Gaussian distributions

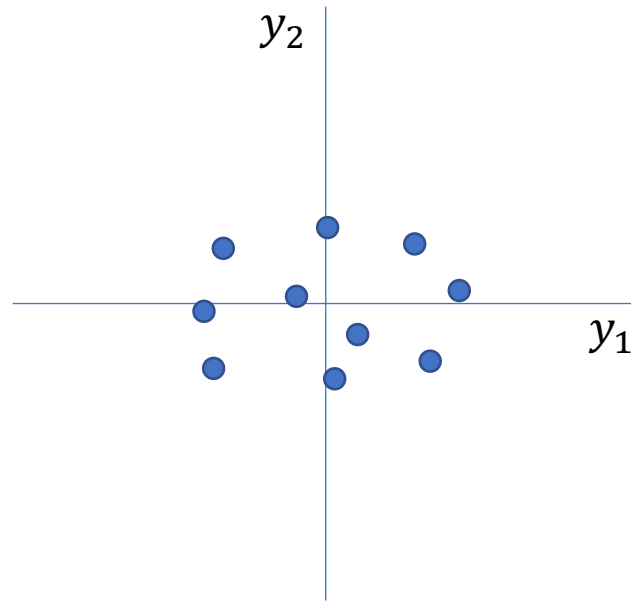
Isotropic



Variance:  $\sigma^2$

$$p(y) \propto e^{-\frac{1}{2\sigma^2}y^T y}$$

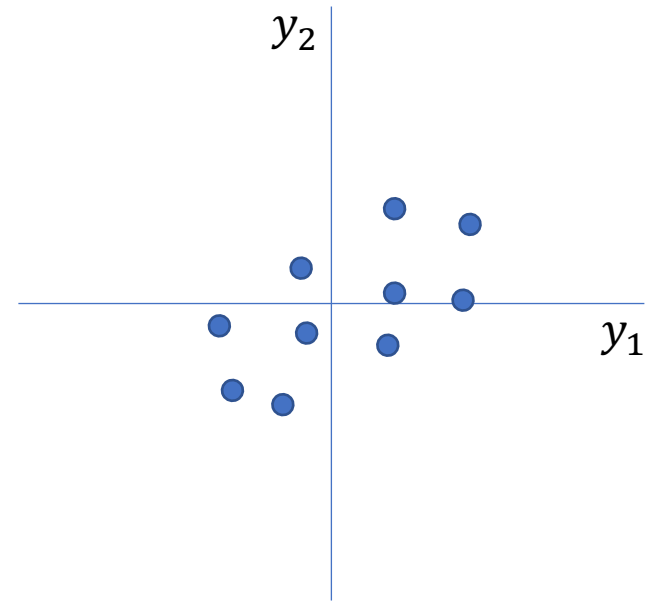
Anisotropic (no correlation)



Variance matrix:  $s = \begin{pmatrix} \sigma_1^2 & 0 \\ 0 & \sigma_2^2 \end{pmatrix}$

$$p(y) \propto e^{-\frac{1}{2}y^T s^{-1} y}$$

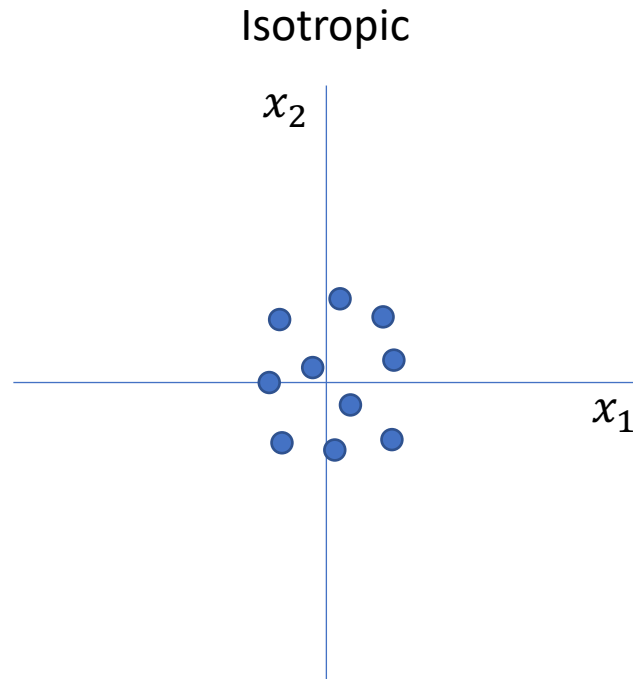
Anisotropic (with correlation)



Co-variance matrix:  $\Sigma = \begin{pmatrix} \sigma_1^2 & \sigma_{1,2} \\ \sigma_{2,1} & \sigma_2^2 \end{pmatrix}$

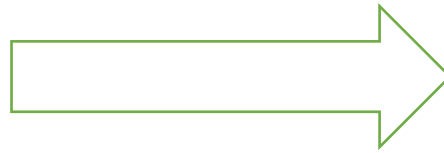
$$p(y) \propto e^{-\frac{1}{2}y^T \Sigma^{-1} y}$$

# Transforming multivariate Gaussian distributions



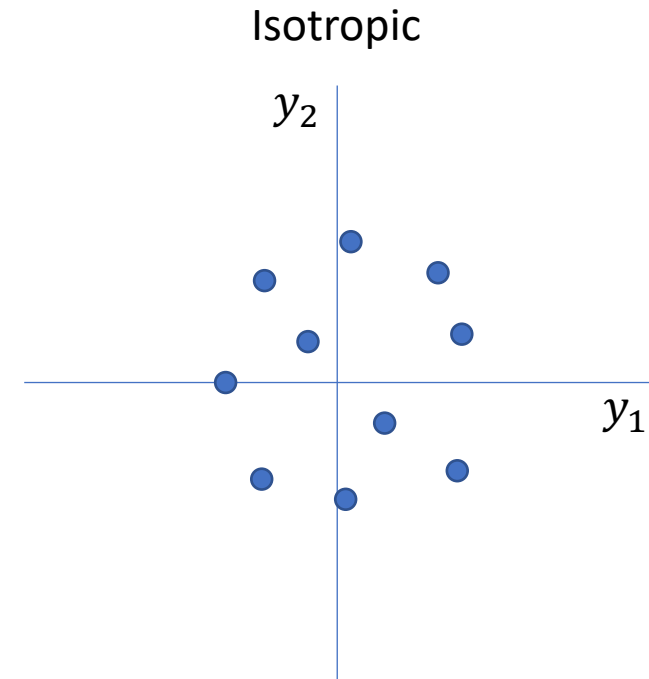
Variance:  $\sigma_x^2 = 1$

$$p(y) \propto e^{-\frac{1}{2}x^T x}$$



$$y = \sigma_y x$$

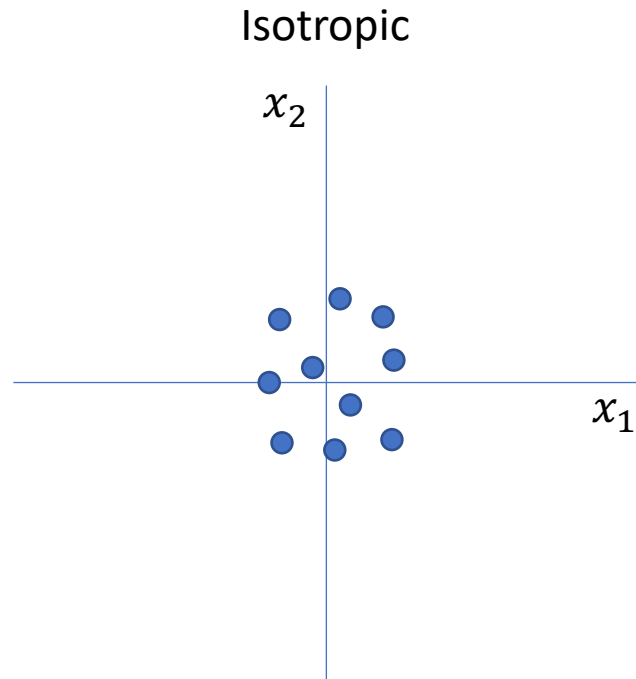
$$x = \sigma_y^{-1} y$$



Variance:  $\sigma_y^2$

$$p(y) \propto e^{-\frac{1}{2\sigma_y^2}y^T y}$$

# Transforming multivariate Gaussian distributions



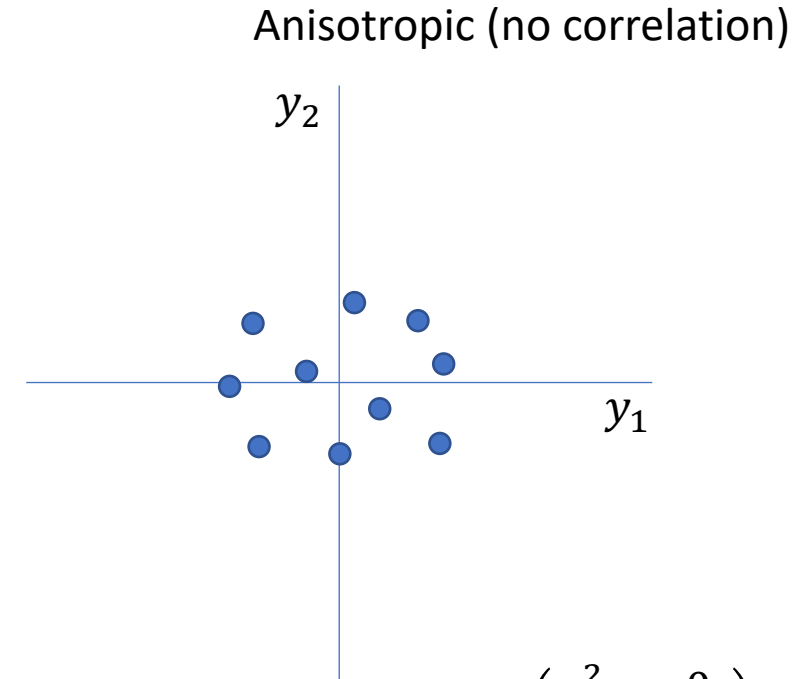
Variance:  $\sigma_x^2 = 1$

$$p(y) \propto e^{-\frac{1}{2}x^T x}$$



$$y = S_y x$$

$$x = S_y^{-1} y$$



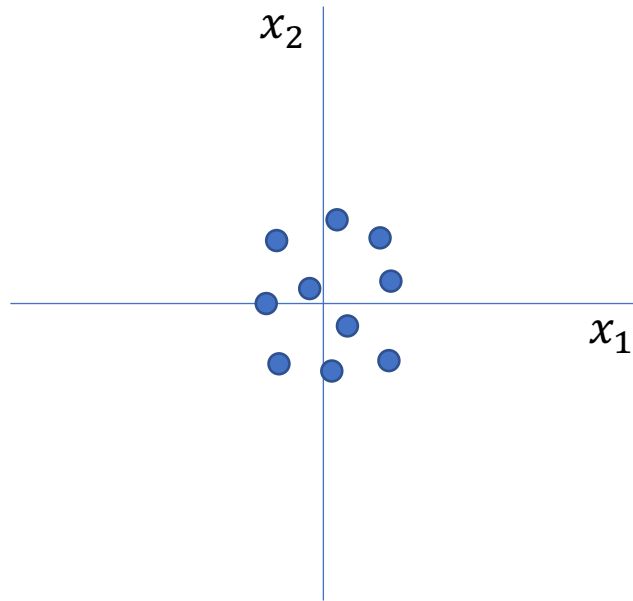
Variance matrix:  $S_y^2 = \Lambda = \begin{pmatrix} \sigma_{y_1}^2 & 0 \\ 0 & \sigma_{y_2}^2 \end{pmatrix}$

$$p(y) \propto e^{-\frac{1}{2}y^T S_y^{-2} y} = e^{-\frac{1}{2}y^T \Lambda^{-1} y}$$

$$\Lambda^{-1} = \begin{pmatrix} \sigma_{y_1}^{-2} & 0 \\ 0 & \sigma_{y_2}^{-2} \end{pmatrix}$$

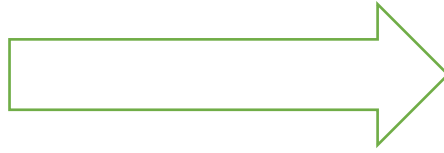
# Transforming multivariate Gaussian distributions

Isotropic



Variance:  $\sigma_x^2 = 1$

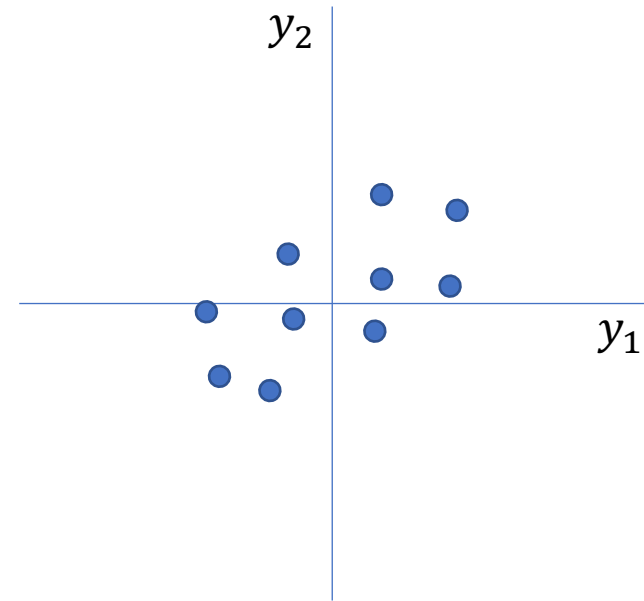
$$p(y) \propto e^{-\frac{1}{2}x^T x}$$



$$y = \Phi S_y \Phi^T x$$

$$x = \Phi S_y^{-1} \Phi^T y$$

Anisotropic (with correlation)



$$\text{Covariance matrix: } \Sigma = \begin{pmatrix} \sigma_{y_1}^2 & \sigma_{y_1 y_2} \\ \sigma_{y_2 y_1} & \sigma_{y_2}^2 \end{pmatrix} = \Phi S_y^2 \Phi^T = \Phi \Lambda \Phi^T$$

$$p(y) \propto e^{-\frac{1}{2}y^T \Sigma^{-1} y}$$

$$\Sigma^{-1} = \Phi \Lambda^{-1} \Phi^T$$

# Eigen-decomposition of the covariance matrix

- A covariance matrix can be seen as a linear mapping between an isotropic Gaussian distribution and a non-isotropic Gaussian distribution with correlated features.
  - Note that all distributions are centered at the origin of coordinates (i.e., feature means are zero)
- Eigen-decomposition of the covariance matrix provides:
  - Eigenvectors: The basis vectors of the rotated space
  - Eigenvalues: The variance of the data is the direction of the basis vectors

# Principal component analysis

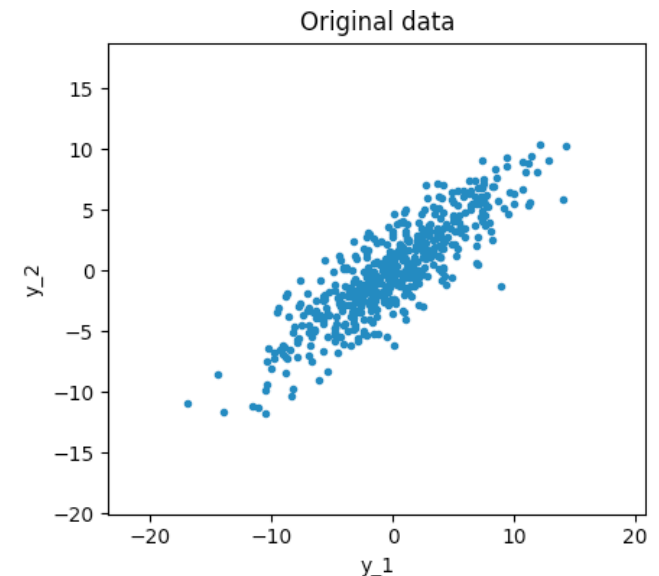
- PCA: data analysis in a rotated feature space (so the features are not correlated) calculated using the eigenvectors of the covariance matrix

**Example: Data analysis in the rotated space using the eigenvectors of the covariance matrix**

```
# Creating dataset with feature correlations
y1 = np.random.normal(0, 5, 500)
noise = np.random.normal(0, 2, 500)
y2 = 0.7 * y1 + noise
data = np.array([y1, y2]).T

covariance = np.cov(data.T)
print(f'Covariance of original data: {covariance}')
```

Covariance of original data:  
[[28.07275458 20.33218298]  
 [20.33218298 18.98284443]]



# Principal component analysis

**Example: Data analysis in the rotated space using the eigenvectors of the covariance matrix**

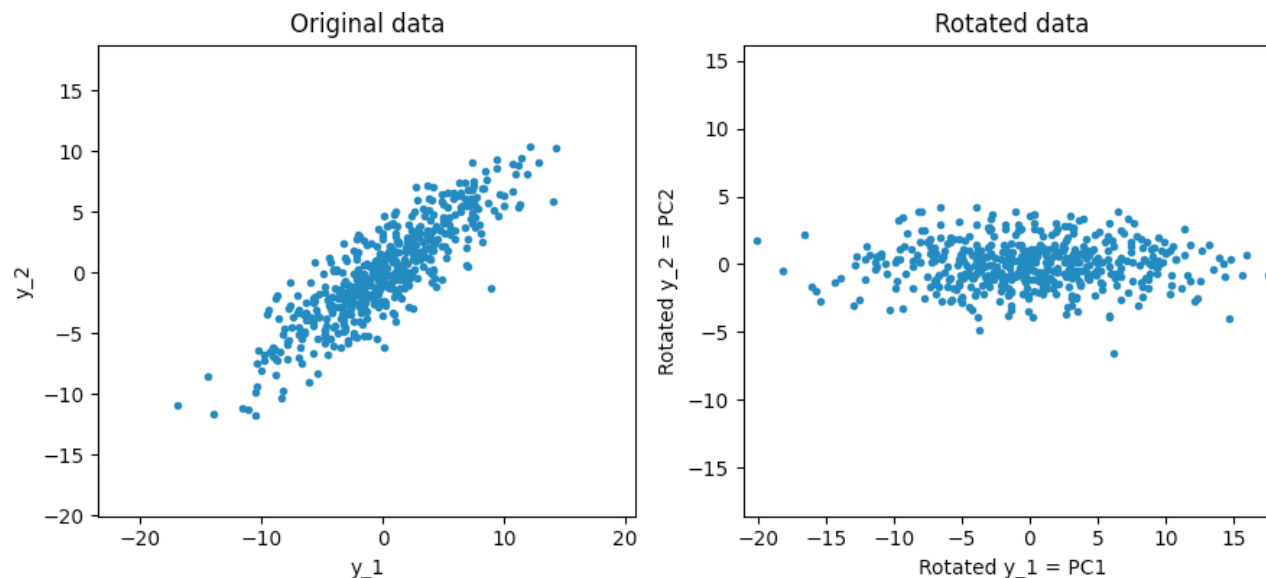
```
# Calculating eigenvalues and eigenvectors
eigenValues, eigenVectors = np.linalg.eig(covariance)
eigenVectors = eigenVectors.T
print(f'Eigenvalues: {eigenValues}')
print(f'Eigenvectors: {eigenVectors}')

# Rotating the data
yRotated = data @ eigenVectors.T
yRotatedCovariance = np.cov(yRotated.T)
print(f'Covariance of y rotated: {yRotatedCovariance}')
```

Eigenvalues:  
[44.36176891 2.6938301 ]

Eigenvectors:  
[[ 0.78043295 0.62523948]  
 [-0.62523948 0.78043295]]

Covariance of y rotated:  
[[4.43617689e+01 9.45580732e-16]  
 [9.45580732e-16 2.69383010e+00]]





# Principal component analysis

Data analysis in the rotated space using the eigenvectors of the covariance matrix

```
# Creating dataset with non-linear fea
x = np.random.normal(0, 5, 500)
noise = np.random.normal(0, 2, 500)
y = 0.7 * x**2 - x + noise
y -= np.mean(y)
data = np.array([x, y]).T
```

Covariance of original data:

```
[[ 24.39693751 -22.91257623]
 [-22.91257623 673.74145237]]
```

Eigenvalues:

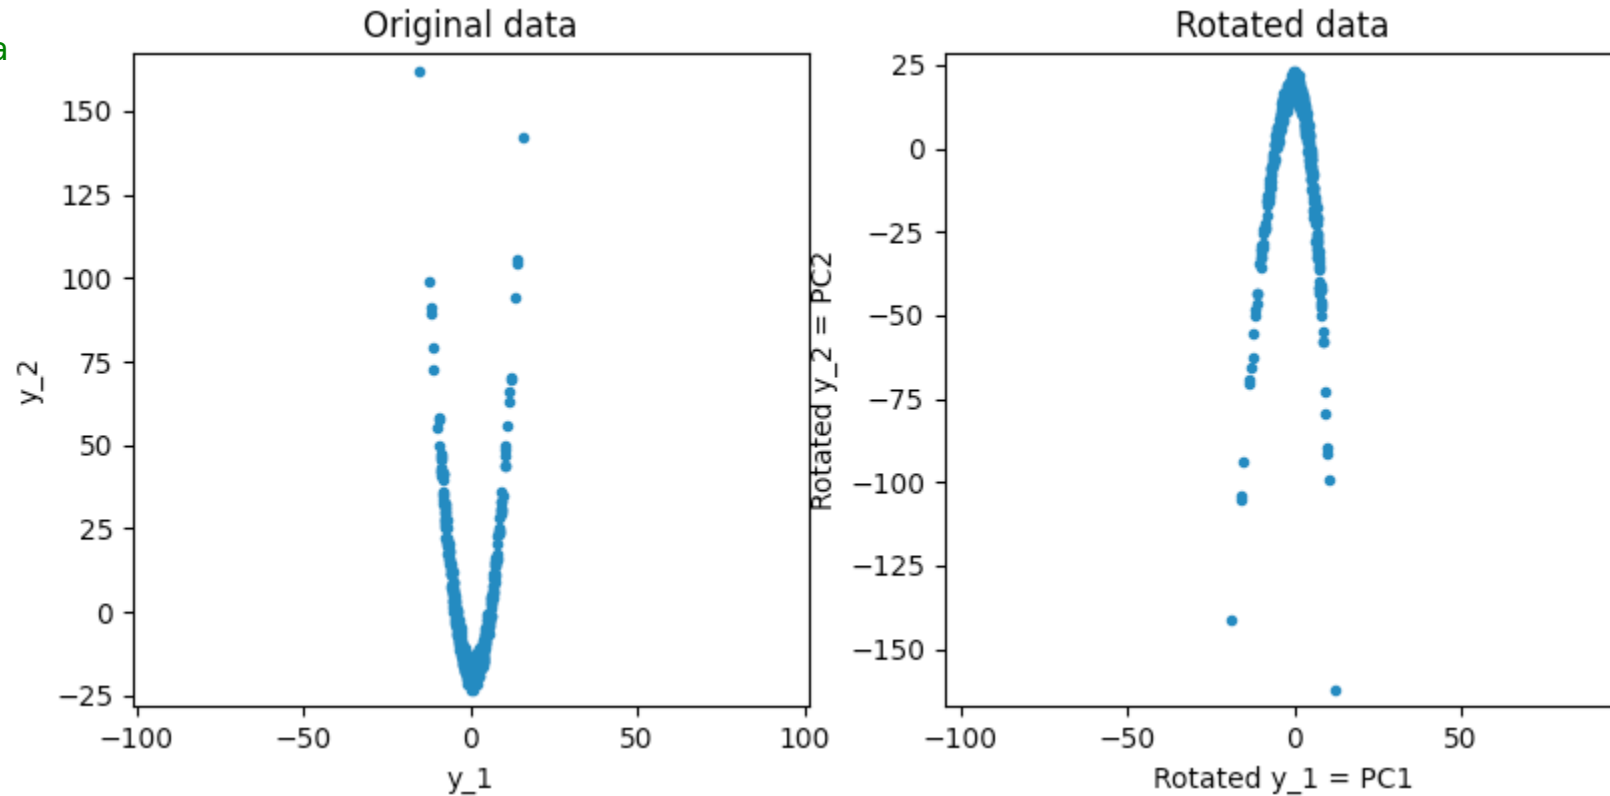
```
[ 23.58945533 674.54893455]
```

Eigenvectors:

```
[[-0.99937958 -0.03522001]
 [ 0.03522001 -0.99937958]]
```

Covariance of y rotated:

```
[[ 2.35894553e+01 -1.27584427e-14]
 [-1.27584427e-14 6.74548935e+02]]
```



PCA is just a rotation...

# Principal component analysis

... but a very powerful rotation!

Example of noise elimination via dimensionality reduction

```
nSamples = 400
nFeatures = 200

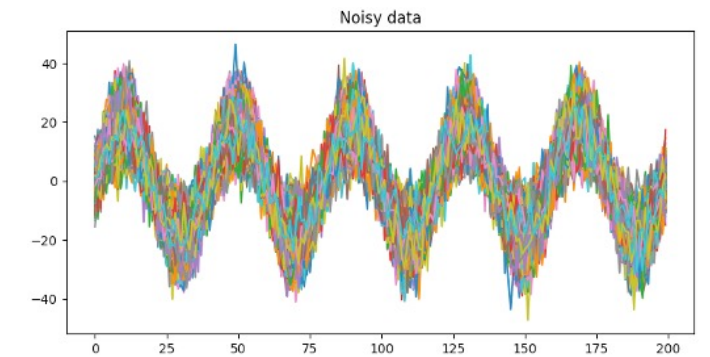
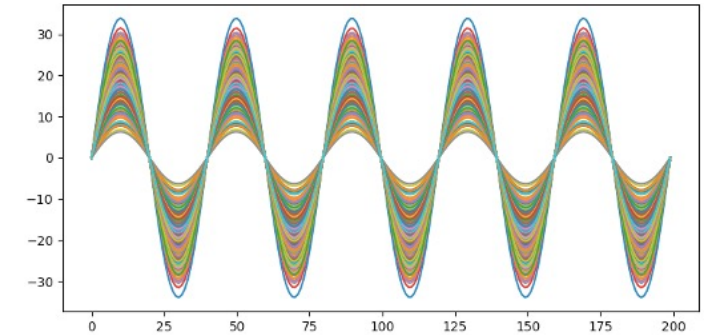
time = np.linspace(0, 1, nFeatures).reshape(1,-1)
temporalSamples = np.repeat((2 * np.pi * time * 5), nSamples, axis=0)
magnitude = np.random.normal(20, 5, nSamples).reshape(-1, 1)
data = magnitude * np.sin(temporalSamples)

noisyData = data + np.random.normal(0, 5, (nSamples, nFeatures))

meanData = np.mean(noisyData, axis=0, keepdims=True)
covariance = np.cov((noisyData - meanData).T)

eigenValues, eigenVectors = np.linalg.eig(covariance)
eigenVectors = eigenVectors.T

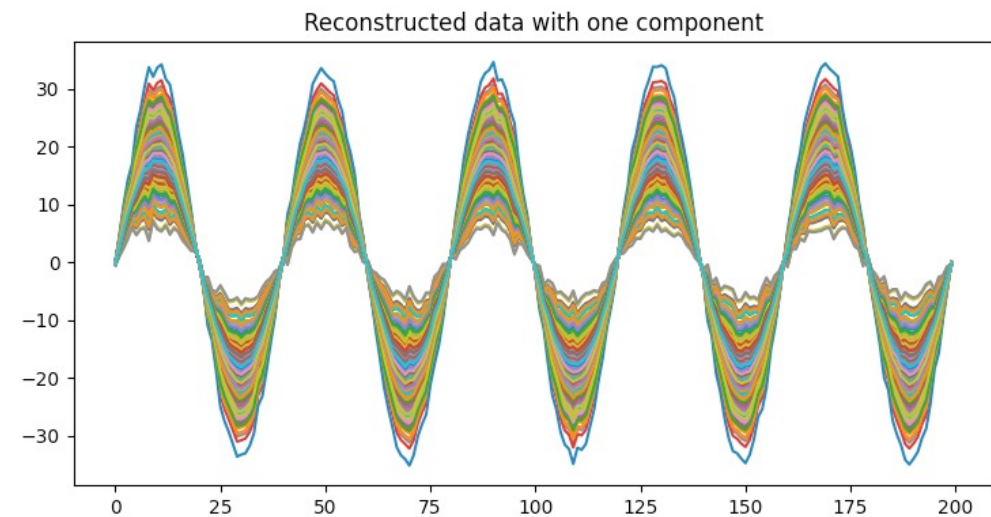
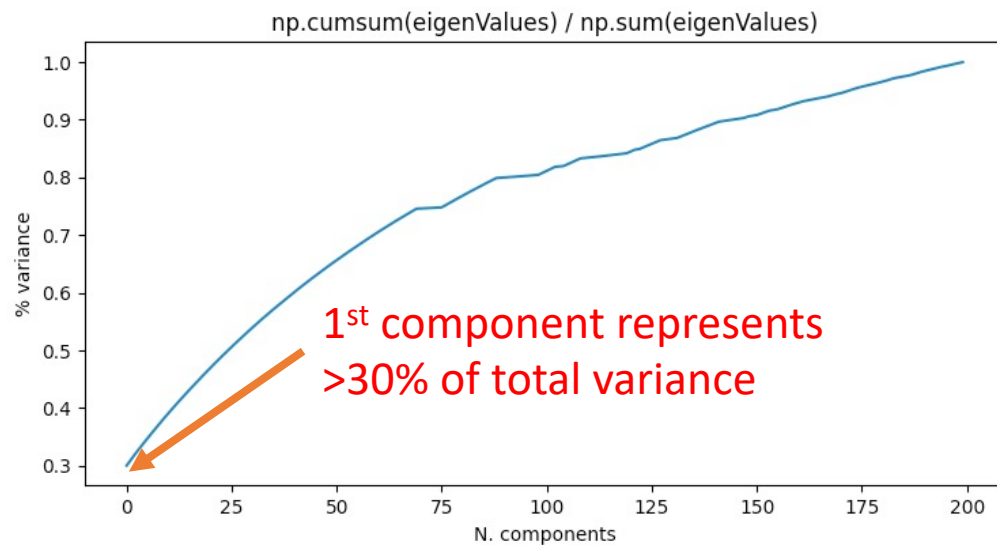
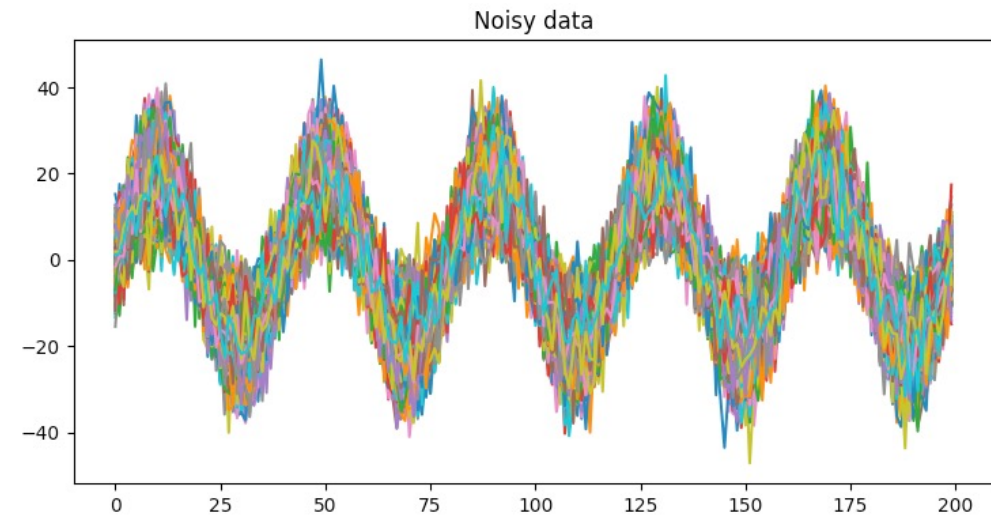
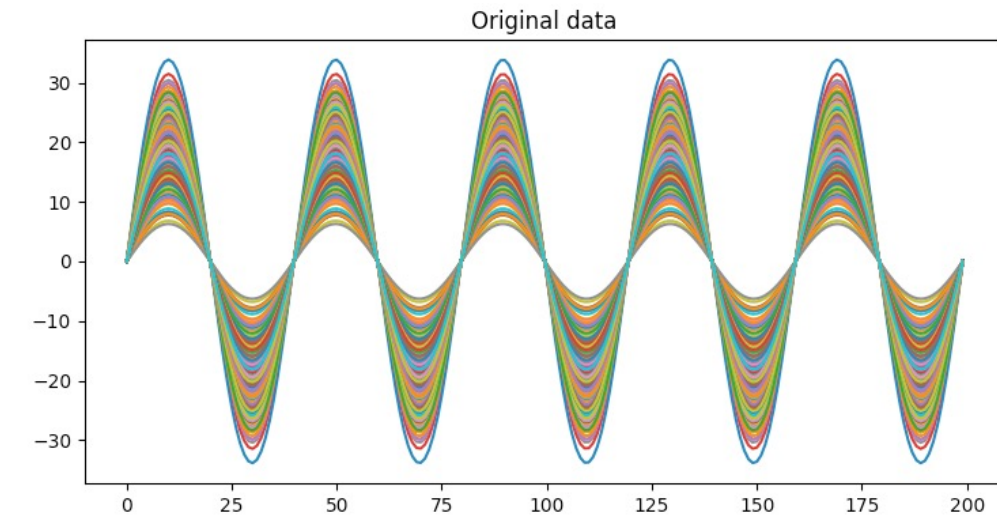
rotatedData = (noisyData - meanData) @ eigenVectors.T
reconstructedData = rotatedData[:, :1] @ eigenVectors[:, 1, :] + meanData
```



Remember that the data needs to be centered at the origin of coordinates

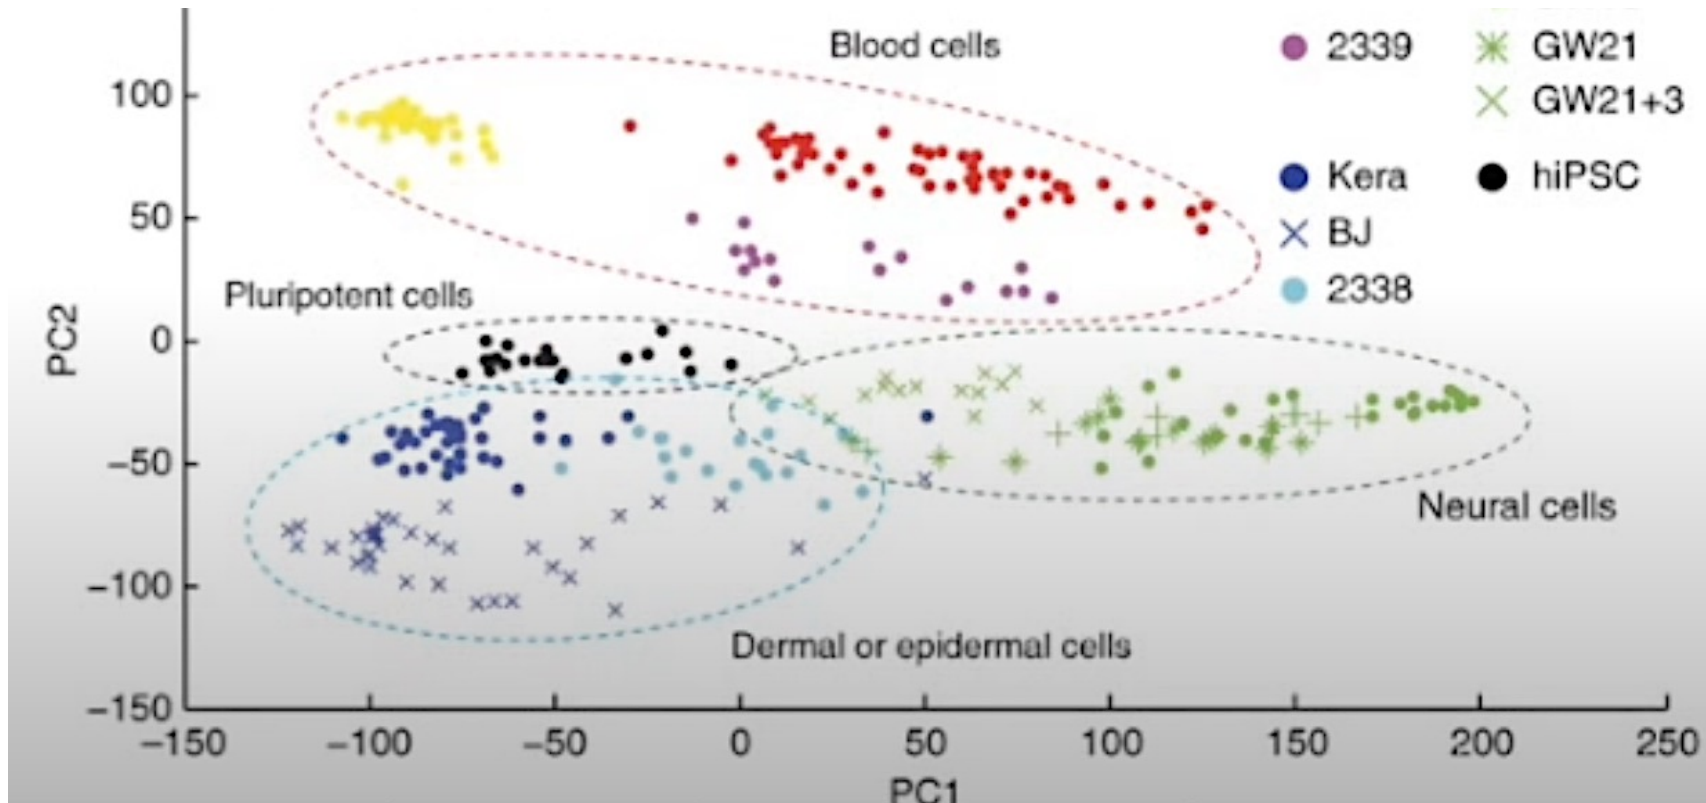
# Principal component analysis

## Example of noise elimination via dimensionality reduction



# Principal component analysis

**Dimensionality reduction is RNA sequencing data**  
(every cell is represented by ~10,000 transcribed genes)

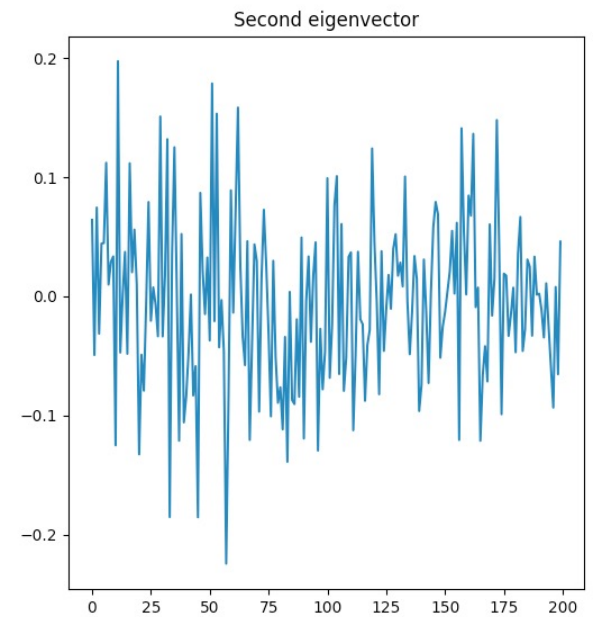
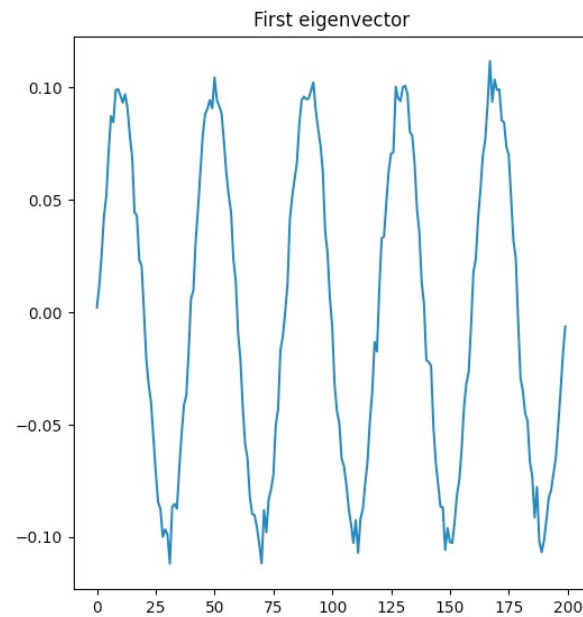
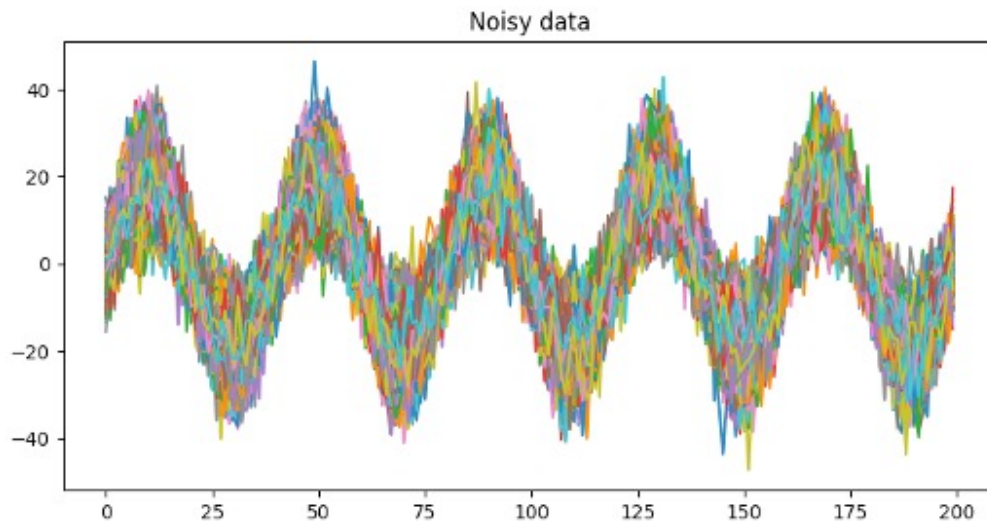


Note: not observing a specific structure in the first principal components does not mean that there is no structure. After all, **PCA is only a linear rotation**

[Pollen et al., Nature Biotechnology, 2014]

# Principal component analysis

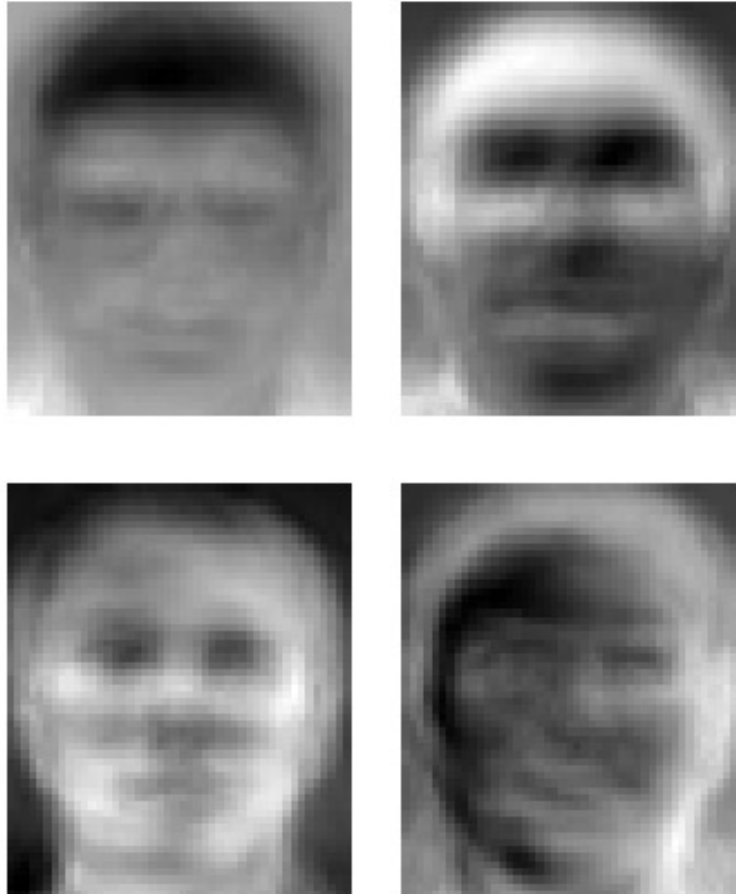
The eigenvectors can be interpreted visually



# Principal component analysis

The eigenvectors can be interpreted visually

Eigenfaces



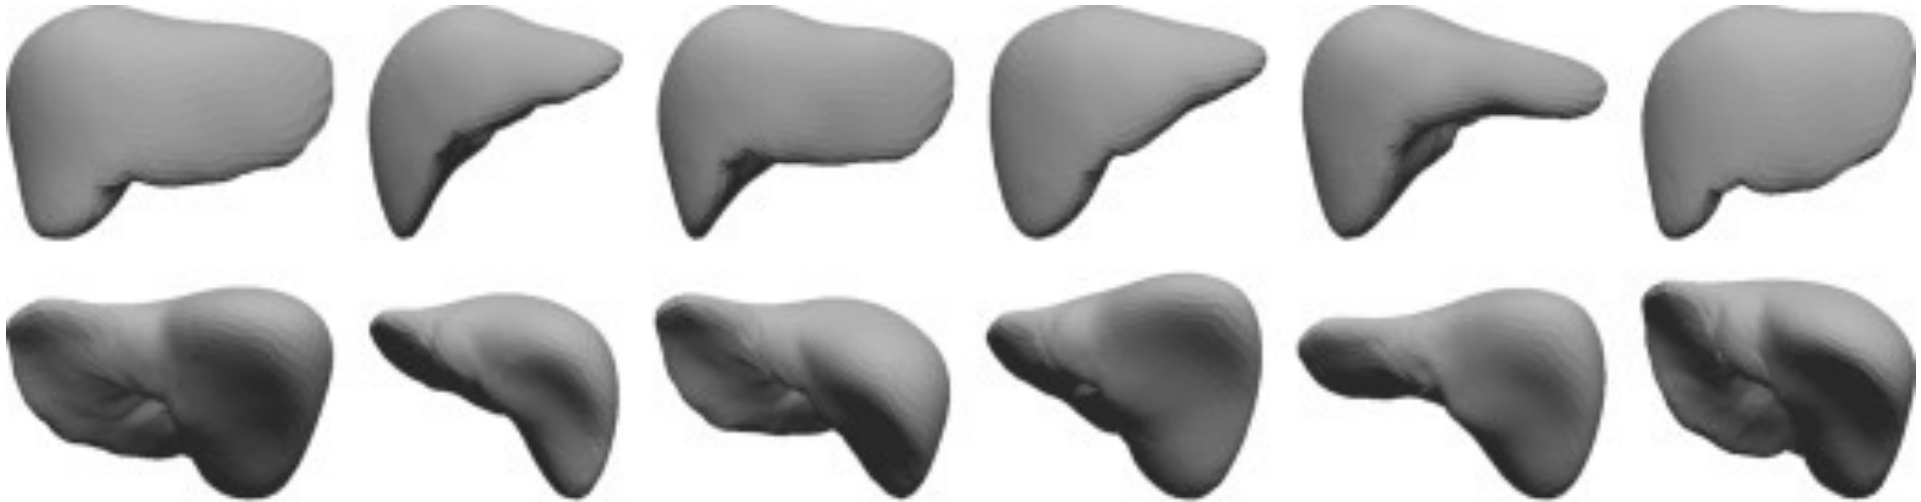
•M. Kirby; L. Sirovich (1990). "Application of the Karhunen-Loeve procedure for the characterization of human faces". *IEEE Transactions on Pattern Analysis and Machine Intelligence*. **12** (1): 103–108. [doi:10.1109/34.41390](https://doi.org/10.1109/34.41390).



# Principal component analysis

- Since PCA features are orthogonal and normally distributed, and the eigenvalues represent the variance, PCA can be used to generate synthetic data

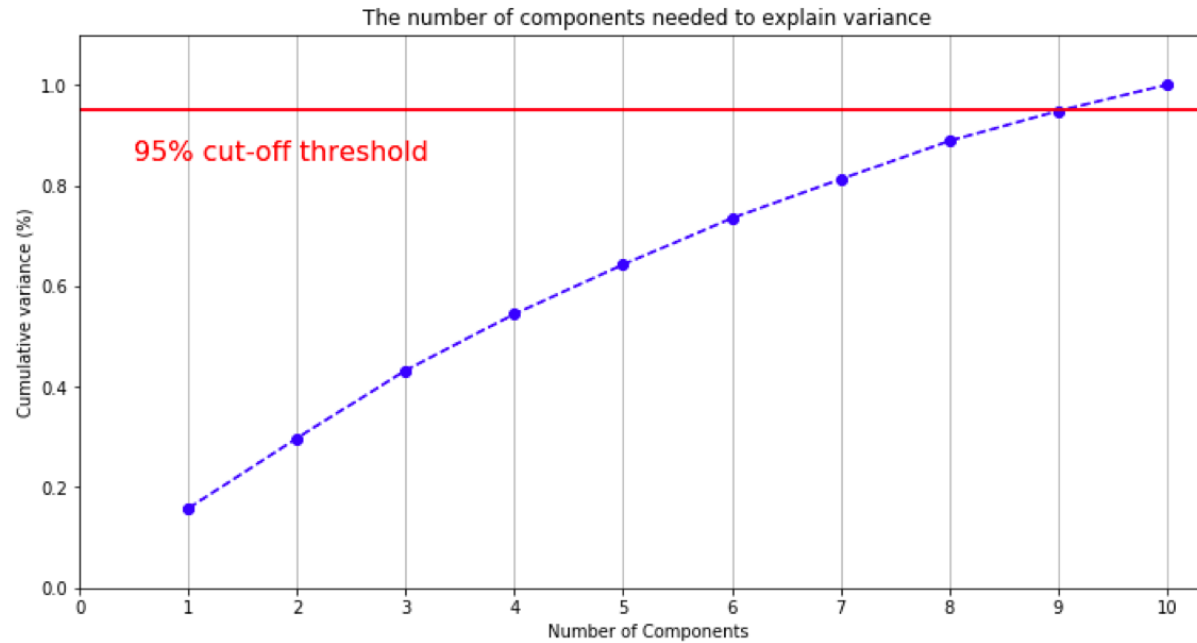
$$y = \bar{x} + p\Phi^T, p \in N(0, \Lambda)$$



# Principal component analysis

## ❑ Selecting number of components:

- Thresholds of between 80% and 98% are common but there is no golden rule.



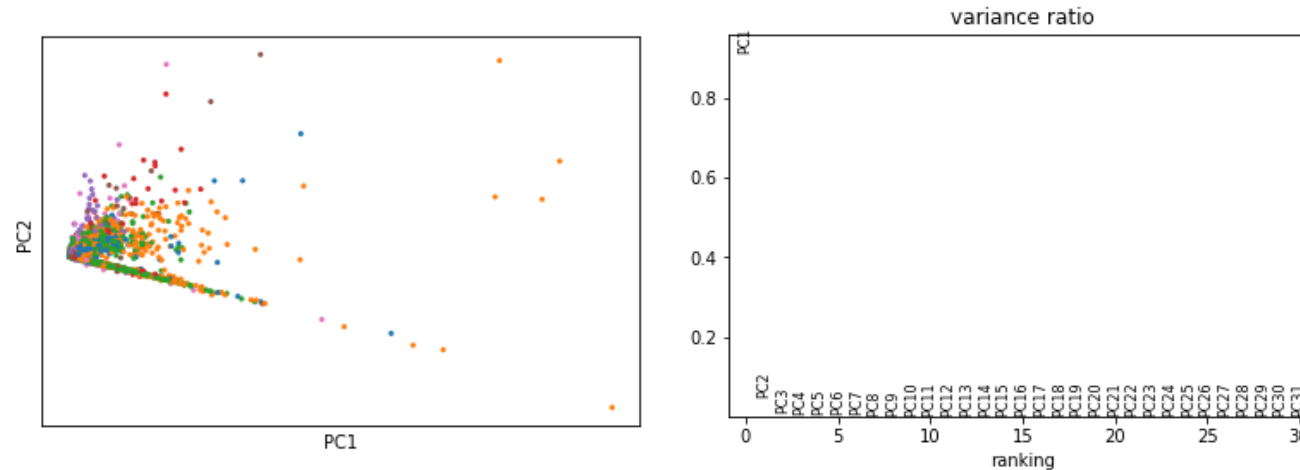
- ## ❑ Component selection is also important for applications such as modeling: The last few components represent data noise and tend to distort the generated data.



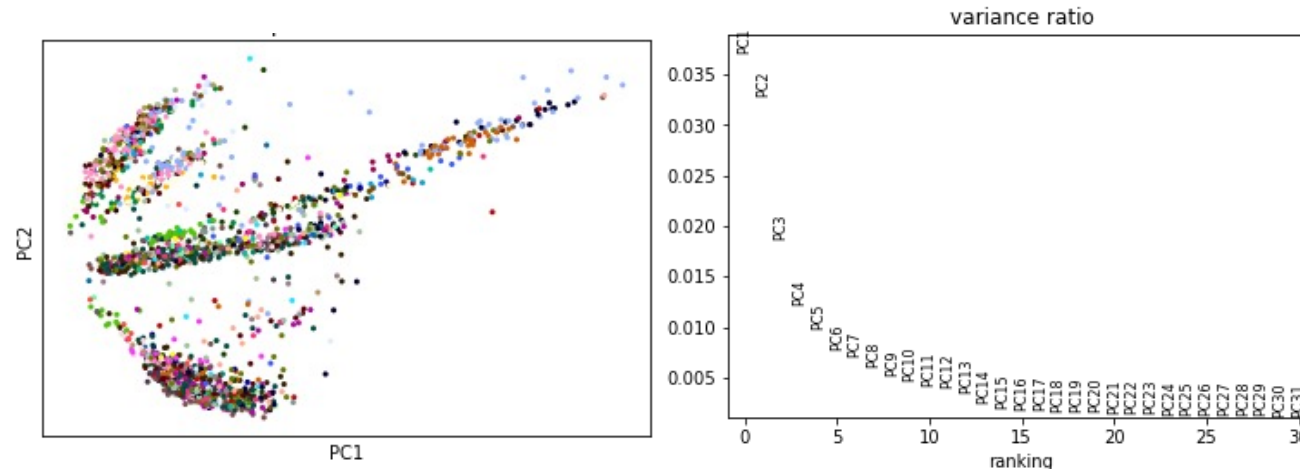
# Principal component analysis

- ❑ Feature normalization is essential for adequate modeling.
  - Example: single cell RNA sequencing data to identify tissue types in mice

Before normalization



After normalization



<https://tabula-muris-senis.ds.czbiohub.org>

# Principal component analysis

## ❑ Limitations

- Not useful for categorical variables
- Dimensionality reduction comes at the expense of accuracy
- Variance is often not equivalent to feature importance
- Assumes normal data distributions
  - Check: [Generalized PCA](#)
- Feature distributions are often not orthogonal
  - Check: [Independent component analysis](#)
- Feature correlations are often non-linear
  - Check: [Kernel PCA](#) (remember the kernel trick?)

# Linear discriminant analysis (LDA)

## □ Supervised dimensionality reduction

- The outcome of the machine learning task is known in the training dataset and it can be leveraged to extract a lower dimensional representation
- Instead of finding a lower-dimensional representation that maximizes data variance, we seek the maximization of class separability

## □ Linear discriminant analysis (LDA)

...or Fisher's linear discriminant

- Goal: find the **linear** combination of features that maximizes the separability between different classes
- Approach: maximization of between-class variance and minimization of within-class variance

# Linear discriminant analysis (LDA)

## □ Linear Discriminant Analysis

- Linear combination of features:  $y = w^T x$
- Objective:  $\max \left( \frac{S_b}{S_{w_1} + S_{w_2}} \right)$ 
  - Between-class variance (points to  $S_b$ )
  - Within-class variance (points to  $S_{w_1} + S_{w_2}$ )

Between class variance:

Original space:  $S_b = (\mu_1 - \mu_2)(\mu_1 - \mu_2)^T$

Transformed space:  $S_b^* = (w^T \mu_1 - w^T \mu_2)(w^T \mu_1 - w^T \mu_2)^T = w^T (\mu_1 - \mu_2)(\mu_1 - \mu_2)^T w = w^T S_b w$

Within class variance:

Original space:  $S_{w_i} = \sum_{x_j \in l_i} (x_j - \mu_j)(x_j - \mu_j)^T$

Transformed space:  $S_{w_i}^* = \sum_{x_j \in l_i} (w^T x_j - w^T \mu_j)(w^T x_j - w^T \mu_j)^T = \sum_{x_j \in l_i} w^T (x_j - \mu_j)(x_j - \mu_j)^T w = \sum_{x_j \in l_i} w^T S_{w_i} w$

# Linear discriminant analysis (LDA)

- Fisher's criterion:

$$J(w) = \frac{w^T S_b w}{w^T S_w w}$$

$$w = \arg \max(J(w))$$

- Solving for  $w$

$$\frac{dJ(w)}{dw} = S_w^{-1} S_b w - J(w) w = 0$$

$$S_w^{-1} S_b w = J(w) w$$

Generalized eigenvalue problem

$$w = S_w^{-1} (\mu_1 - \mu_2)$$

Fisher's discriminant: not really a discriminant but a direction for data projections

# Linear discriminant analysis (LDA)

- LDA generalizes to separating  $C$  classes using  $C-1$  projections

$$w \rightarrow W = [w_1 | w_2 \dots | w_{C-1}]$$

$$y = W^T x$$

$$S_w = \sum_{i=1}^{C-1} S_{w_i}$$

$$S_b = \sum_{i=1}^{C-1} (\mu_1 - \mu)(\mu_i - \mu)^T$$

$$J(w) = \frac{W^T S_b W}{W^T S_w W} \rightarrow S_w^{-1} S_b w_i = \lambda_i w$$

$w_i$  are the eigenvectors with the highest eigenvalues

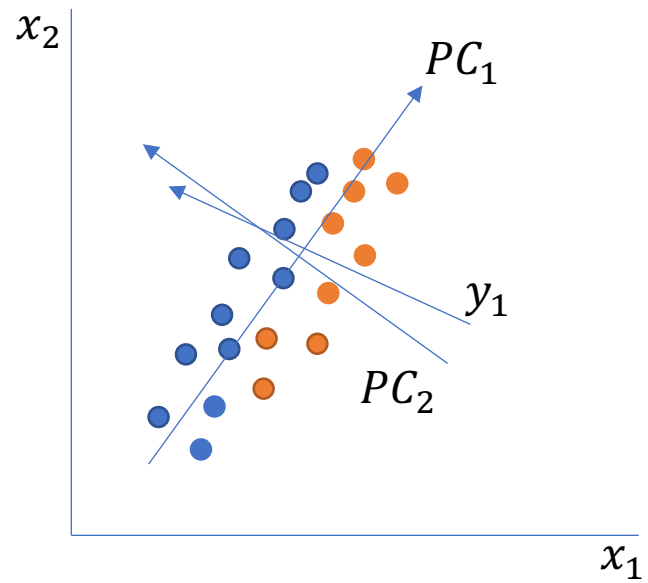
# Linear discriminant analysis (LDA)

## □ Limitations

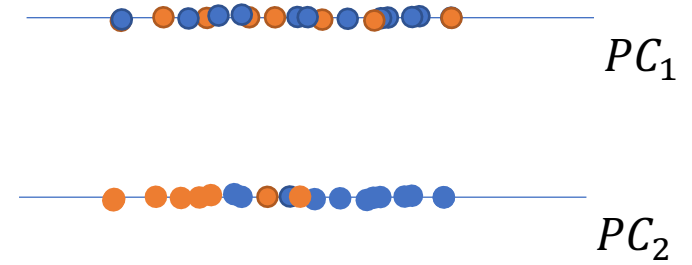
- Assumes linearly separable classes
  - Remember the kernel trick? Check: [Generalized or kernel discriminant analysis](#)
- Homoscedasticity: uniform variances
- Assumes normal distributions
- The maximum number of projections is limited by the number of classes

# LDA vs PCA

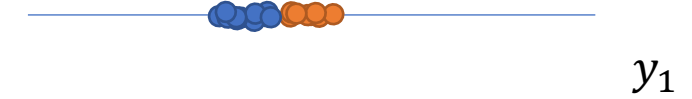
## □ LDA vs PCA



PCA



LDA





# Next class

## ❑ Next class is a flipped lecture

- Unsupervised learning – clustering

## ❑ Next week

- Tuesday 10/25: practical class on unsupervised learning
- Thursday 10/27: no in-person class

## ❑ Week after

- Paper presentations