

# BIOS 7747: BIOS 7747: Machine Learning for Biomedical Applications

## Training neural networks

Antonio R. Porras ([antonio.porras@cuanschutz.edu](mailto:antonio.porras@cuanschutz.edu))

Department of Biostatistics and Informatics

Colorado School of Public Health

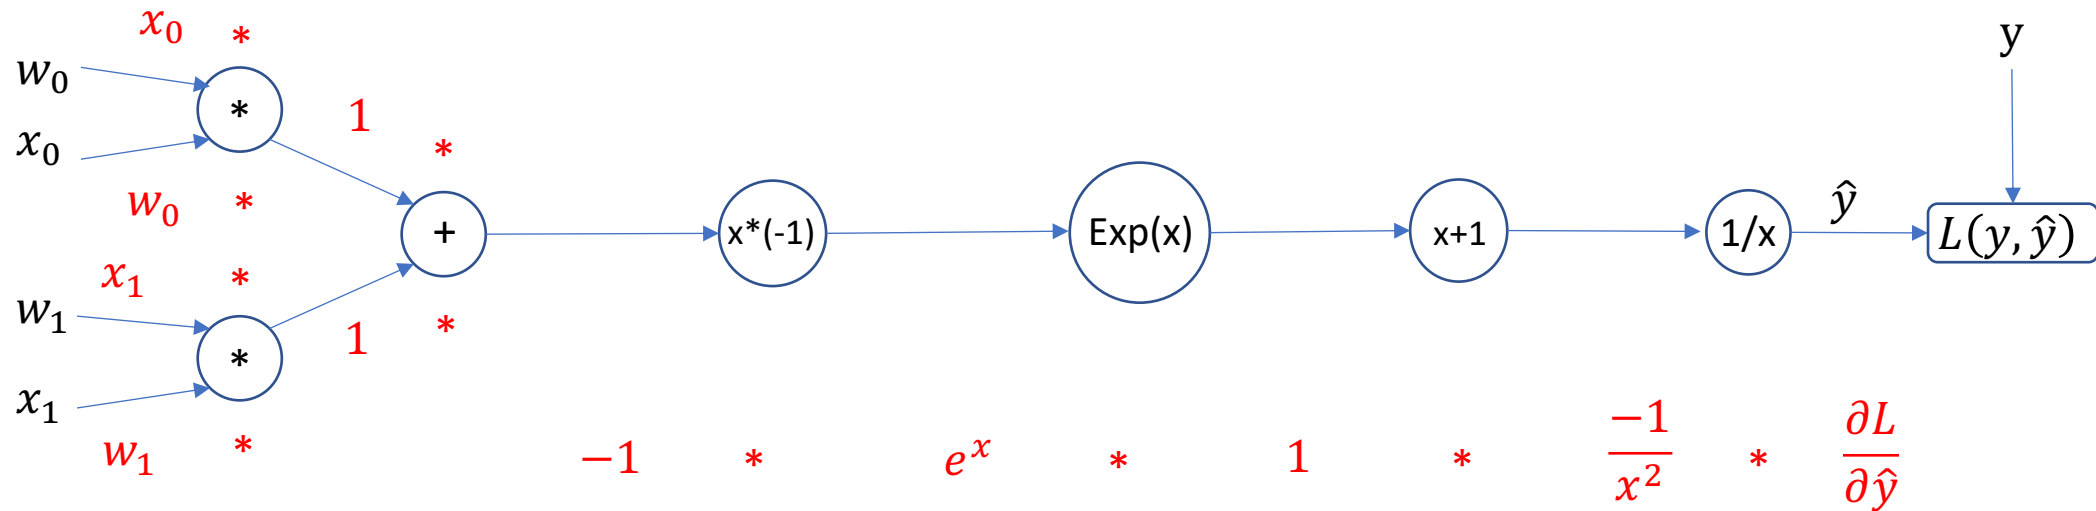
University of Colorado Anschutz Medical Campus

# Outline

- ❑ Vanishing and exploding gradients
- ❑ Activation functions
- ❑ Batch normalization
- ❑ Weight initialization
- ❑ Regularization
- ❑ Optimization and stop criteria

# Vanishing and exploding gradients

□ Let  $\hat{y} = f(x; w, b) = \frac{1}{1 + \exp(-(w_0 x_0 + w_1 x_1 + b))}$ , and  $L(y, \hat{y}) = y - \log(\hat{y}) - (1 - y)\log(1 - \hat{y})$



Situation 1: the gradient of a node is very large  $\longrightarrow$  Backpropagated gradients can “explode” and cause numerical instability

Situation 2: the gradient of a node is very small  $\longrightarrow$  Backpropagated gradients can “vanish” and prevent optimization

**It is very important to keep gradients in optimal ranges**

# Vanishing and exploding gradients

- ❑ Gradients are typically very variable between neurons, which makes it very hard to train deep neural networks
  - Both exploding and vanishing problems can occur during training on different neurons
- ❑ To prevent vanishing and exploding gradients:
  - All neurons have input/outputs within similar ranges
  - All neurons have gradients within similar ranges

} All data should have similar distributions


# Activation functions

## □ Sigmoid

In Pytorch  
`torch.nn.Sigmoid`

$$f(x) = \frac{1}{1 + e^{-x}}$$

$$f'(x) = f(x)(1 - f(x))$$



- Limited output range:  $[0, 1]$
- Zero-gradient problem if  $x$  is not in the linear region

# Activation functions


## □ Tanh (hyperbolic tangent)

In Pytorch  
`torch.nn.Tanh`

$$f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

$$f'(x) = \text{sech}(x)^2 = \frac{1}{\cosh(x)^2}$$

- Limited output range:  $[-1, 1]$
- Output data are centered around 0
- Zero-gradient problem if  $x$  is not in the linear region



# Activation functions


## □ ReLU

In Pytorch  
`torch.nn.ReLU`

$$f(x) = \max(0, x)$$

$$f'(x) = \begin{cases} 0 & \text{if } x < 0 \\ 1 & \text{otherwise} \end{cases}$$

- No gradient issues on activated neurons
- Computationally efficient
- Constant gradient speeds up convergence
- Zero-gradient problem on deactivated neurons



# Activation functions

## □ Leaky ReLU


$$f(x) = \begin{cases} x & \text{if } x \geq 0 \\ 0.01x & \text{otherwise} \end{cases}$$

$$f'(x) = \begin{cases} 1 & \text{if } x \geq 0 \\ 0.01 & \text{otherwise} \end{cases}$$

- Solves the problem of “dead” neurons with zero-gradient

### In Pytorch

`torch.nn.LeakyReLU(negative_slope=0.01)`





# Activation functions

## □ Parametric ReLU


$$f(x) = \begin{cases} x & \text{if } x \geq 0 \\ \alpha x & \text{otherwise} \end{cases}$$

$$f'(x) = \begin{cases} 1 & \text{if } x \geq 0 \\ \alpha & \text{otherwise} \end{cases}$$

- $\alpha$  can be learned for an “optimal” gradient from deactivated neurons

### In Pytorch

`torch.nn.PReLU(num_parameters=1, init=0.25)`



# Activation functions

## □ ELU (Exponential linear unit)


$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha(e^x - 1) & \text{otherwise} \end{cases}$$

$$f'(x) = \begin{cases} 1 & \text{if } x \geq 0 \\ \alpha e^x & \text{otherwise} \end{cases}$$

- $\alpha$  is a constant (hyperparameter)
- Unlike leaky ReLU, the deactivated values converge to a constant: less prone to overfitting

### In Pytorch

`torch.nn.ELU(alpha=1.0)`



# Batch normalization

## ❑ The problem of data variability:

- Operations with inputs or features with high variance will produce outputs of high variance and gradients with high variance
  - Training may be challenging
- Some activation functions only work well within limited ranges
  - e.g., sigmoid, tanh
- Other activation functions will not reduce the high variance in the data
  - e.g., ReLU, ELU
- Batch normalization provides a tool to control the variance in every layer

# Batch normalization

- ❑ Batch normalization attempts to create feature maps with zero-mean and unit-variance
- ❑ If  $x = [x^1, x^2, \dots, x^d]$  is a set of activated features in a specific layer, the batch-normalized feature values are:

$$\hat{x}^k = \frac{x^k - E[x^k]}{\text{std}(x^k)}$$

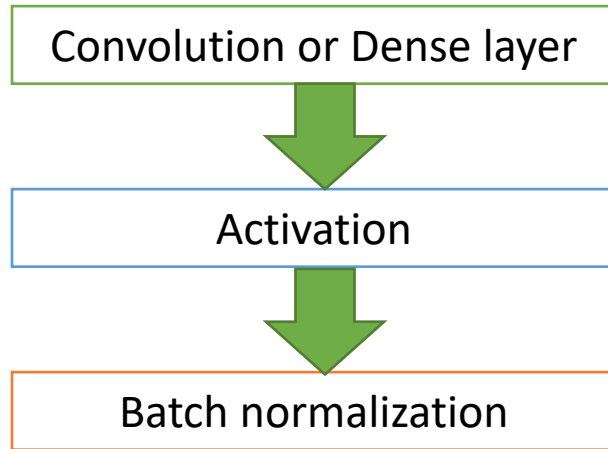
Note:  $E[x^k]$  and  $\text{std}(x^k)$  are the average and standard deviation of feature  $x^k$  in the training dataset

- ❑ Running mean and variances are normally used instead of minibatch-only data during training
  - $\text{RunninValue} = \text{momentum} * \text{learnedValue} + (1 - \text{momentum}) * \text{batchValue}$

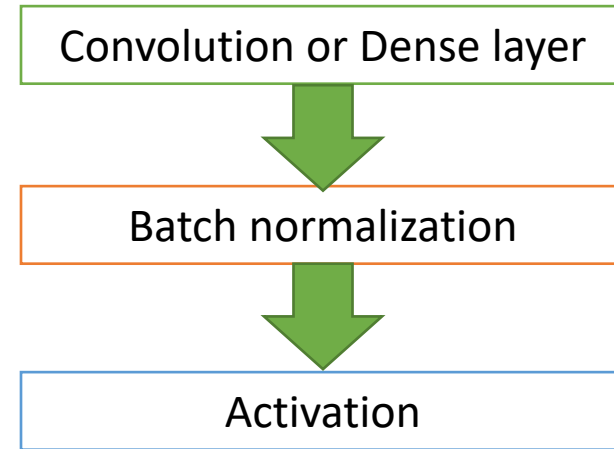
## In Pytorch

`torch.nn.BatchNorm1d(num_features, momentum=0.1)`

# Batch normalization



- Useful to limit the activation range (e.g., linear activations)



- Useful to keep non-linear activation functions in the adequate range (e.g., sigmoid, tanh)
- Will also limit the range of linear activations

# Batch normalization

- ❑ Helps preventing exploding and vanishing gradients by keeping stability in data ranges
- ❑ Reduces overfitting by changing feature values during training
  - Adds “noise” during training
- ❑ Because of changing values during training, optimization becomes harder

# Weight initialization

## □ Weight initialization

- Random initialization using normal or uniform distributions
  - If initialized weights are too small, then the variance of the input signal shrinks as it passes through each layer until it's too tiny to be useful [vanishing signal and gradients]
  - If initialized weights are too large, then the variance of the input signal grows as it passes through each layer until it's too large to be useful [exploding signal and gradients]
- How do we achieve a good weight initialization that preserves data variance?

# Weight initialization

## □ Xavier / Glorot initialization

- Goal: to keep the same variance after every layer:

$$\text{var}(y) = \text{var}(w_1x_1 + w_2x_2 + \dots + w_nx_n + b)$$

$$\text{var}(w_ix_i) = E(x_i)^2\text{var}(w_i) + E(w_i)^2\text{var}(x_i) + \text{var}(w_i)\text{var}(x_i)$$

Assuming that inputs and weights follow a Gaussian distributions

$$\text{var}(w_ix_i) = \text{var}(w_i)\text{var}(x_i)$$

Assuming zero-mean

$$\text{var}(y) = \text{var}(w_1)\text{var}(x_1) + \text{var}(w_2)\text{var}(x_2) + \dots$$

Substituting (b is a constant with zero variance)

$$\text{var}(y) = N_{\text{input}} * \text{var}(w_i)\text{var}(x_i)$$

Since all are identically distributed

$$\text{var}(w_i) = \frac{1}{N_{\text{input}}}$$

Since we want  $\text{var}(y) = \text{var}(x_i)$



# Weight initialization

## ❑ Xavier / Glorot initialization

- Normal Xavier initialization that keeps a stable variance in the forward pass:

$$w \in N\left(0, \frac{1}{N_{input}}\right)$$

- If we want to preserve the variance during backpropagation too:

$$w \in N\left(0, \frac{2}{N_{input} + N_{outputs}}\right)$$

- Xavier uniform distribution:

$$w \in [-x, x], \quad x = \sqrt{\frac{6}{N_{input} + N_{outputs}}}$$

## In Pytorch

`torch.nn.init.xavier_uniform_()` # Initialize weights


`torch.nn.init.xavier_normal_()` # Initialize weights

# Weight initialization

## □ Transfer learning:


- Reuse of previously trained weights for a different task

Task 1



Weight transfer

Task 2




# Weight initialization

## ❑ Transfer learning:

- Two options after initialization:

Train last layers only




- Fewer weights to train will reduce overfitting
- Features in the first layers may not be too relevant for the current problem

# Weight initialization

## ❑ Transfer learning:

- Two options after initialization:

Train entire network




- More weights to train will be more likely to cause overfitting
- Features will be problem-specific and may provide higher accuracy

# Reducing overfitting: Regularization

## □ Regularization

- Goal: to improve the performance of the network on unseen data (generalizability) by reducing overfitting
- The performance improvement on unseen data usually comes at the expense of a worse training performance



# Reducing overfitting: Regularization

## □ Weight decay:

$$\arg \min_W \frac{1}{N} \sum_{n=1}^N L(\hat{y}_n, y_n) + \frac{\lambda}{\beta} R(W)$$

Regularization weight  $\lambda$  points to the numerator of the regularization term.

Number of weights  $\beta$  points to the denominator of the regularization term.

$R(W)$  is boxed and labeled as the Regularization term.

- L2 (Ridge) regularization:  $R(W) = \sum_{\forall i} w_i^2$  —————> Keeps weights uniformly small
- L1 (Lasso) regularization:  $R(W) = \sum_{\forall i} |w_i|$  —————> Tends to eliminate weights (deactivate neurons)
- L2 is more commonly used

# Reducing overfitting: Regularization

## □ Weight decay:

### In Pytorch

#### L1-normalization

```
l1_penalty = torch.nn.L1Loss()  
reg_loss = 0  
for param in model.parameters():  
    reg_loss += l1_penalty(param)  
Loss += reg_loss
```

#### L2-normalization

```
torch.optim.SGD(..., weight_decay=0)
```

# Reducing overfitting: Dropout

## □ Dropout

- Goal: To reduce overfitting by minimizing co-adaptation
- Co-adaptation: when multiple neurons extract the same or very similar information from the previous layer
- The problems of co-adaptation:
  - High correlation between extracted features (less meaningful independent information)
  - Extracting the same information in different neurons gives a higher significance to specific features that may only be that significant in the training dataset (promotes overfitting)




# Reducing overfitting: Dropout

## □ Dropout


- Dropout consists in eliminating (zeroing) the values of random neurons at each training step
- Dropout rate ( $r_D$ ): fraction of neurons deactivated
- Keep rate:  $r_K = 1 - r_D$
- To preserve data variance, the activated values of the neurons kept is divided by  $r_K$
- At test time all neurons are preserved

# Reducing overfitting: Dropout

## □ Dropout



(a) Standard Neural Net



(b) After applying dropout.

# Reducing overfitting: Dropout

## □ Dropout

- When enough training data are available, a dropout rate of 50% is optimal
- Dropout can make training much more difficult
- One normally must find a balance between the amount of dropout at each layer, the training and the test performances

### In Pytorch

`torch.nn.Dropout(p=0.5)`


# Optimization algorithms

## □ Regular gradient descent approach:

$$w_t = w_{t-1} - \boxed{\alpha} \nabla L(w_{t-1})$$

Learning rate


- May provide good results if gradient estimations are accurate



# Optimization algorithms

## □ Stochastic gradient descent (SGD):

- Gradients are computed independently for every minibatch and training may be difficult
- The smaller the minibatch, the harder the training



**In Pytorch**

`torch.optim.SGD(params, lr)`

# Optimization algorithms

## □ Stochastic gradient descent (SGD) with momentum:

$$V_t = \beta V_{t-1} + (1 - \boxed{\beta}) \nabla L(W_t)$$

Momentum

$$w_{t+1} = w_t - \alpha V_t$$


- Higher momentum achieves more stable gradients
- $\beta = 0$  is equivalent to regular SGD
- Typical values  $\geq 0.9$

### In Pytorch

`torch.optim.SGD(params, lr, momentum=0)`

# Optimization algorithms

- Stochastic gradient descent (SGD) with momentum:



# Optimization algorithms

## □ Root mean square propagation (RMSProp):

$$S_t = \beta S_{t-1} + (1 - \beta)(\nabla L(w_t) \nabla L(w_t))$$

$$w_{t+1} = w_t - \alpha \frac{\nabla L(w_t)}{\sqrt{S_t + \epsilon}}$$

- Gradients are normalized to a uniform scale that is updated with momentum

In Pytorch

$\beta$   
`torch.optim.RMSprop(params, lr=0.01, alpha=0.99)`



# Optimization algorithms

## □ Adaptive momentum estimation (Adam):

- Combines RMSProp and SGD with momentum

$$V_t = \beta_1 V_{t-1} + (1 - \beta_1) \nabla L(w_t)$$

$$S_t = \beta_2 S_{t-1} + (1 - \beta_2) (\nabla L(w_t) \nabla L(w_t))$$


$$w_{t+1} = w_t - \alpha \frac{V_t}{\sqrt{S_t + \epsilon}}$$

### In Pytorch

`torch.optim.Adam(params, lr=0.001, betas=(0.9, 0.999))`

# Optimization algorithms


## □ Example



# Optimization algorithms

## Thoughts:


- Lower training loss is not always the best-case scenario
- One optimizer may provide lower training loss at the expense of higher degree of overfitting
- Many variables and design decisions affect training and comparisons are difficult (normalization, activation functions, dropout [rate], regularization, etc.)



# Optimization algorithms

## □ Learning rate

- Hyperparameter that controls how much of the gradient is used to update parameters
- High values may cause divergence from the optimal solution
- Low values will likely cause getting trapped in local minima



# Optimization algorithms


## □ Learning rate

- Step decay:
  - Reduction by a fraction every  $k$  iterations
- Simple continuous learning rate decay:

$$\alpha_t = \frac{\alpha_0}{1 + kt}$$

- Exponential decay:

$$\alpha_t = \alpha_0 e^{-kt}$$



# Optimization algorithms

## □ Learning rate


- Adaptive learning rate: AdaGrad

$$\alpha^{(i)} = \frac{\alpha_0^{(i)}}{\sqrt{\epsilon I + \sum_{\tau=1}^t g_{\tau}^{(i)2}}}$$

Parameter gradient

Accounts for total gradient magnitude during optimization


- There is a different learning rate for each parameter




# Optimization algorithms

## □ Learning rate


- Cyclical learning rate



Basic schedule



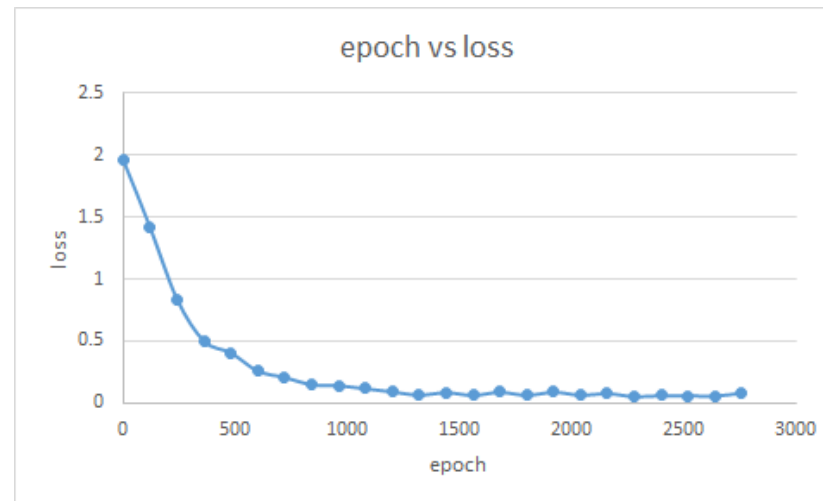
Triangular schedule with exponential decay



# Optimization algorithms

## □ Stop criterion

- Fixed number of epochs
  - Problem-, model- and dataset-specific approach
  - Large number can ensure convergence






# Optimization algorithms

## □ Stop criterion

- Threshold on loss function changes
  - The number of iterations will be different on every model, problem and dataset
  - Can significantly reduce the number of iterations



# Optimization algorithms

## □ Stop criterion

- Maximum in validation accuracy
  - Requires a representative validation dataset
  - Reduces the size of the dataset used for training (overfitting is more likely) at the expense of having a direct estimation of the amount of overfitting

