

# BIOS 7747: Machine Learning for Biomedical Applications

## Convolutional neural networks

Antonio R. Porras ([antonio.porras@cuanschutz.edu](mailto:antonio.porras@cuanschutz.edu))

Department of Biostatistics and Informatics  
Colorado School of Public Health  
University of Colorado Anschutz Medical Campus

# Outline

## □ Introduction to convolutional neural networks

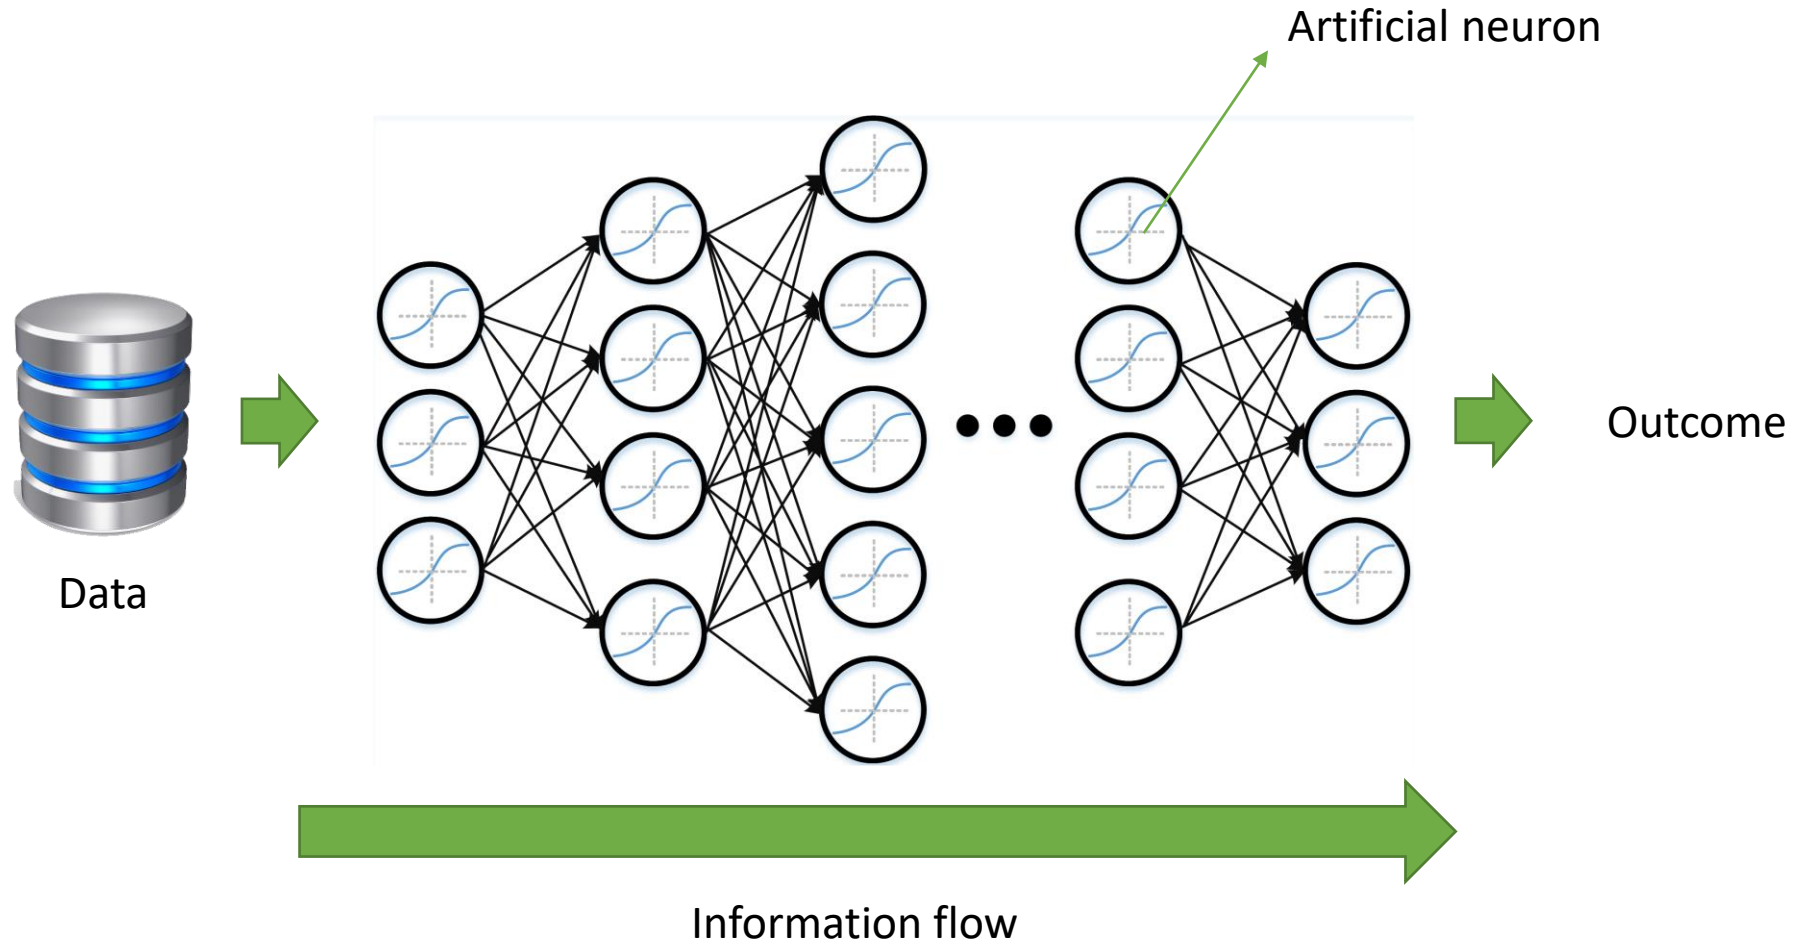
- Convolutions
- Downsampling
- Activation
- Full architecture
- Training

## □ Architectures and basic design concepts

## □ Fully convolutional networks

# Introduction to convolutional neural networks

- In previous class: artificial neural networks



# Introduction to convolutional neural networks

## ❑ What happens if we have too much data?

- Thousands or millions of observations?
  - It does not affect the network architecture
  - It will likely decrease overfitting and build more robust models
- Thousands or millions of features?
  - It will require a much wider network
  - It will likely increase overfitting (many more weights)



# Introduction to convolutional neural networks

- ❑ Images (and temporal signals) usually have a high number of pixels (temporal samples)
  - 2D images with size 200x200: 40,000 features per image
  - 3D images with size 200x200x200: 8,000,000 features per image
  - Signal sampled at 120Hz for 5 minutes: 36,000 features per signal

- ❑ Number of parameters needed only in the first hidden layer:

$$\underbrace{\#Neurons * \#Features}_{w} + \underbrace{\#Neurons}_{b}$$

- ❑ To evaluate one image using a network with only 1,000 neurons in one hidden layer we would need :

$$4 * (40,000 + 1,000 * 40,000 + 1,000 + 1,000 * 1 + 1) \approx 160MB$$

image                      hidden layer                      output neuron

- ❑ More realistic scenario of a narrow 8-layer network :

- $4 * (40,000 + 40,000 * 40,000 + 40,000 + 40,000 * 30,000 + 30,000 + 30,000 * 20,000 + 20,000 + 20,000 * 10,000 + 10,000 + 10,000 * 2,000 + 2,000 + 2,000 * 512 + 512 + 512 * 128 + 128 + 128 * 1 + 1) \approx 14GB$

# Introduction to convolutional neural networks

## □ Memory needs:

Memory for model parameters

Memory for outputs

Memory for parameter gradients

Memory for error and losses

Memory for optimizer's momentum

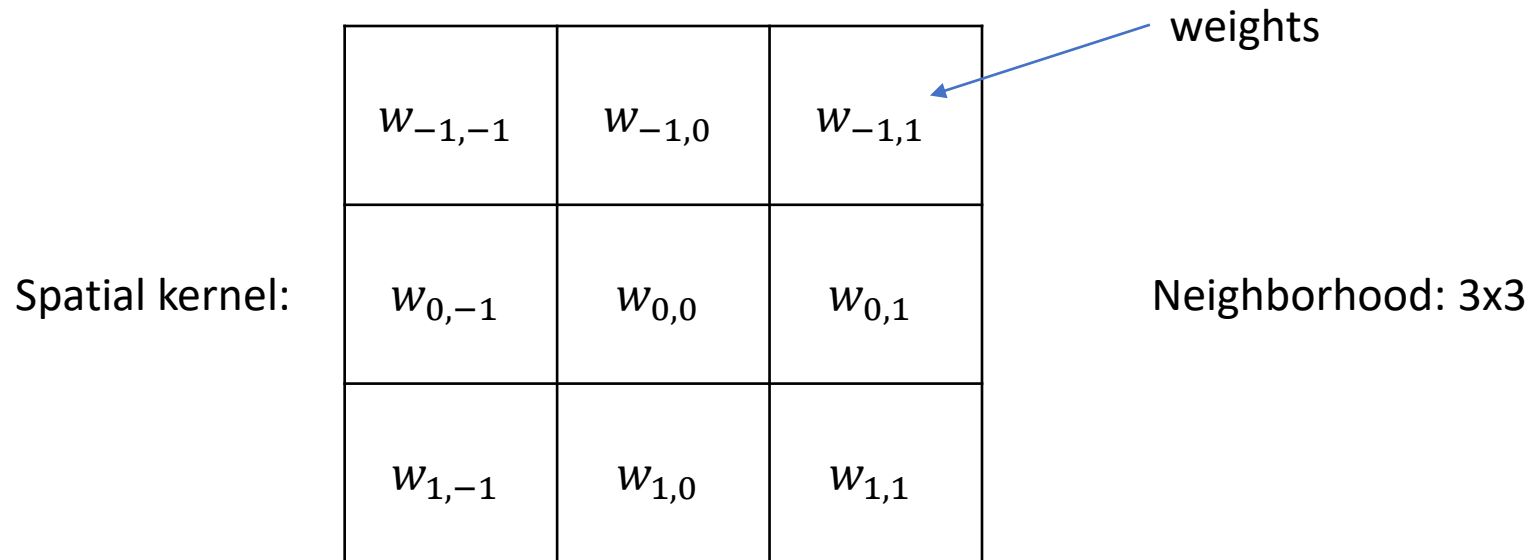
Memory for operations

Library overhead, other resources, OS management, etc.



# Convolutions

- ❑ Spatial operation that aggregates the information in a specific neighborhood
- ❑ Defined using a base kernel and a convolution operation



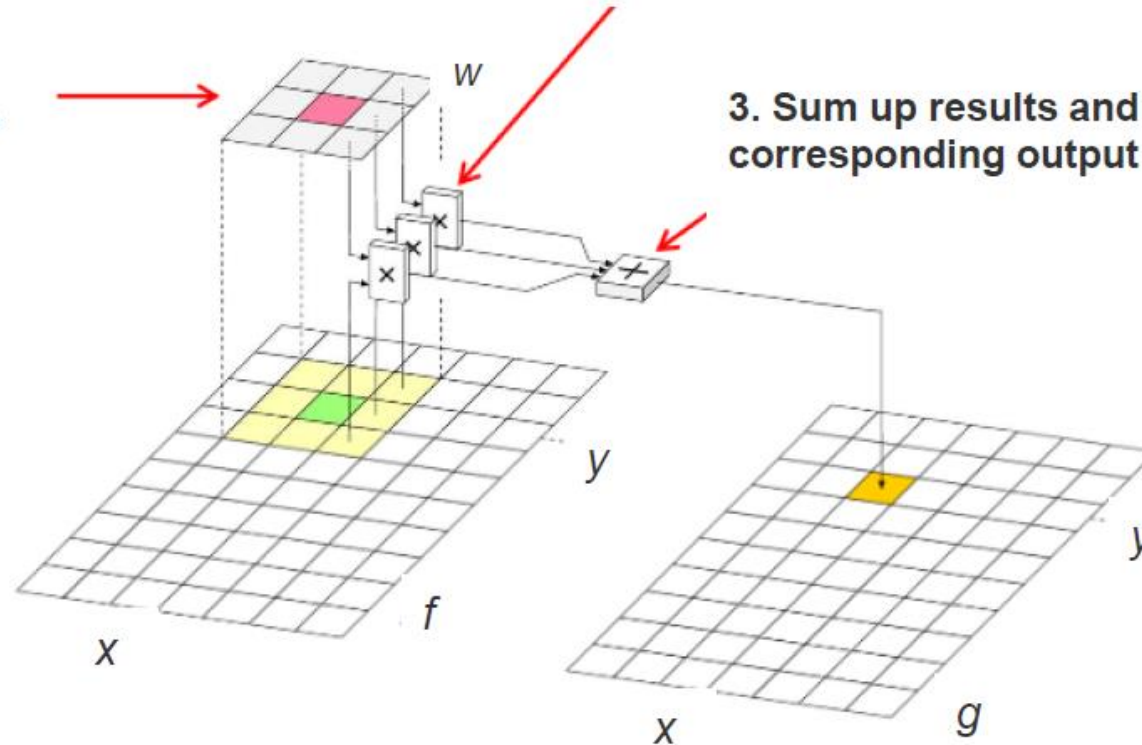
# Convolutions

For each image position  $(x,y)$ :

1. Move filter matrix  $w$  over image such that  $w(0,0)$  coincides with current image position  $(x,y)$

2. Multiply all filter coefficients  $w(s,t)$  with corresponding pixel  $f(x+s,y+t)$

3. Sum up results and store sum in corresponding output image  $g(x,y)$

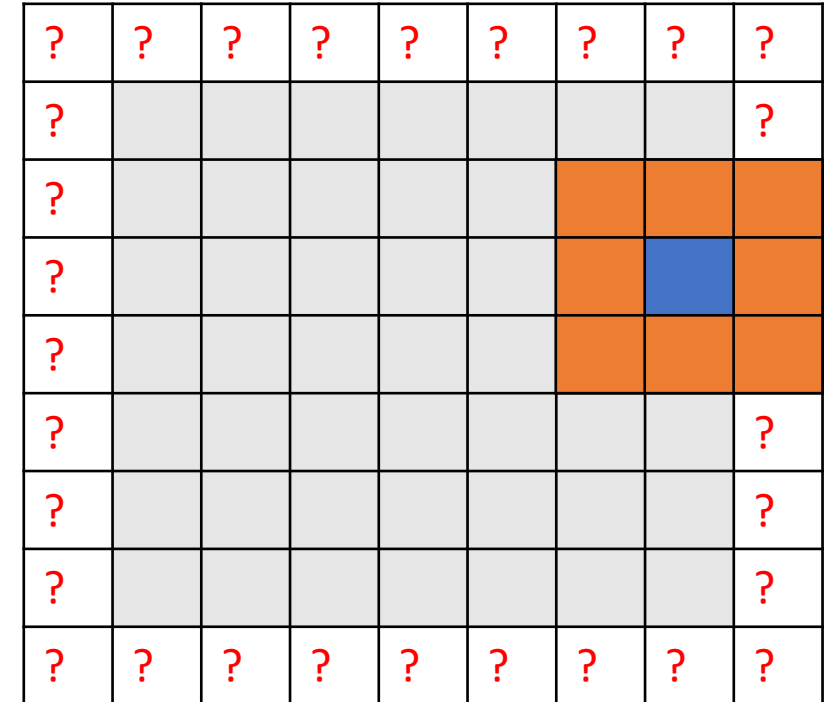
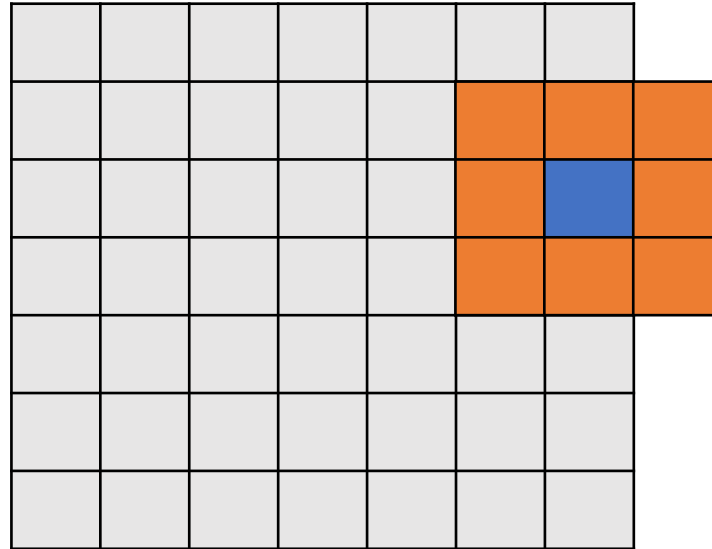
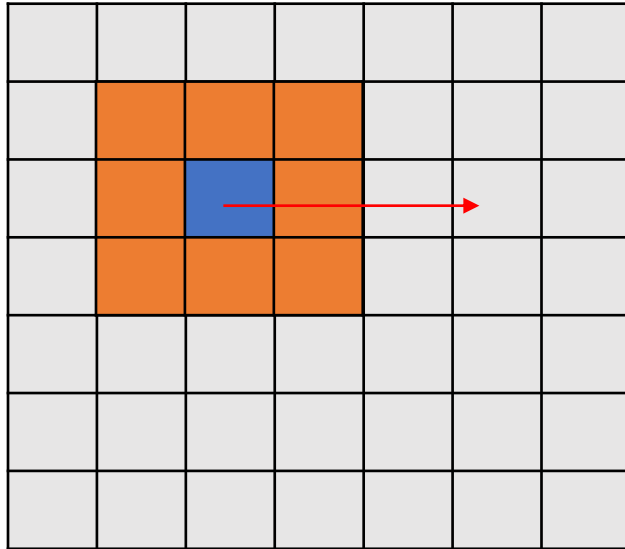


Cross-correlation of image  $f$  and filter  $w$ :  $g(x,y) = \sum_{i,j=-1}^1 f(x+i,y+j) * w(i,j)$



# Convolutions

## □ Padding



- Zero-padding
- Mirroring
- Replication
- Circle-padding
- Border removal

# Convolutions

- ❑ Visual (or temporal) patterns consists in:
  - Intensity value in a continuous region (brighter or darker)
  - Intensity change between regions (edge strength)
- ❑ Convolutions can provide information about:
  - Regional intensity  
[aka smoothing or (weighted) average filter]
  - Edge information  
[aka sharpening or differential filter]

$$\frac{1}{9} * \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline -\frac{1}{8} & -\frac{1}{8} & -\frac{1}{8} \\ \hline -\frac{1}{8} & 1 & -\frac{1}{8} \\ \hline -\frac{1}{8} & -\frac{1}{8} & -\frac{1}{8} \\ \hline \end{array}$$

# Convolutions

## ❑ Convolution vs. cross-correlation

Cross-correlation:  $g(x, y) = \sum_{i,j=-1}^1 f(x + i, y + j) * w(i, j)$

Convolution:  $g(x, y) = \sum_{i,j=-1}^1 f(x + i, y + j) * w(-i, -j)$

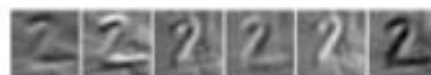
- ❑ To implement a convolution, a 180° rotation must be applied to the kernel
- ❑ They are only equivalent in symmetric kernels
- ❑ Most libraries implement cross-correlations, not convolutions

# Convolutions

## □ Example of convolutions



conv1



conv1



conv1

# Convolutions

- Convolutional neural networks use convolutional filters to calculate spatial or temporal features

- Convolution operations replace the linear decision function of the perceptron

$$z = wx + b \quad \longrightarrow \quad z = w * I + b$$

- Comparison:

- Parameters required for linear function:  $\#Neurons \times \#Features + \#Neurons$
- Parameters required for convolution:  $\#Filters \times FilterSize + \#Filters$

- The number of convolution parameters does not depend on the spatial size (or signal length)

# Convolutions

## ❑ Example: Extraction of 10 features of a 200x200 image

- Fully connected layer with 10 neurons:
  - Number of parameters:  $10 \times (200 \times 200) + 10 = 400,010$  parameters
  - Every perceptron combines all image information (most information will not be relevant so there are very high chances of overfitting)
  - Output: 10 features
- Convolutional network with 10 filters:
  - Region size: 5x5 neighborhood
  - Number of parameters  $10 \times 25 + 10 = 260$  parameters
  - Every feature combines only regional information
  - Output: 10 features at each pixel (10x200x200)

# Downsampling

- ❑ Downsampling: reduces the amount of data (dimensionality reduction)
  - Increases robustness to slight changes in rotation and translation
  - Convolutions of lower resolution images with same kernels aggregate information from larger regions
- ❑ Two main approaches to downsampling

## Convolution with strides

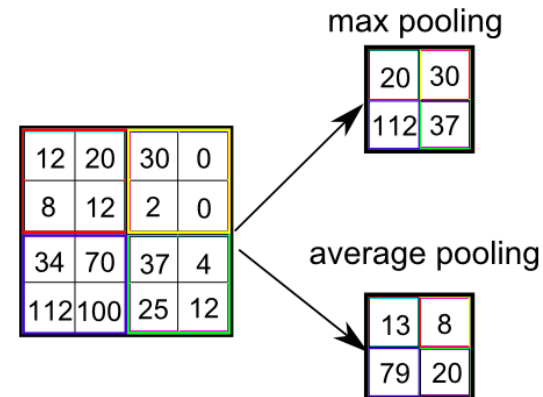
0 <sub>2</sub>	0 <sub>0</sub>	0 <sub>1</sub>	0	0	0	0
0 <sub>1</sub>	2 <sub>0</sub>	2 <sub>0</sub>	3	3	3	0
0 <sub>0</sub>	0 <sub>1</sub>	1 <sub>1</sub>	3	0	3	0
0	2	3	0	1	3	0
0	3	3	2	1	2	0
0	3	3	0	2	3	0
0	0	0	0	0	0	0

2x2 strides

1	6	5
7	10	9
7	10	8

- Faster

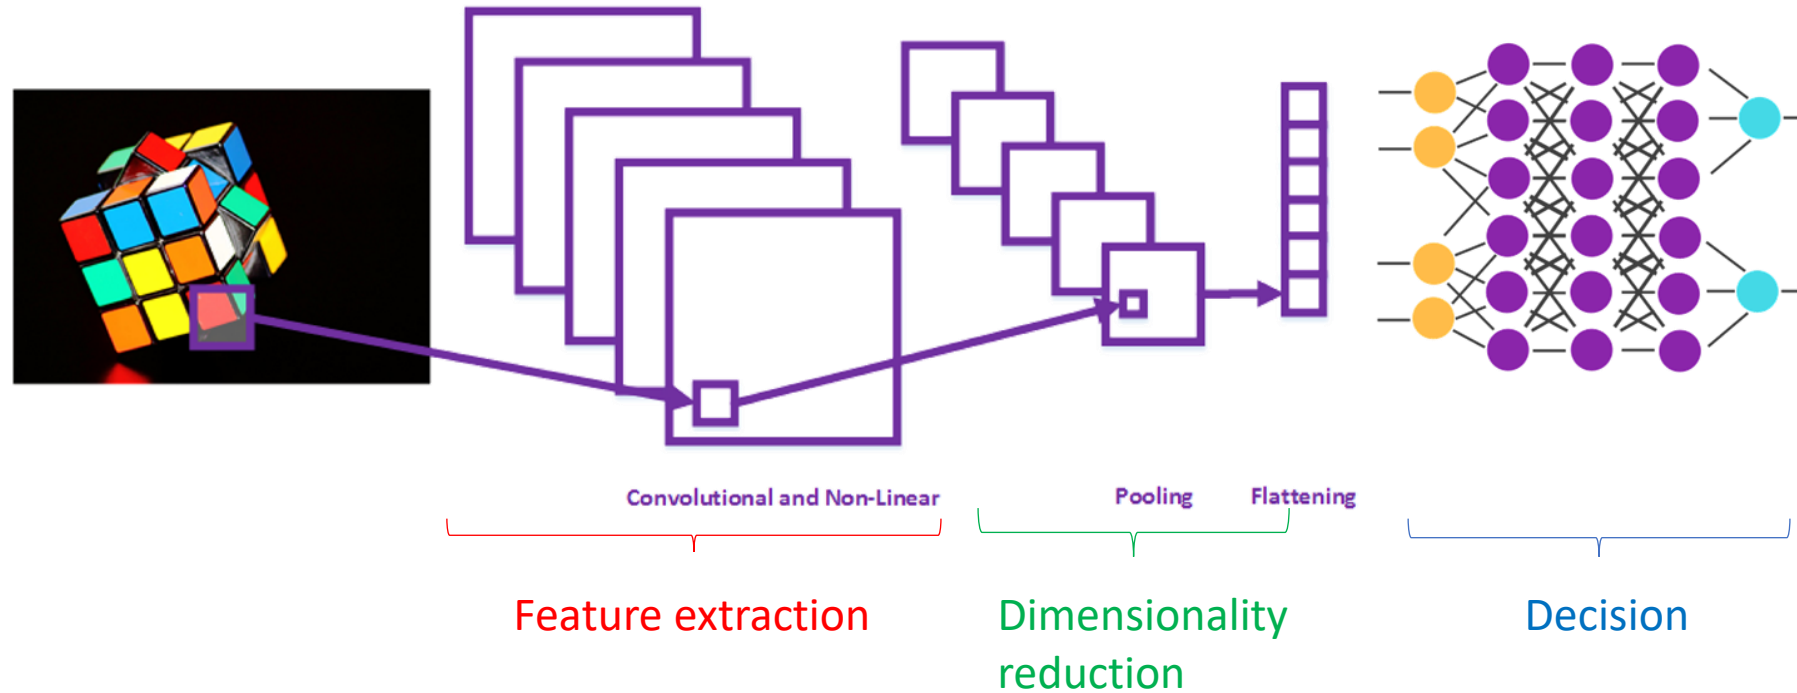
## Pooling (after convolution)



2x2 pooling with 2x2 strides

# Downsampling

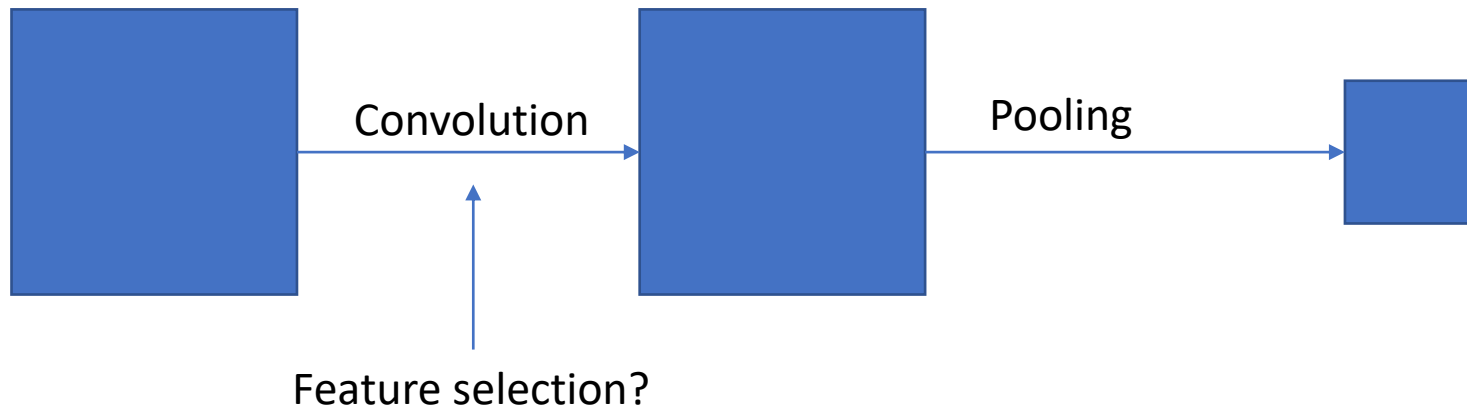
- Typical convolutional neural network



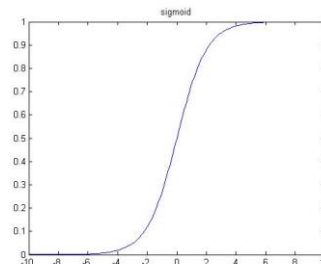


# Activation

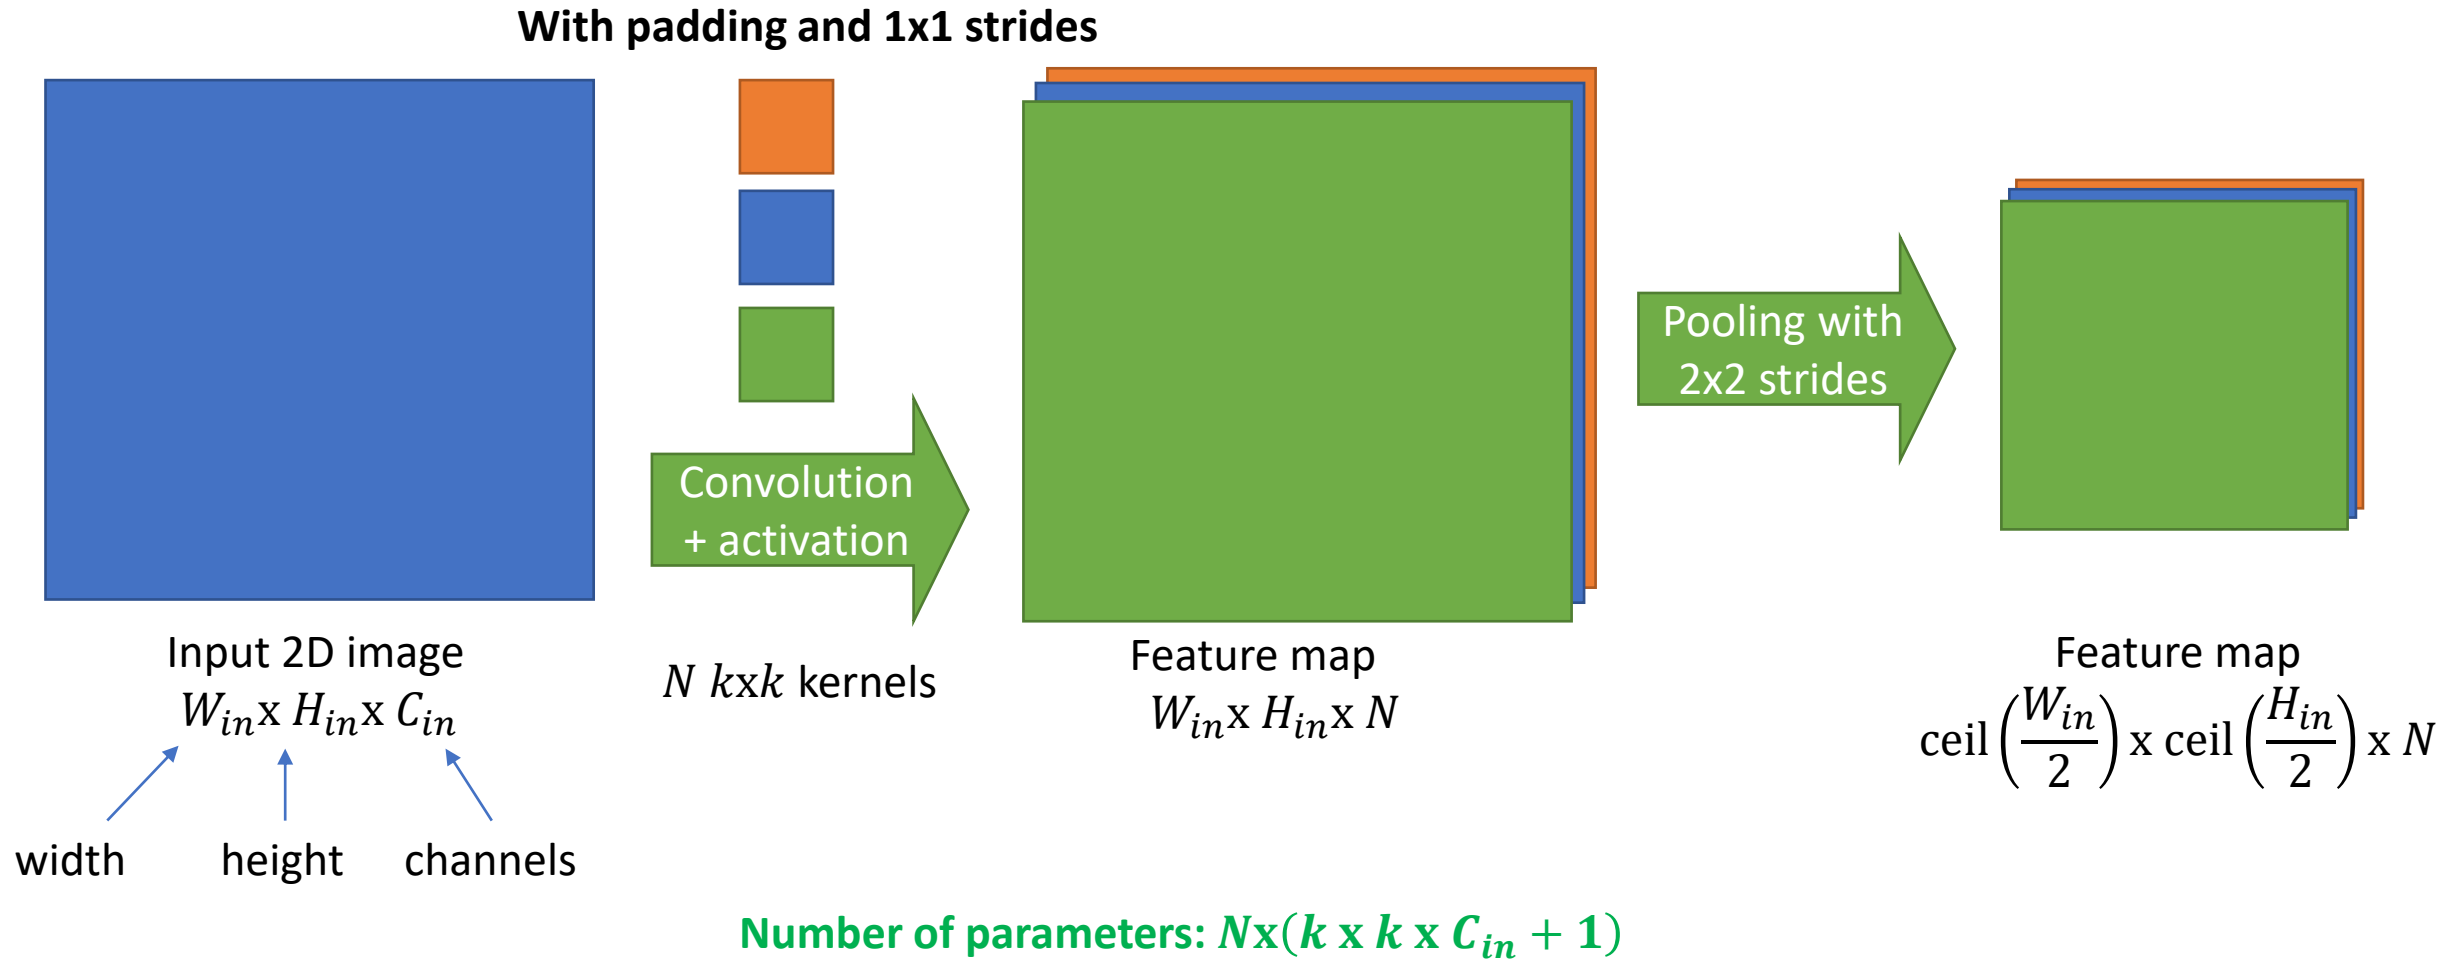
- Although pooling can eliminate meaningless features, it needs to either choose between different potentially meaningful features (max pool) or aggregate potentially meaningless features (average pooling)



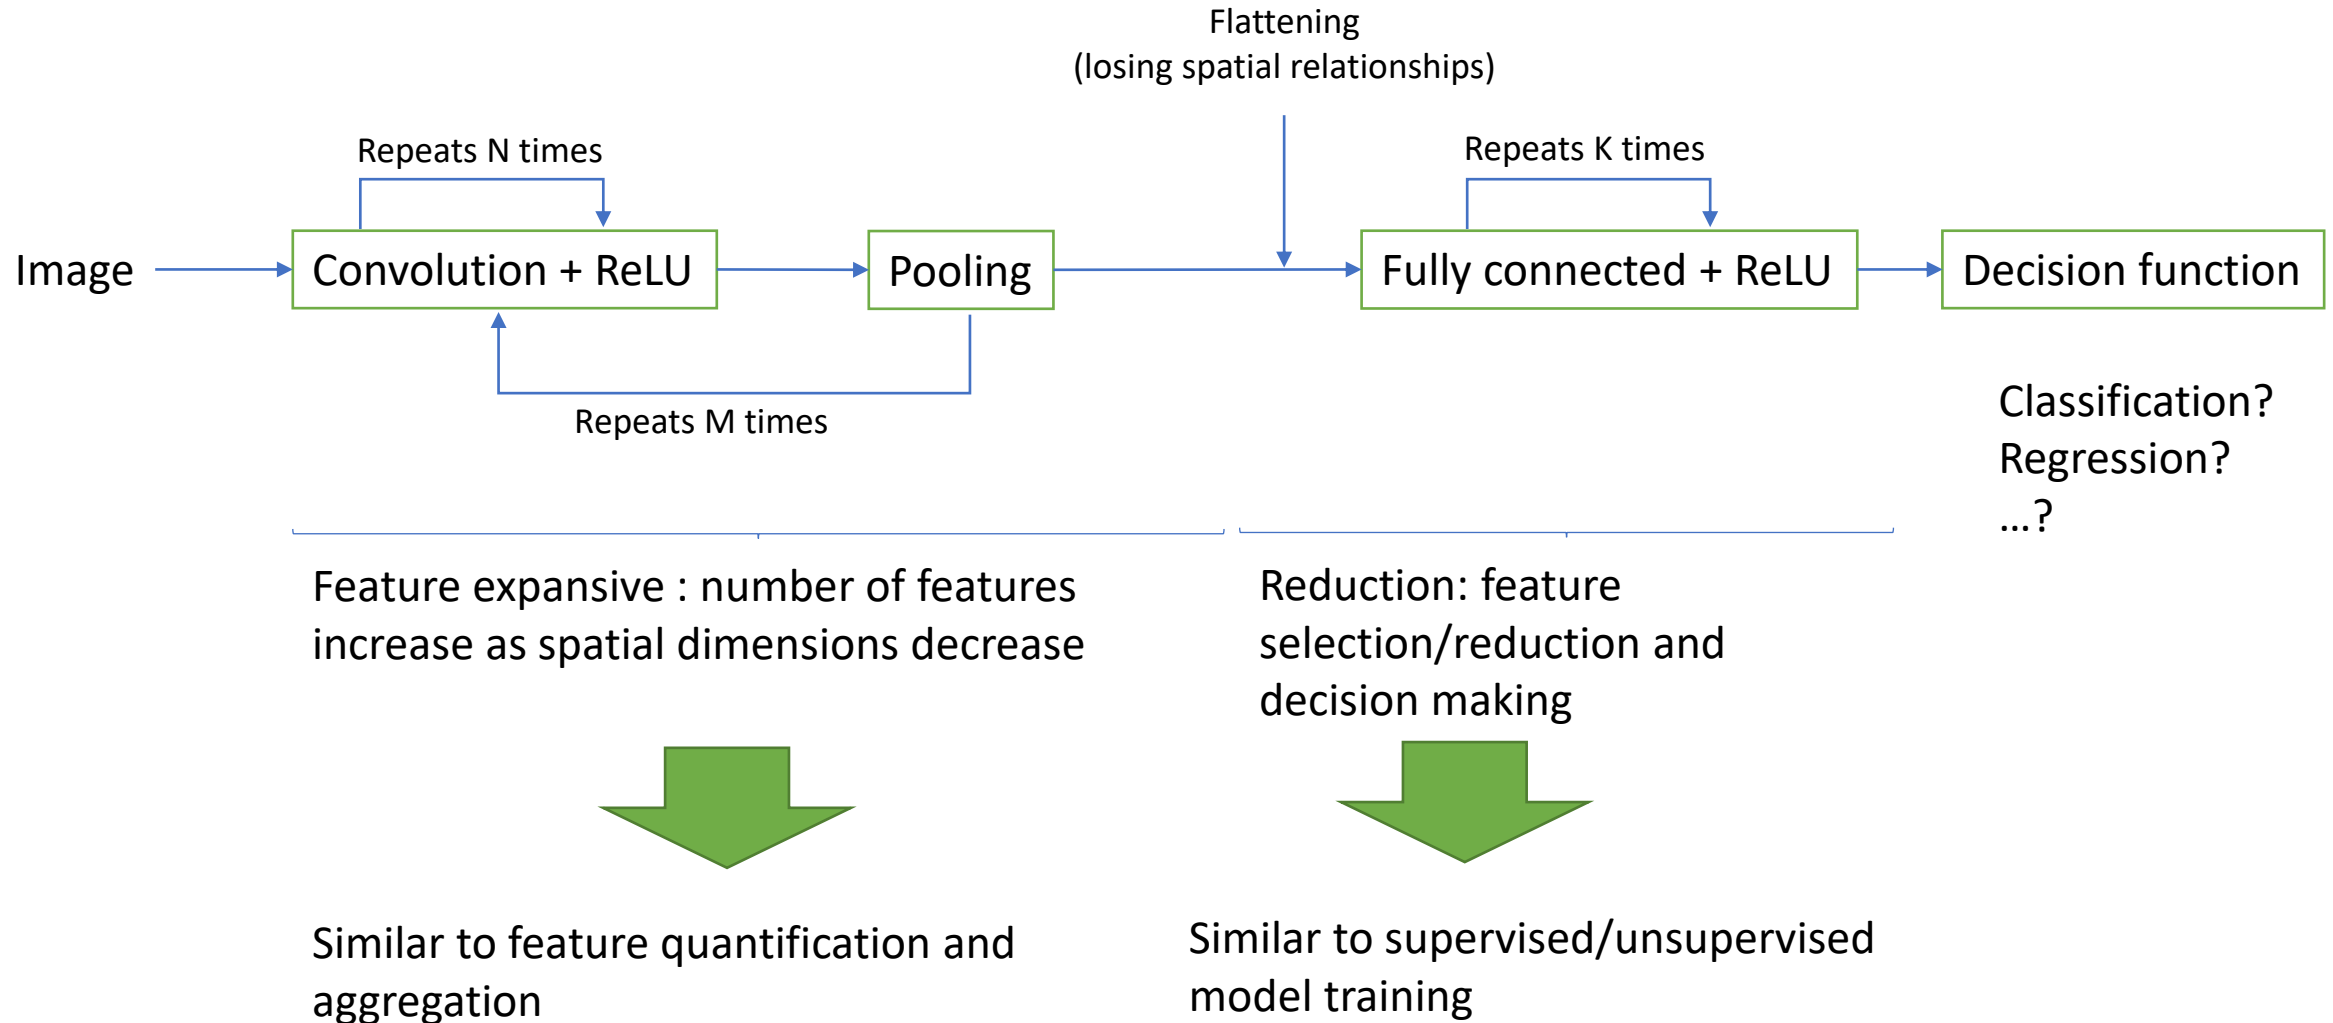
- The activation function
  - An activation enables kernels to learn how to zero-out irrelevant spatial patterns



# Full architecture



# Full architecture



# Training

- How can we backpropagate a ~~convolution~~ cross-correlation operation?

$$z = w * I + b$$

Single channel

3x3 kernel

$$z(i, j) = \sum_{u=-1}^{-1} \sum_{v=-1}^{-1} I(i - u, j - v) w(u, v) + b$$

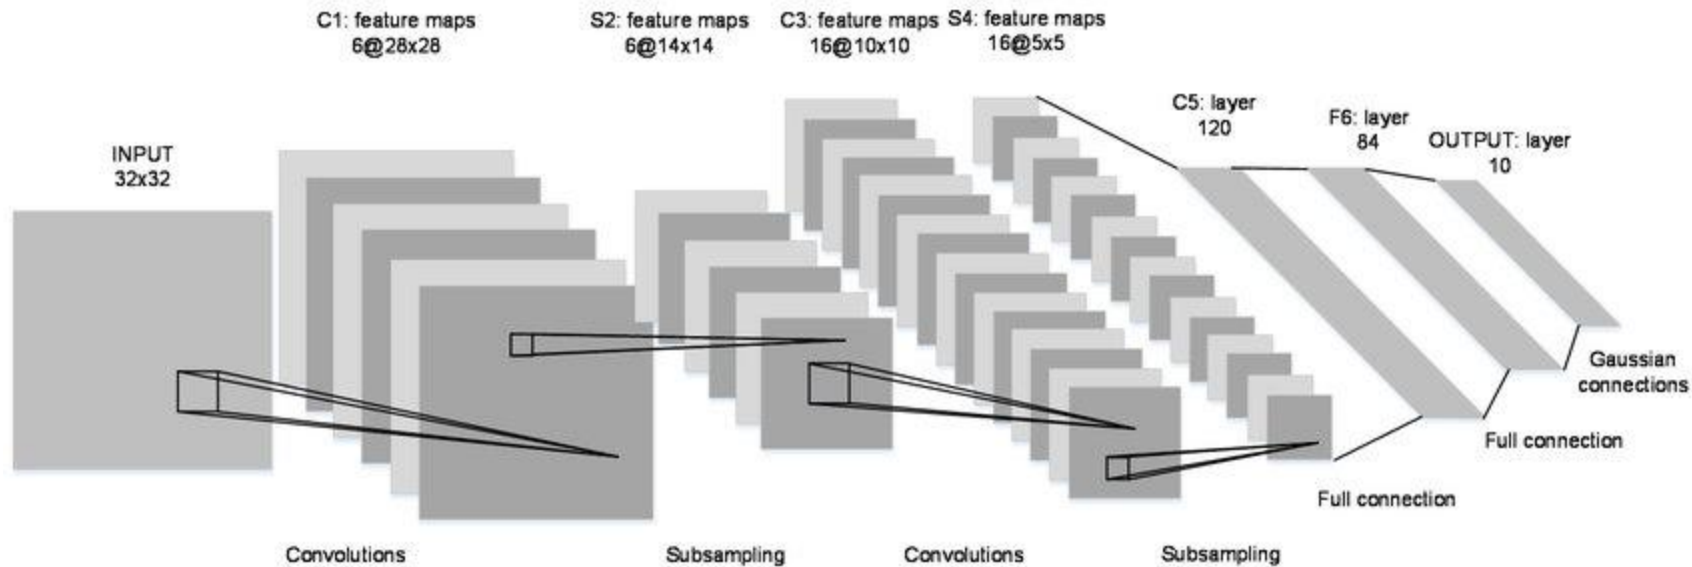
Derivative at location  $(i, j)$ :

$$\frac{\partial L}{\partial w}(i, j) = \frac{\partial L}{\partial z(i, j)} \frac{\partial z(i, j)}{\partial w} = \frac{\partial L}{\partial z(i, j)} I(i - 1 : i + 1, j - 1 : j + 1)$$

$$\frac{\partial L}{\partial b}(i, j) = \frac{\partial L}{\partial z(i, j)}$$

# Architectures and basic design concepts

## □ LeNet-5

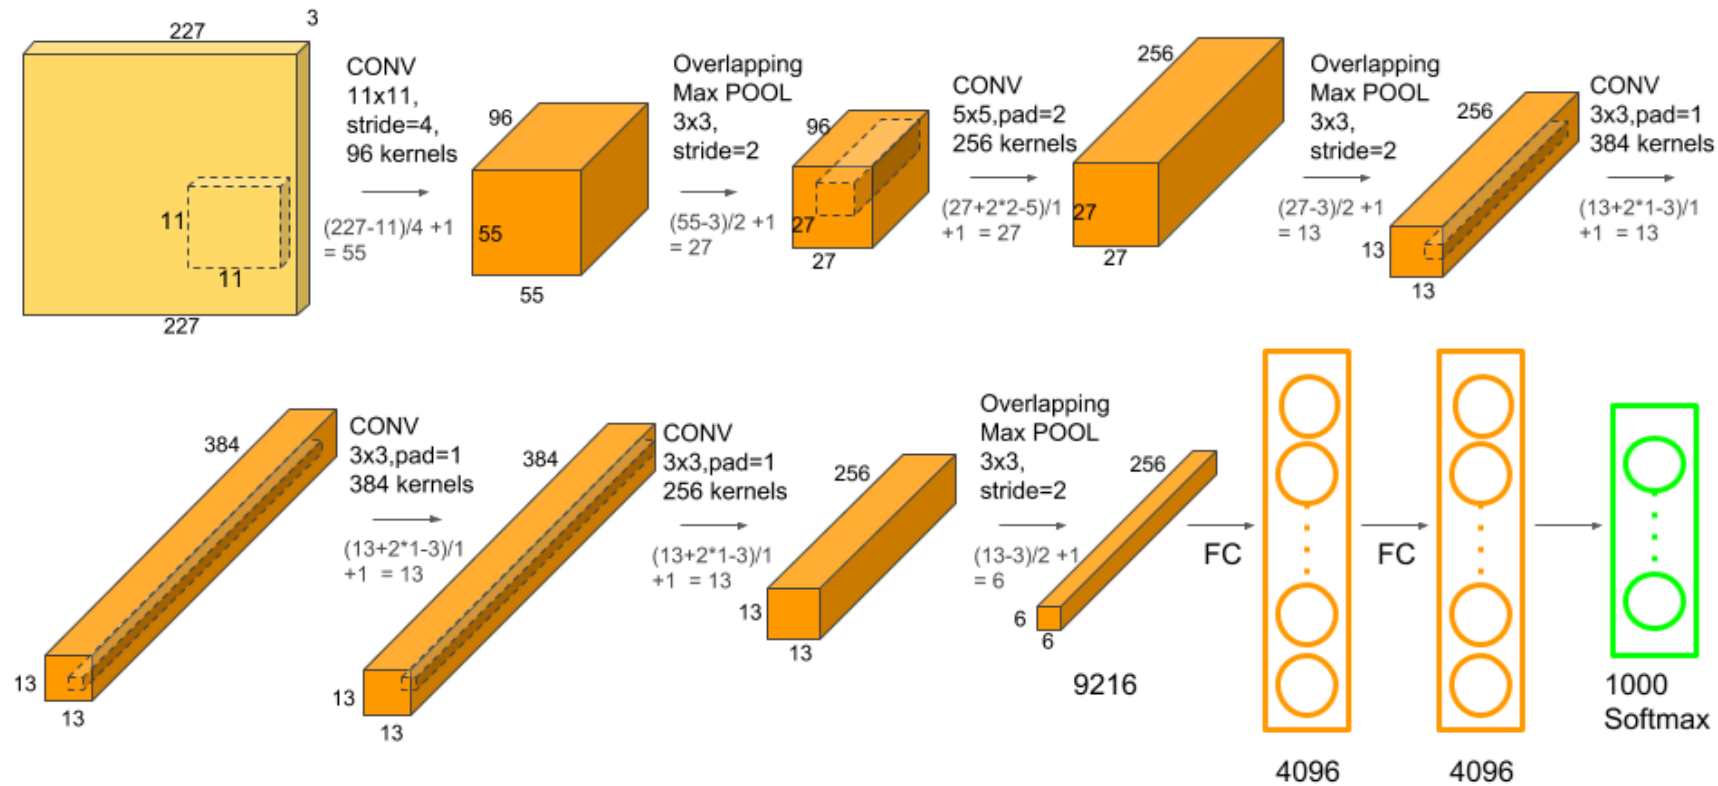


5x5 convolution (1 stride)  
2x2 pooling (2 strides)

Y. LeCun, *et al.*, "Gradient-based learning applied to document recognition", *Proceedings of the IEEE*, 1998

# Architectures and basic design concepts

## □ AlexNet



$$\sigma(\vec{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

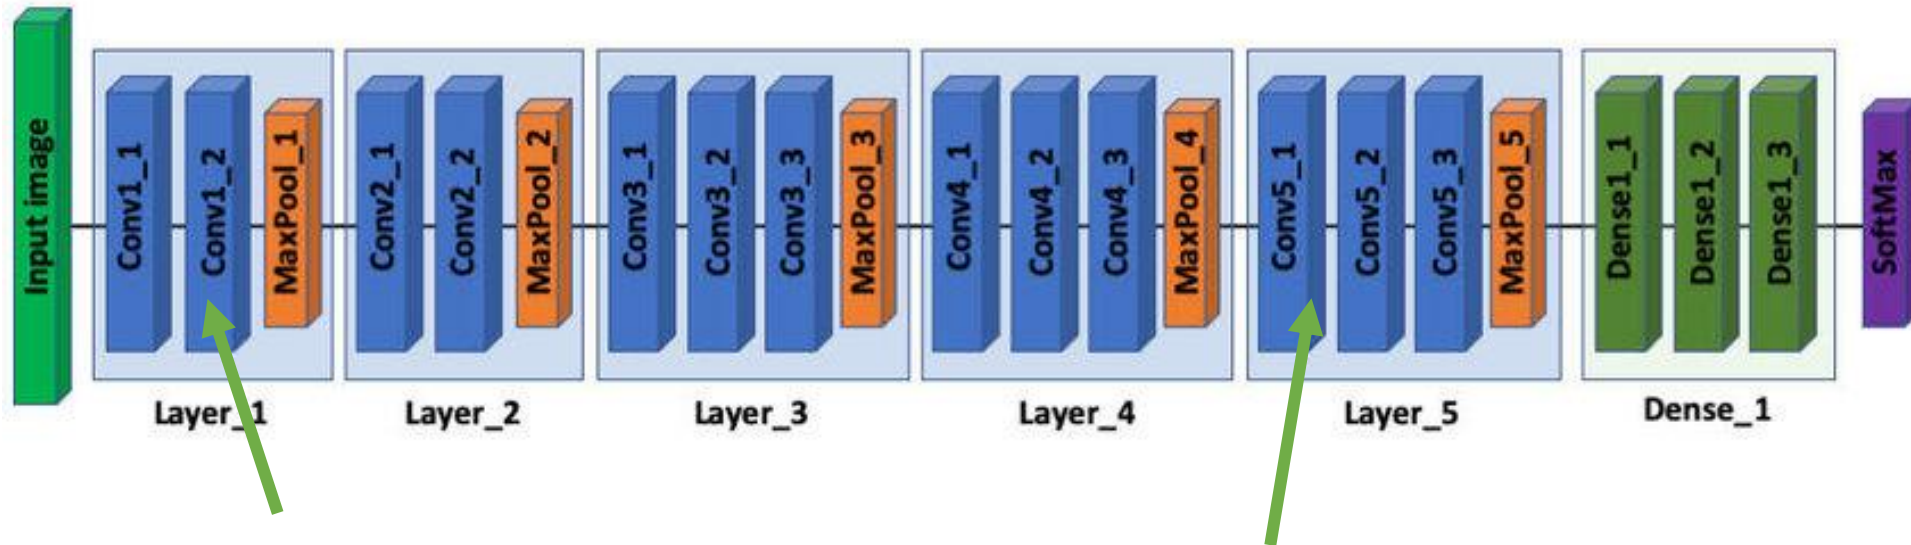
Softmax activation for disjoint classification

A. Krizhevsky et al., "ImageNet Classification with Deep Convolutional Neural Networks", *NIPS*, 2012

# Architectures and basic design concepts

## □ VGG-16

All convolutions are 3x3, 1x1 stride, 1 padding  
All max pooling layers are 2x2 with 2x2 strides



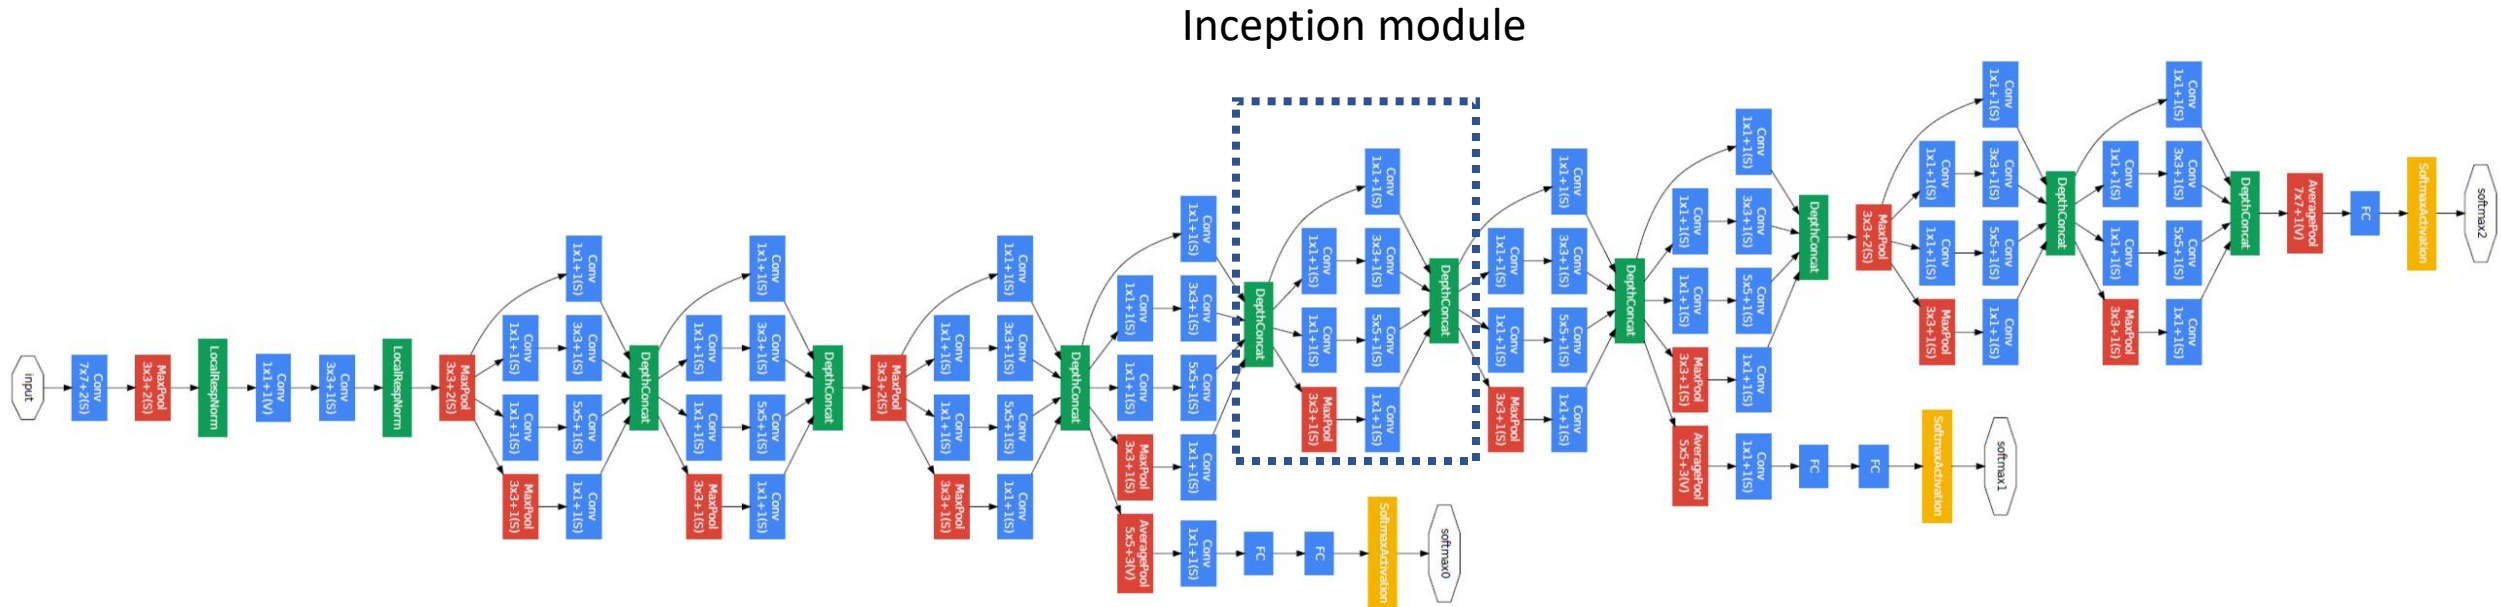
Why two 3x3 kernels? It covers the same space than a 5x5 kernel with less parameters ( $3 \times 3 \times 2 = 18$  vs.  $5 \times 5 = 25$ ).

Three 3x3 kernels (27 parameters) covers the same space than one 7x7 kernel (49 parameters).

K. Simonyan, A. Zisserman, "Very Deep Convolutional Networks for Large-Scale Image Recognition", *ICLR*, 2015

# Architectures and basic design concepts

## □ GoogLeNet



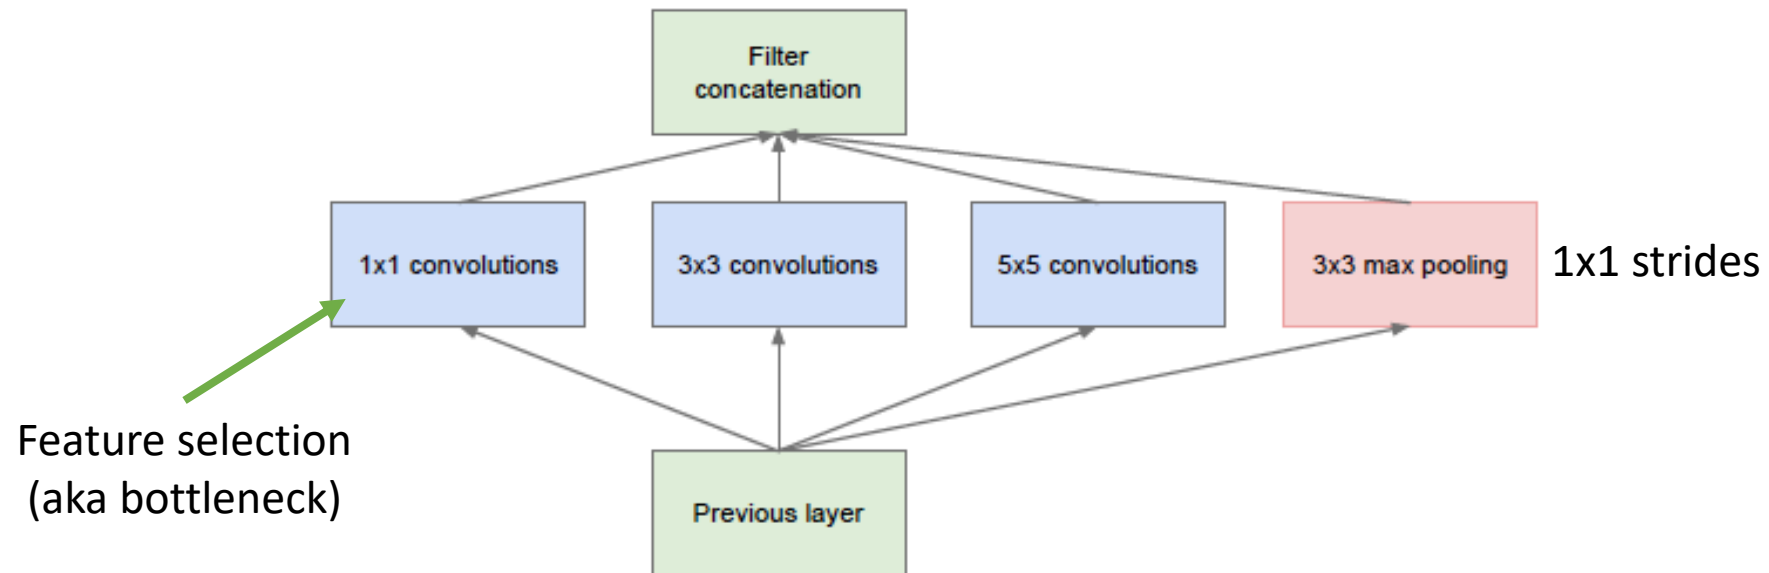
C. Szegedy *et al.*, “Going Deeper with Convolutions”, *CVPR*, 2015



# Architectures and basic design concepts

## □ GoogLeNet

Theoretical inception module: multi-scale filter bank (remember multi-scale Gabor filter banks?)

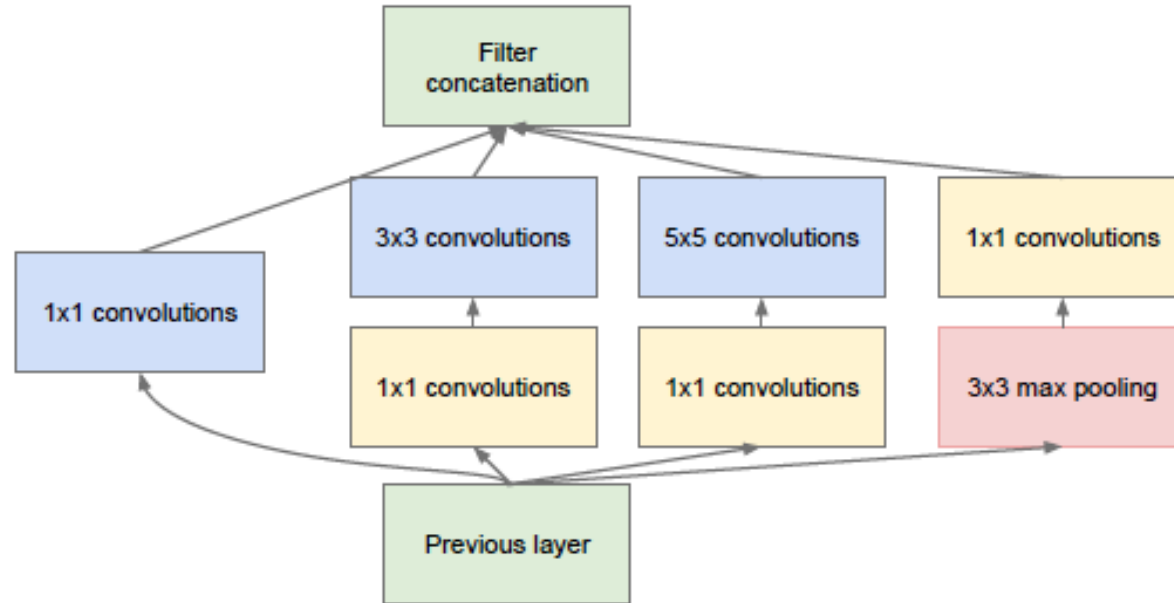


C. Szegedy *et al.*, "Going Deeper with Convolutions", *CVPR*, 2015

# Architectures and basic design concepts

## □ GoogLeNet

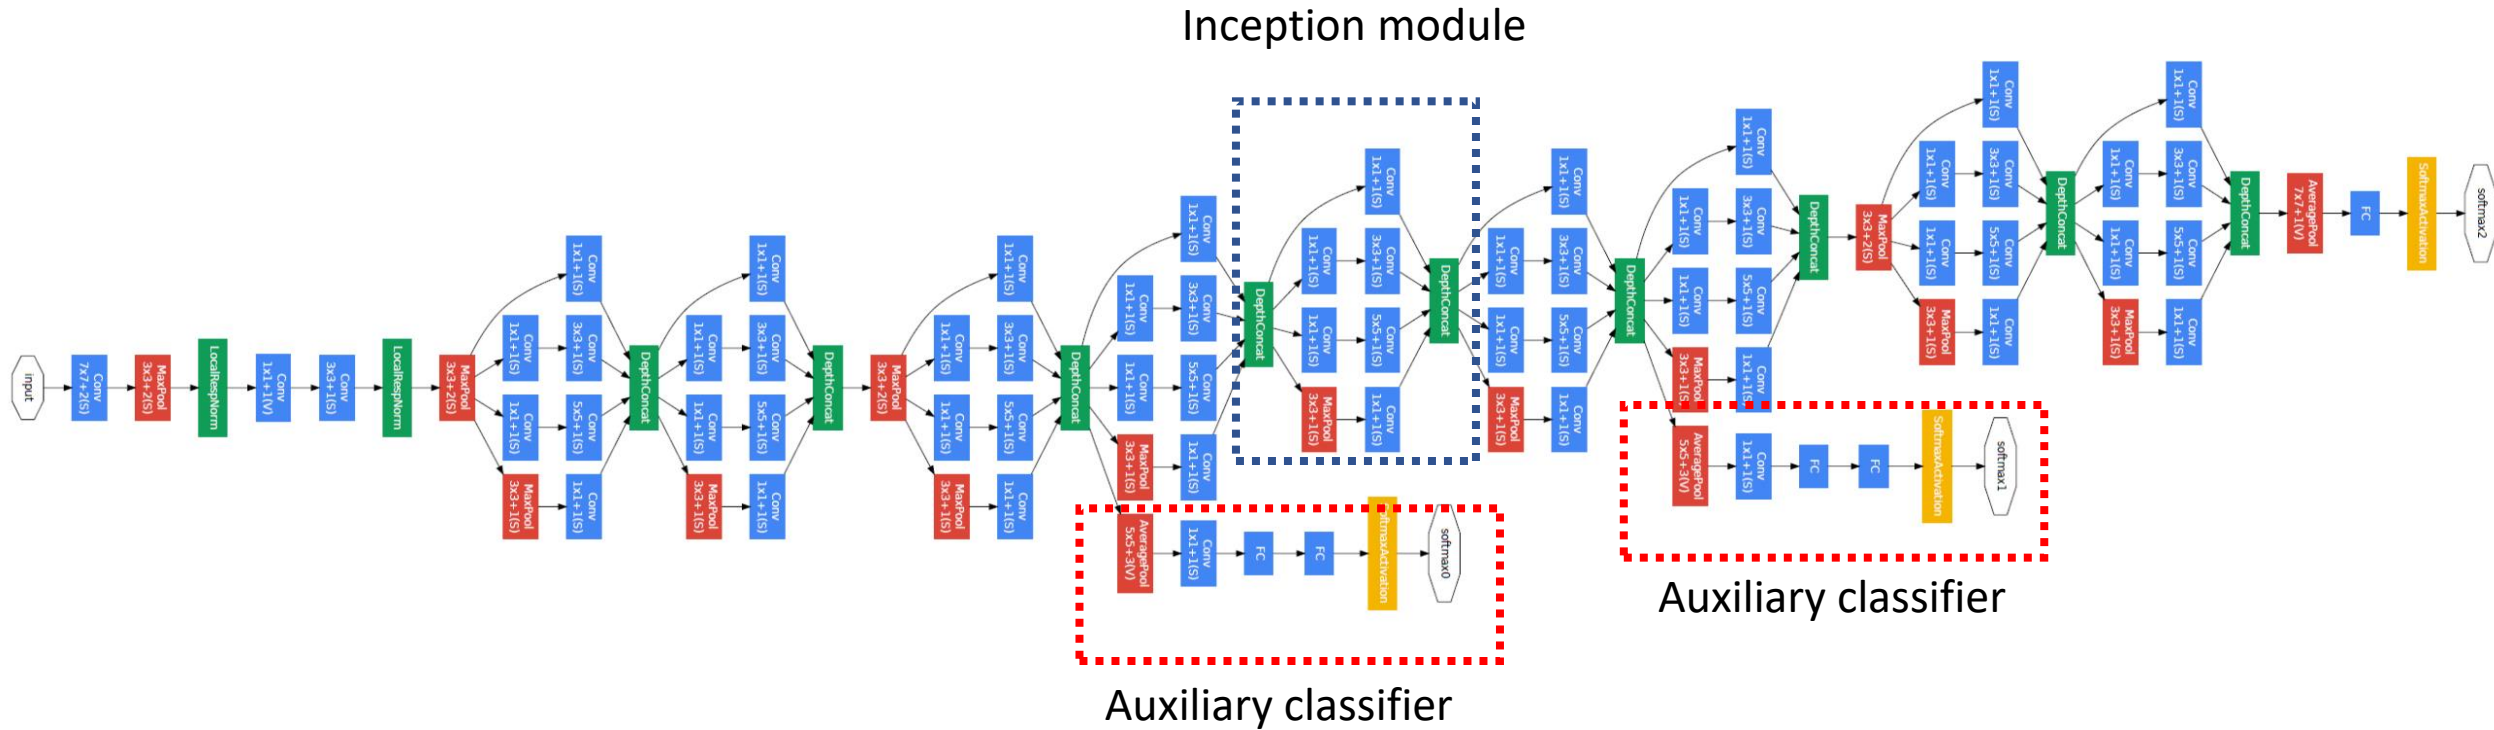
Inception and dimensionality reduction in GoogLeNet



C. Szegedy *et al.*, “Going Deeper with Convolutions”, *CVPR*, 2015

# Architectures and basic design concepts

## □ GoogLeNet

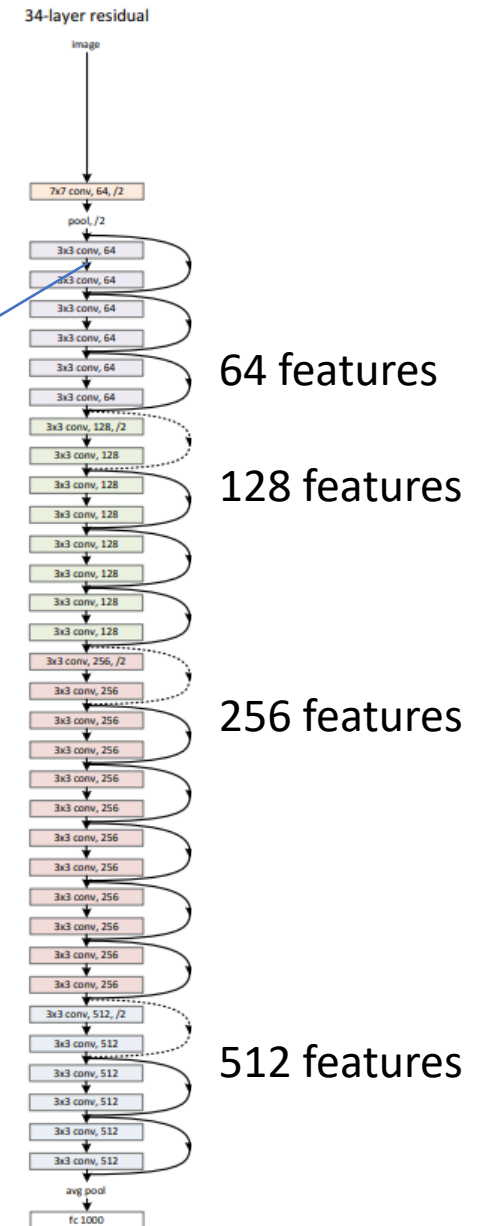
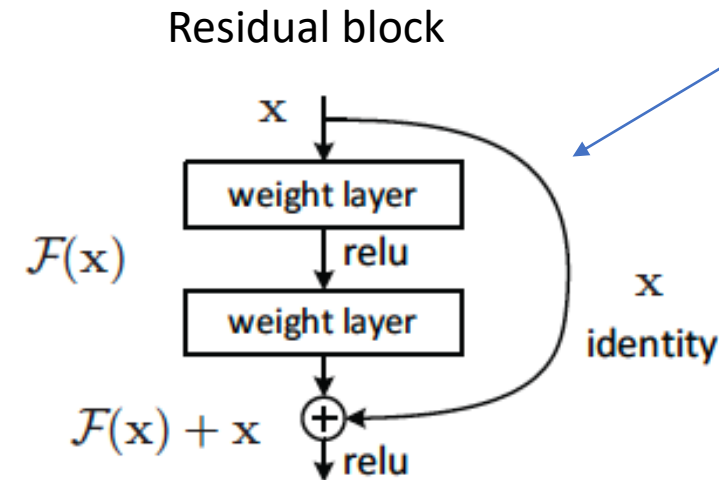


C. Szegedy *et al.*, “Going Deeper with Convolutions”, *CVPR*, 2015

# Architectures and basic design concepts

## □ ResNet

- Deeper networks highly suffer from **vanishing gradient problem**
- Residual blocks allows backpropagation of gradients without vanishing further
- Resnet also showed that it may be optimal to double the number of features as the dimensions halve in deeper layers

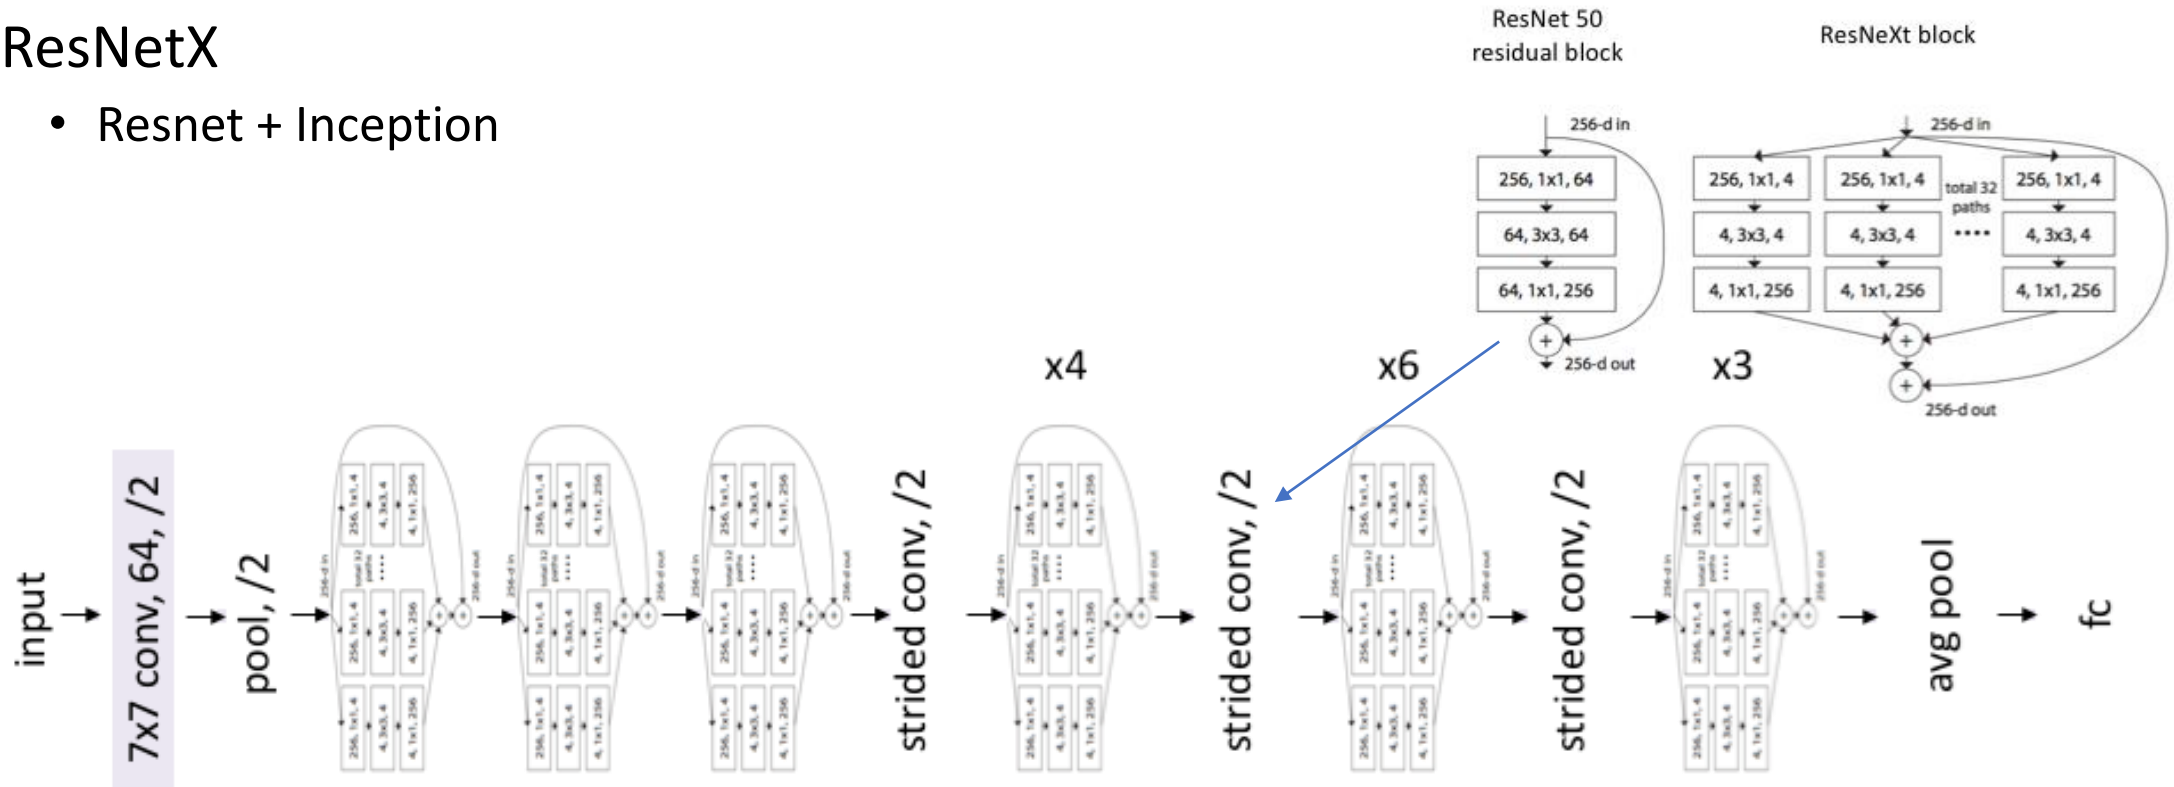


K. He, *et al.*, “Deep Residual Learning for Image Recognition”, *CVPR*, 2016

# Architectures and basic design concepts

## □ ResNetX

- Resnet + Inception

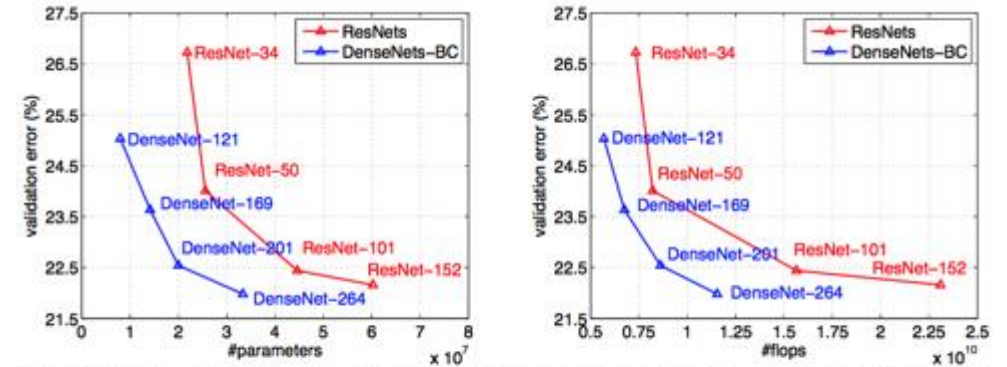


S. Xie, *et al.*, “Aggregated Residual Transformations for Deep Neural Networks”, *CVPR*, 2017

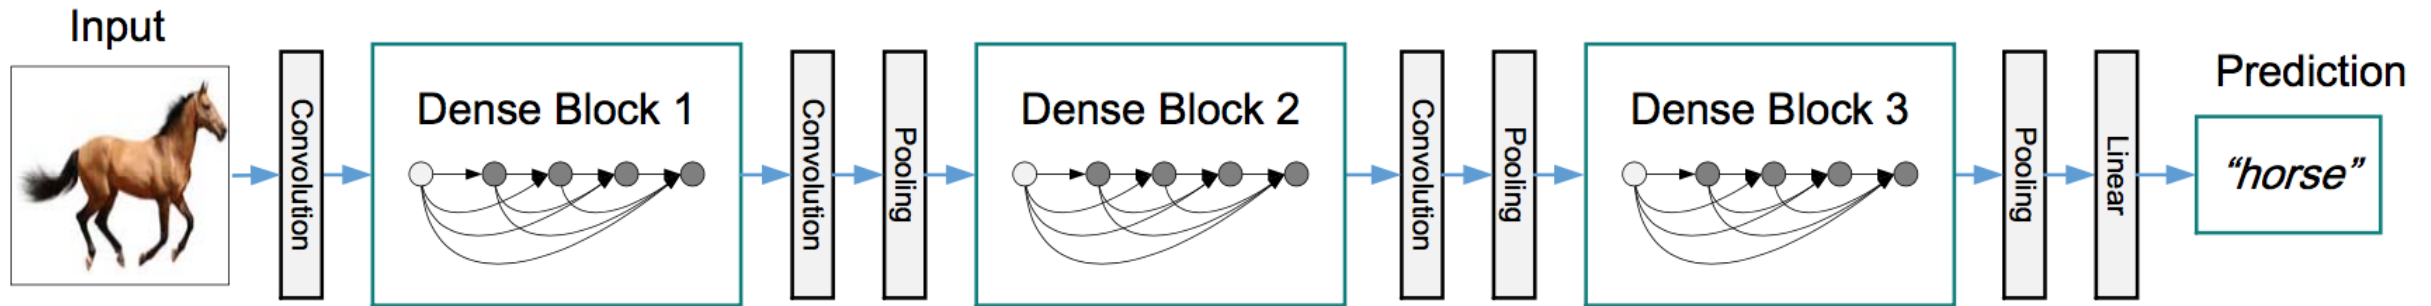
# Architectures and basic design concepts

## □ DenseNet

- Pass all residuals to all layers in every block
- Or how to take Resnet to the extreme...



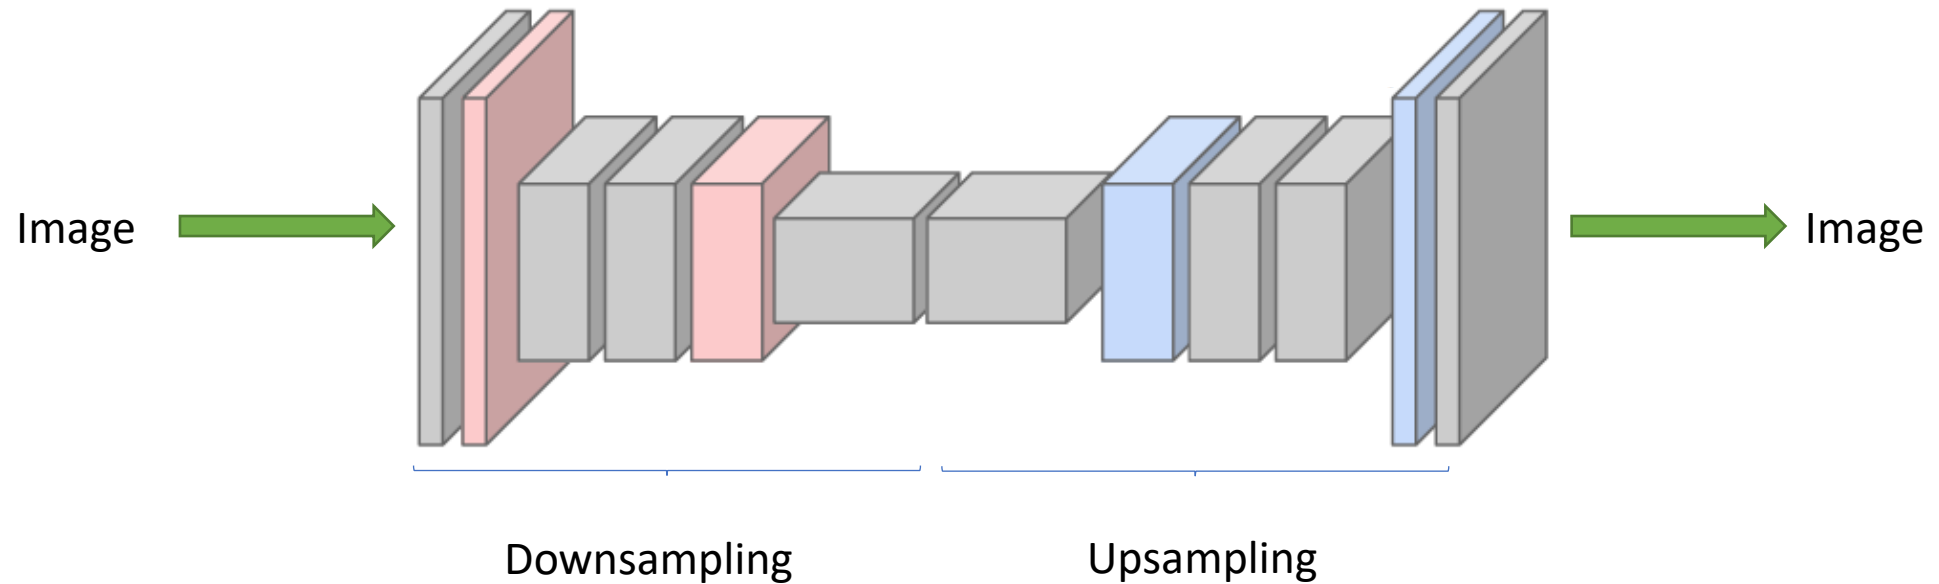
**Figure 3:** Comparison of the DenseNets and ResNets top-1 error rates (single-crop testing) on the ImageNet validation dataset as a function of learned parameters (*left*) and FLOPs during test-time (*right*).



G. Huang, *et al.*, “Densely Connected Convolutional Networks”, *CVPR*, 2017

# Fully convolutional networks

- Fully convolutional networks
  - Normally designed to create an output image



# Fully convolutional networks

- Fully convolutional networks
  - Unpooling

**Nearest Neighbor**

1	2
3	4



1	1	2	2
1	1	2	2
3	3	4	4
3	3	4	4

Input: 2 x 2

Output: 4 x 4

**“Bed of Nails”**

1	2
3	4



1	0	2	0
0	0	0	0
3	0	4	0
0	0	0	0

Input: 2 x 2

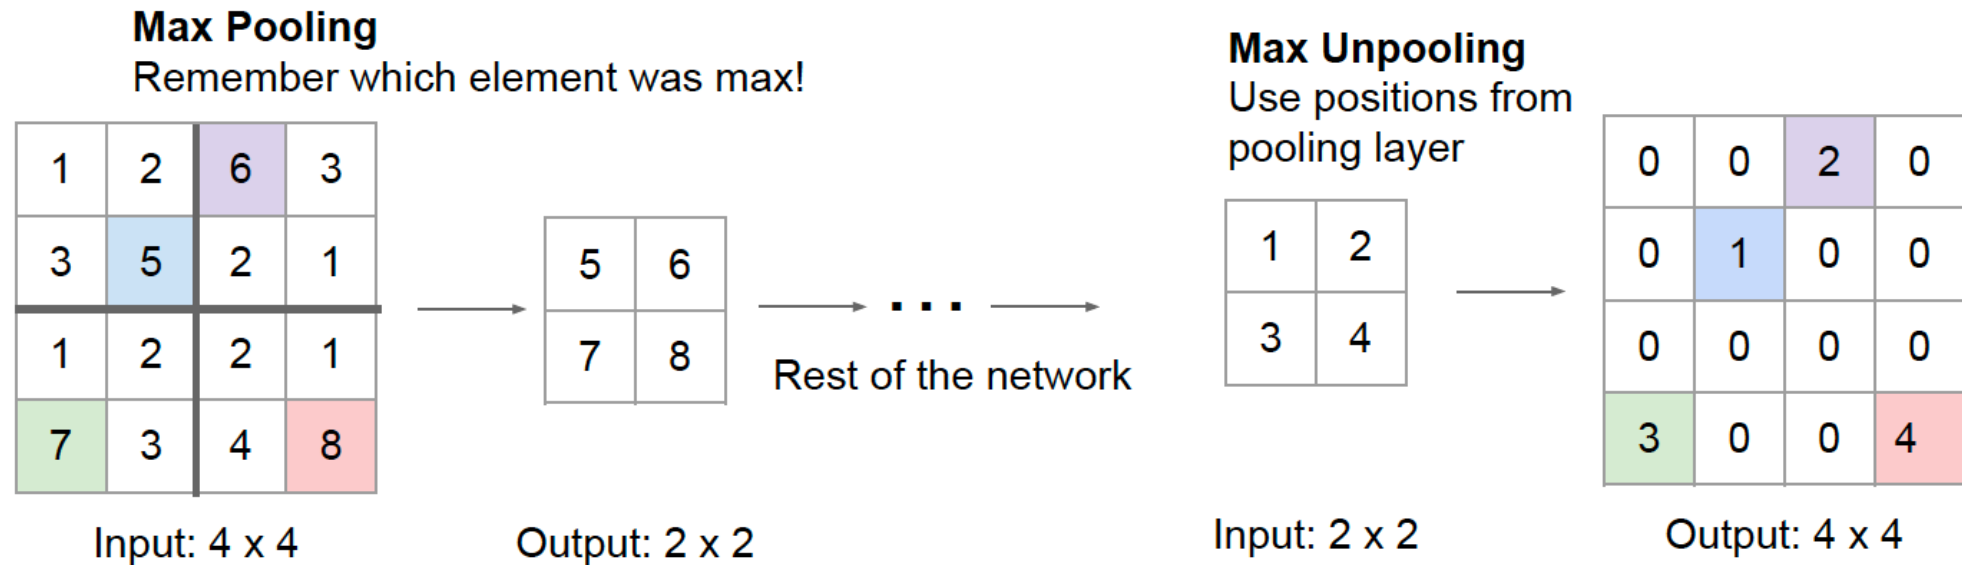
Output: 4 x 4



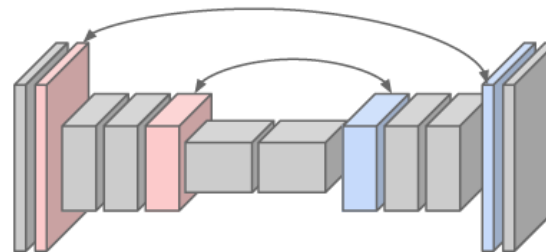
# Fully convolutional networks

## □ Fully convolutional networks

- Unpooling



Corresponding pairs of  
downsampling and  
upsampling layers

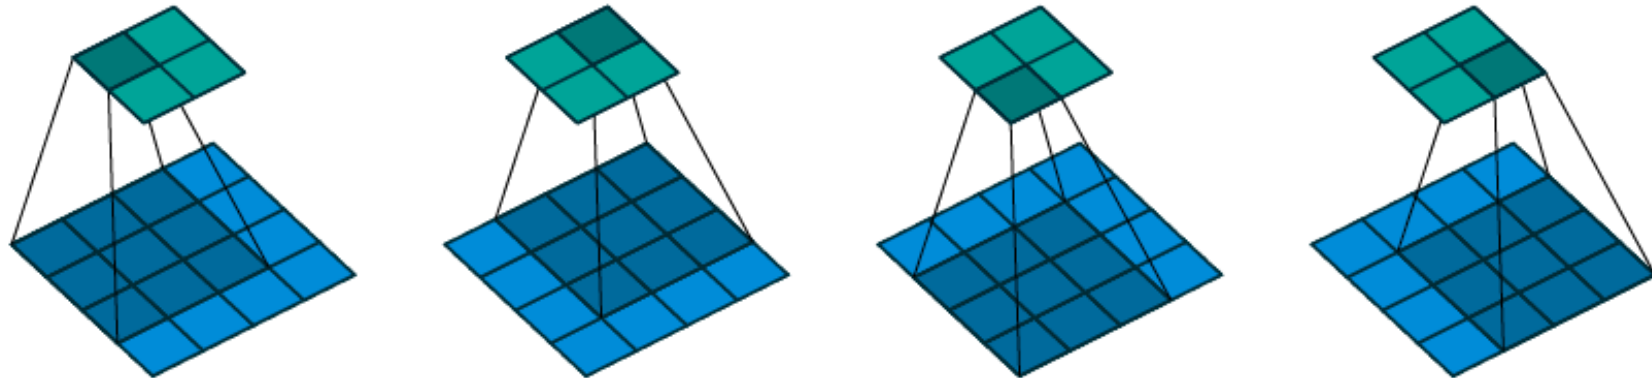


# Fully convolutional networks

## □ Fully convolutional networks

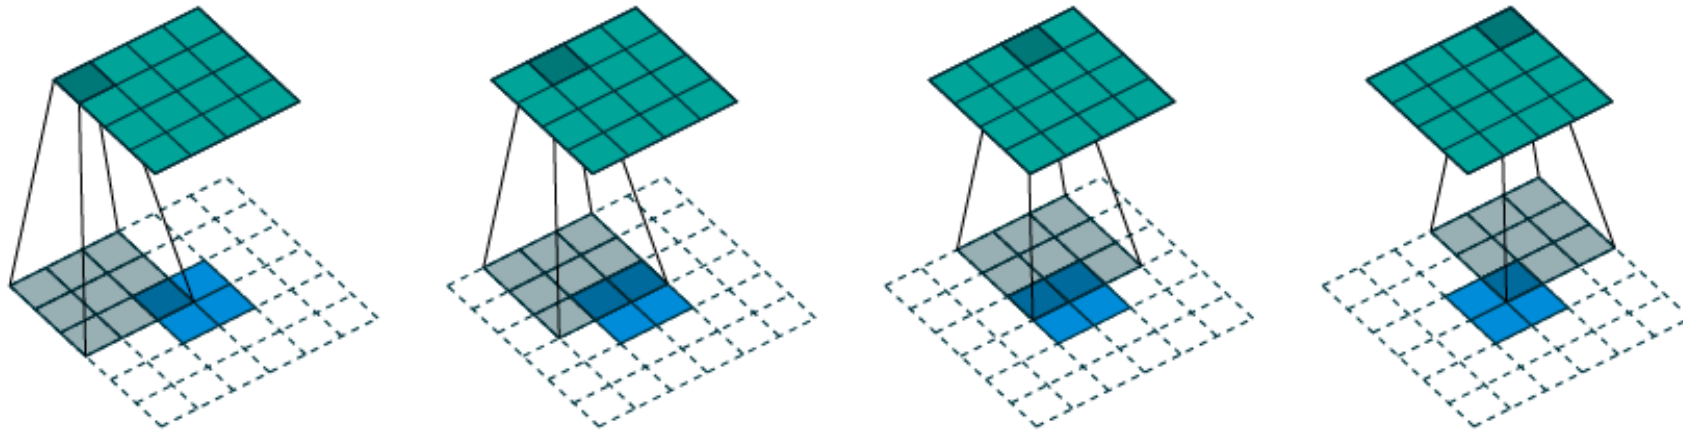
- Transposed convolution (or upconvolution)

Convolution



Input:  $4 \times 4$   
Kernel:  $3 \times 3$   
Stride: 1  
Padding: 0  
Output:  $2 \times 2$

Transposed convolution



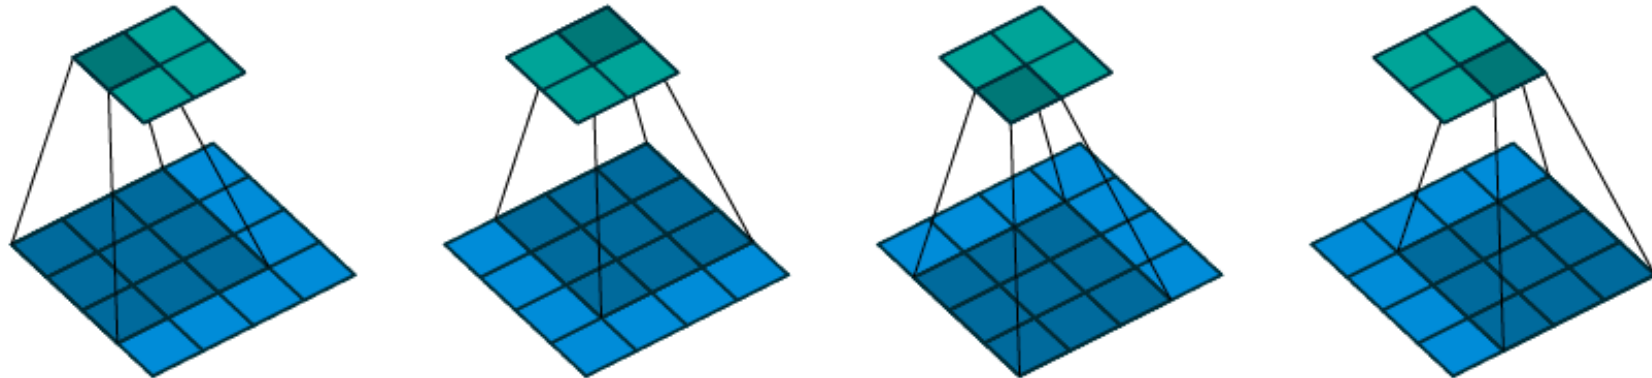
Input:  $2 \times 2$   
Kernel:  $3 \times 3$   
Stride: 1  
Padding: 2  
Output:  $4 \times 4$

# Fully convolutional networks

## □ Fully convolutional networks

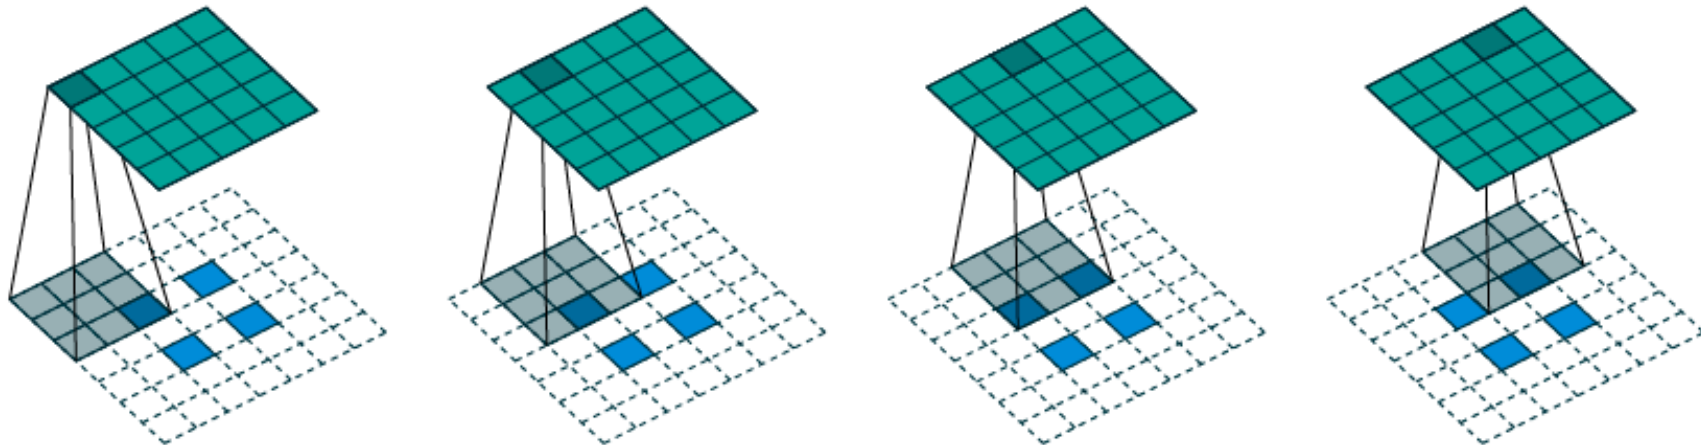
- Transposed convolution (or upconvolution)

Convolution



Input:  $4 \times 4$   
Kernel:  $3 \times 3$   
Stride: 1  
Padding: 0  
Output:  $2 \times 2$

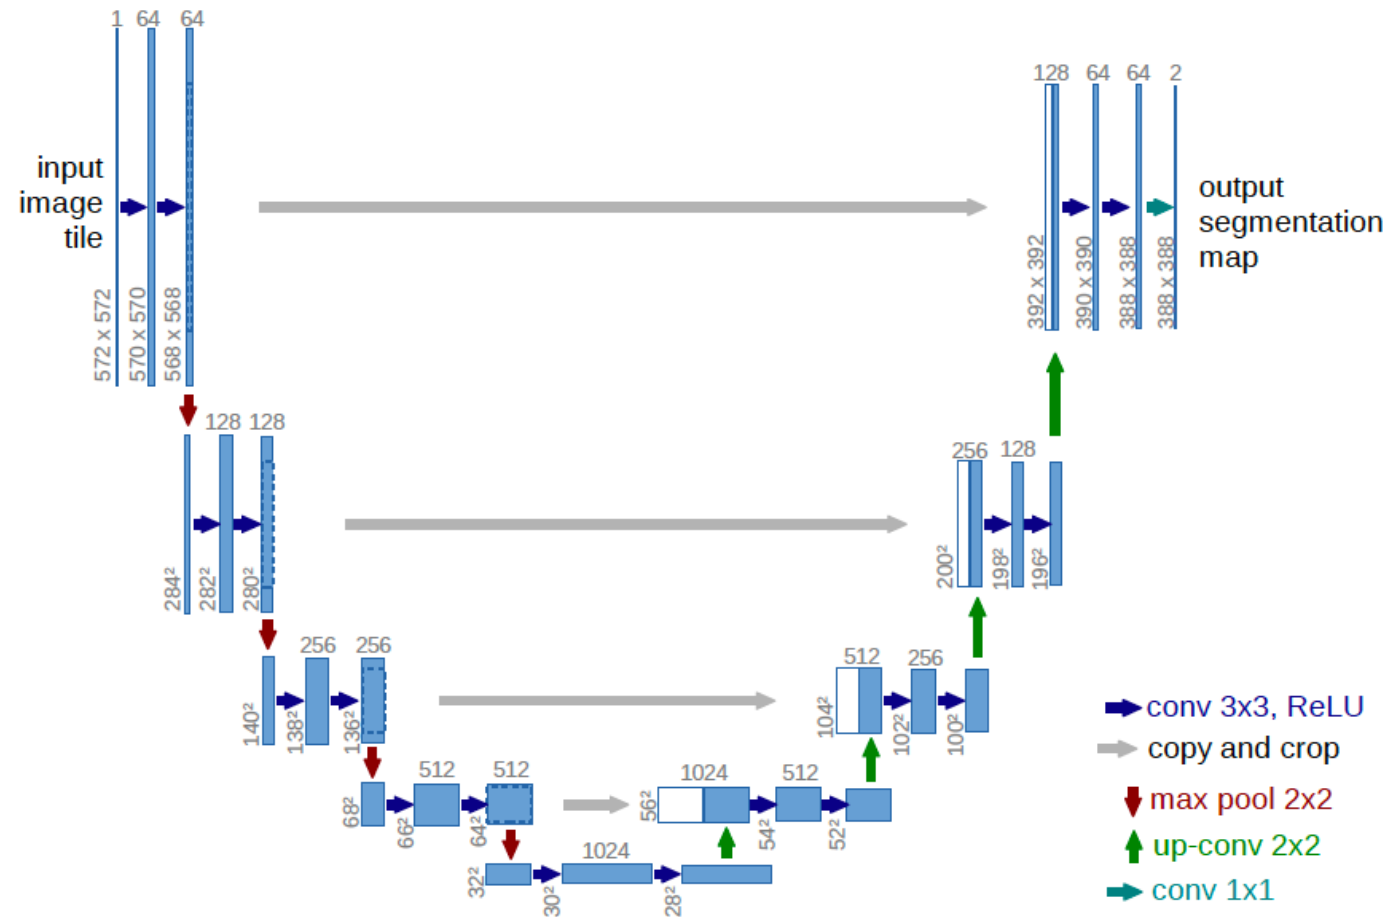
Transposed convolution



Input:  $2 \times 2$   
Kernel:  $3 \times 3$   
Stride: 1  
Padding: 2  
Output:  $5 \times 5$

# Fully convolutional networks

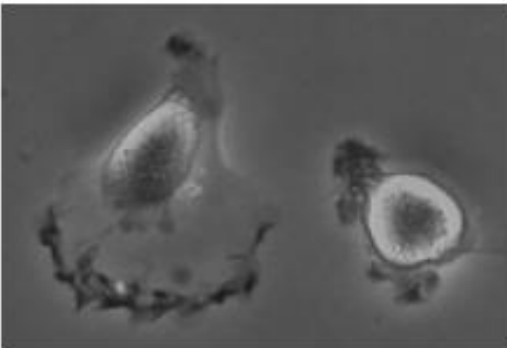
## □ U-net



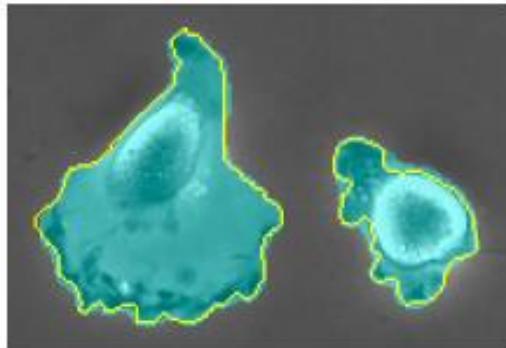
# Fully convolutional networks

## □ U-net

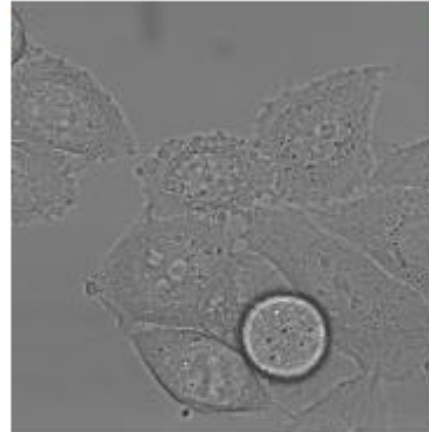
a



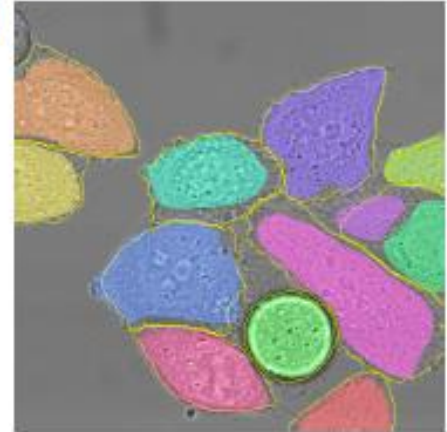
b



c

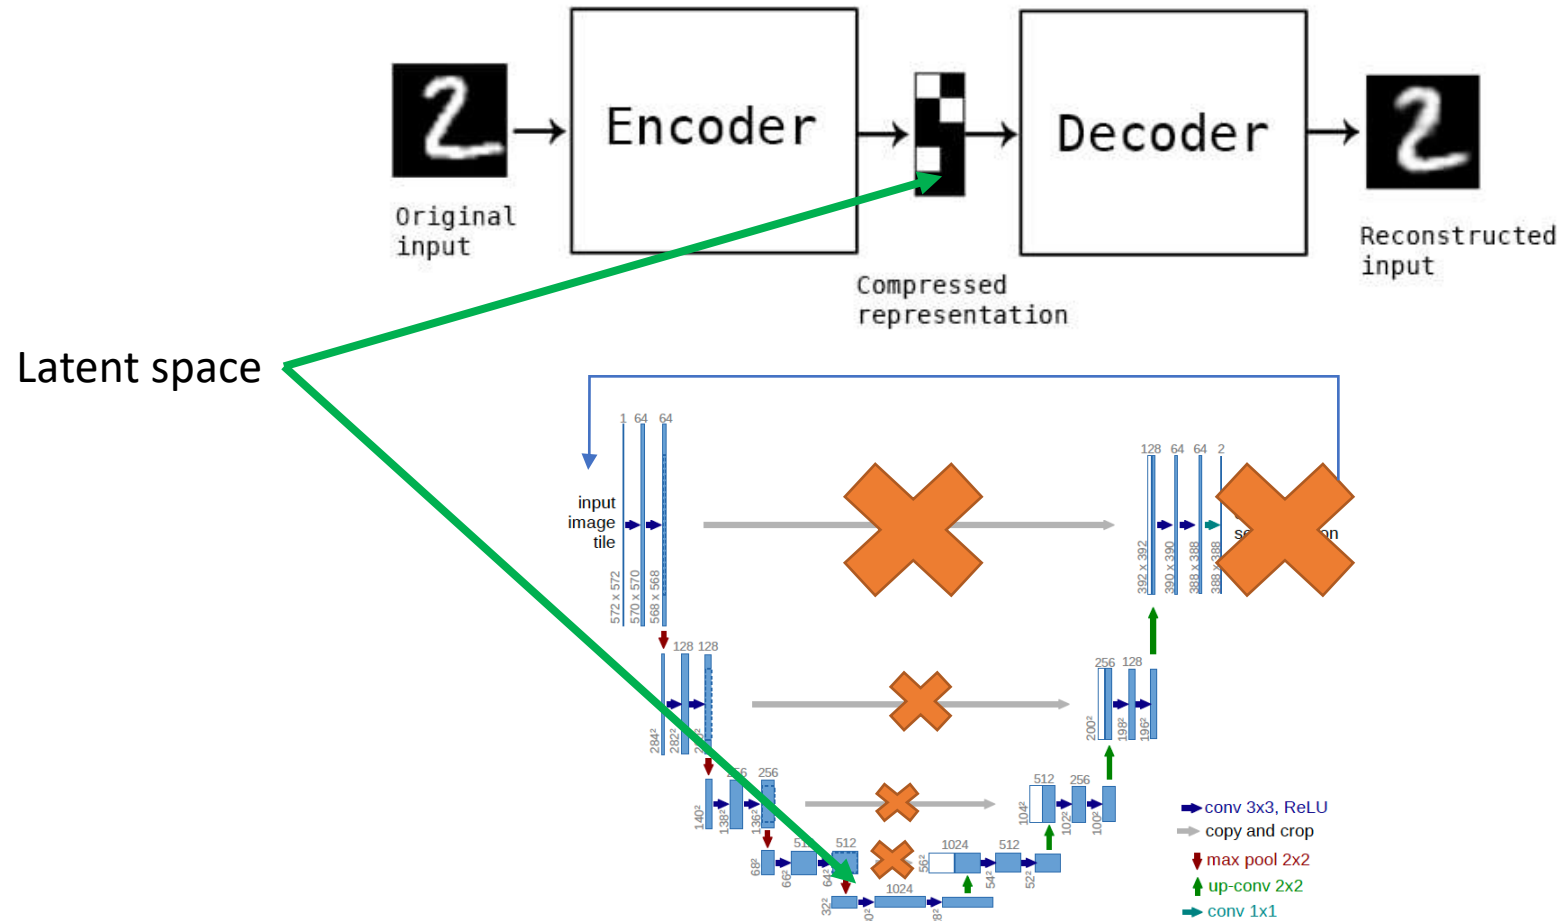


d



# Fully convolutional networks

## Autoencoders



# Fully convolutional networks

- Generative adversarial networks (GANs)
  - Or learning how to fake images

